

REACTIVE MICROSYSTEMS

THE EVOLUTION OF MICROSERVICES AT SCALE

JONAS BONER

@JBONER

WE HAVE BEEN SPOILED BY
THE ALMIGHTY
MONOLITH





KNOCK. KNOCK. WHO'S THERE?

R
P
A
L
I
T
Y

Yesterday	Today
Single machines	Clusters of machines
Single core processors	Multicore processors
Expensive RAM	Cheap RAM
Expensive disk	Cheap disk
Slow networks	Fast networks
Few concurrent users	Lots of concurrent users
Small data sets	Large data sets
Latency in seconds	Latency in milliseconds

A black and white historical photograph captures a scene from the early 20th century. In the foreground, a large, multi-spoked wheel of a horse-drawn carriage is prominent. A man, dressed in a dark jacket and a flat cap, sits in the driver's seat of the carriage, looking towards the right. Behind him, a dark-colored horse is harnessed to the carriage, its head turned slightly forward. The carriage appears to be carrying some goods or equipment. The background features a row of brick buildings with arched windows and decorative cornices. Bare trees stand between the buildings and the street. The overall atmosphere is one of a bygone era.

WE CAN'T MAKE THE HORSE FASTER

A black and white historical photograph of a man driving a vintage car. The man is wearing a flat cap and a dark, high-collared coat. He is seated in the driver's seat of a large, open-top vehicle with spoked wheels and a prominent front grille. The car appears to be from the early 20th century. In the background, there are several large wooden shipping crates stacked in a yard, with some text visible on them such as "DETROIT MOTOR CAR COMPANY" and "NO 18".

WE NEED CARS FOR WHERE WE ARE GOING

BUT DON'T JUST DRINK THE

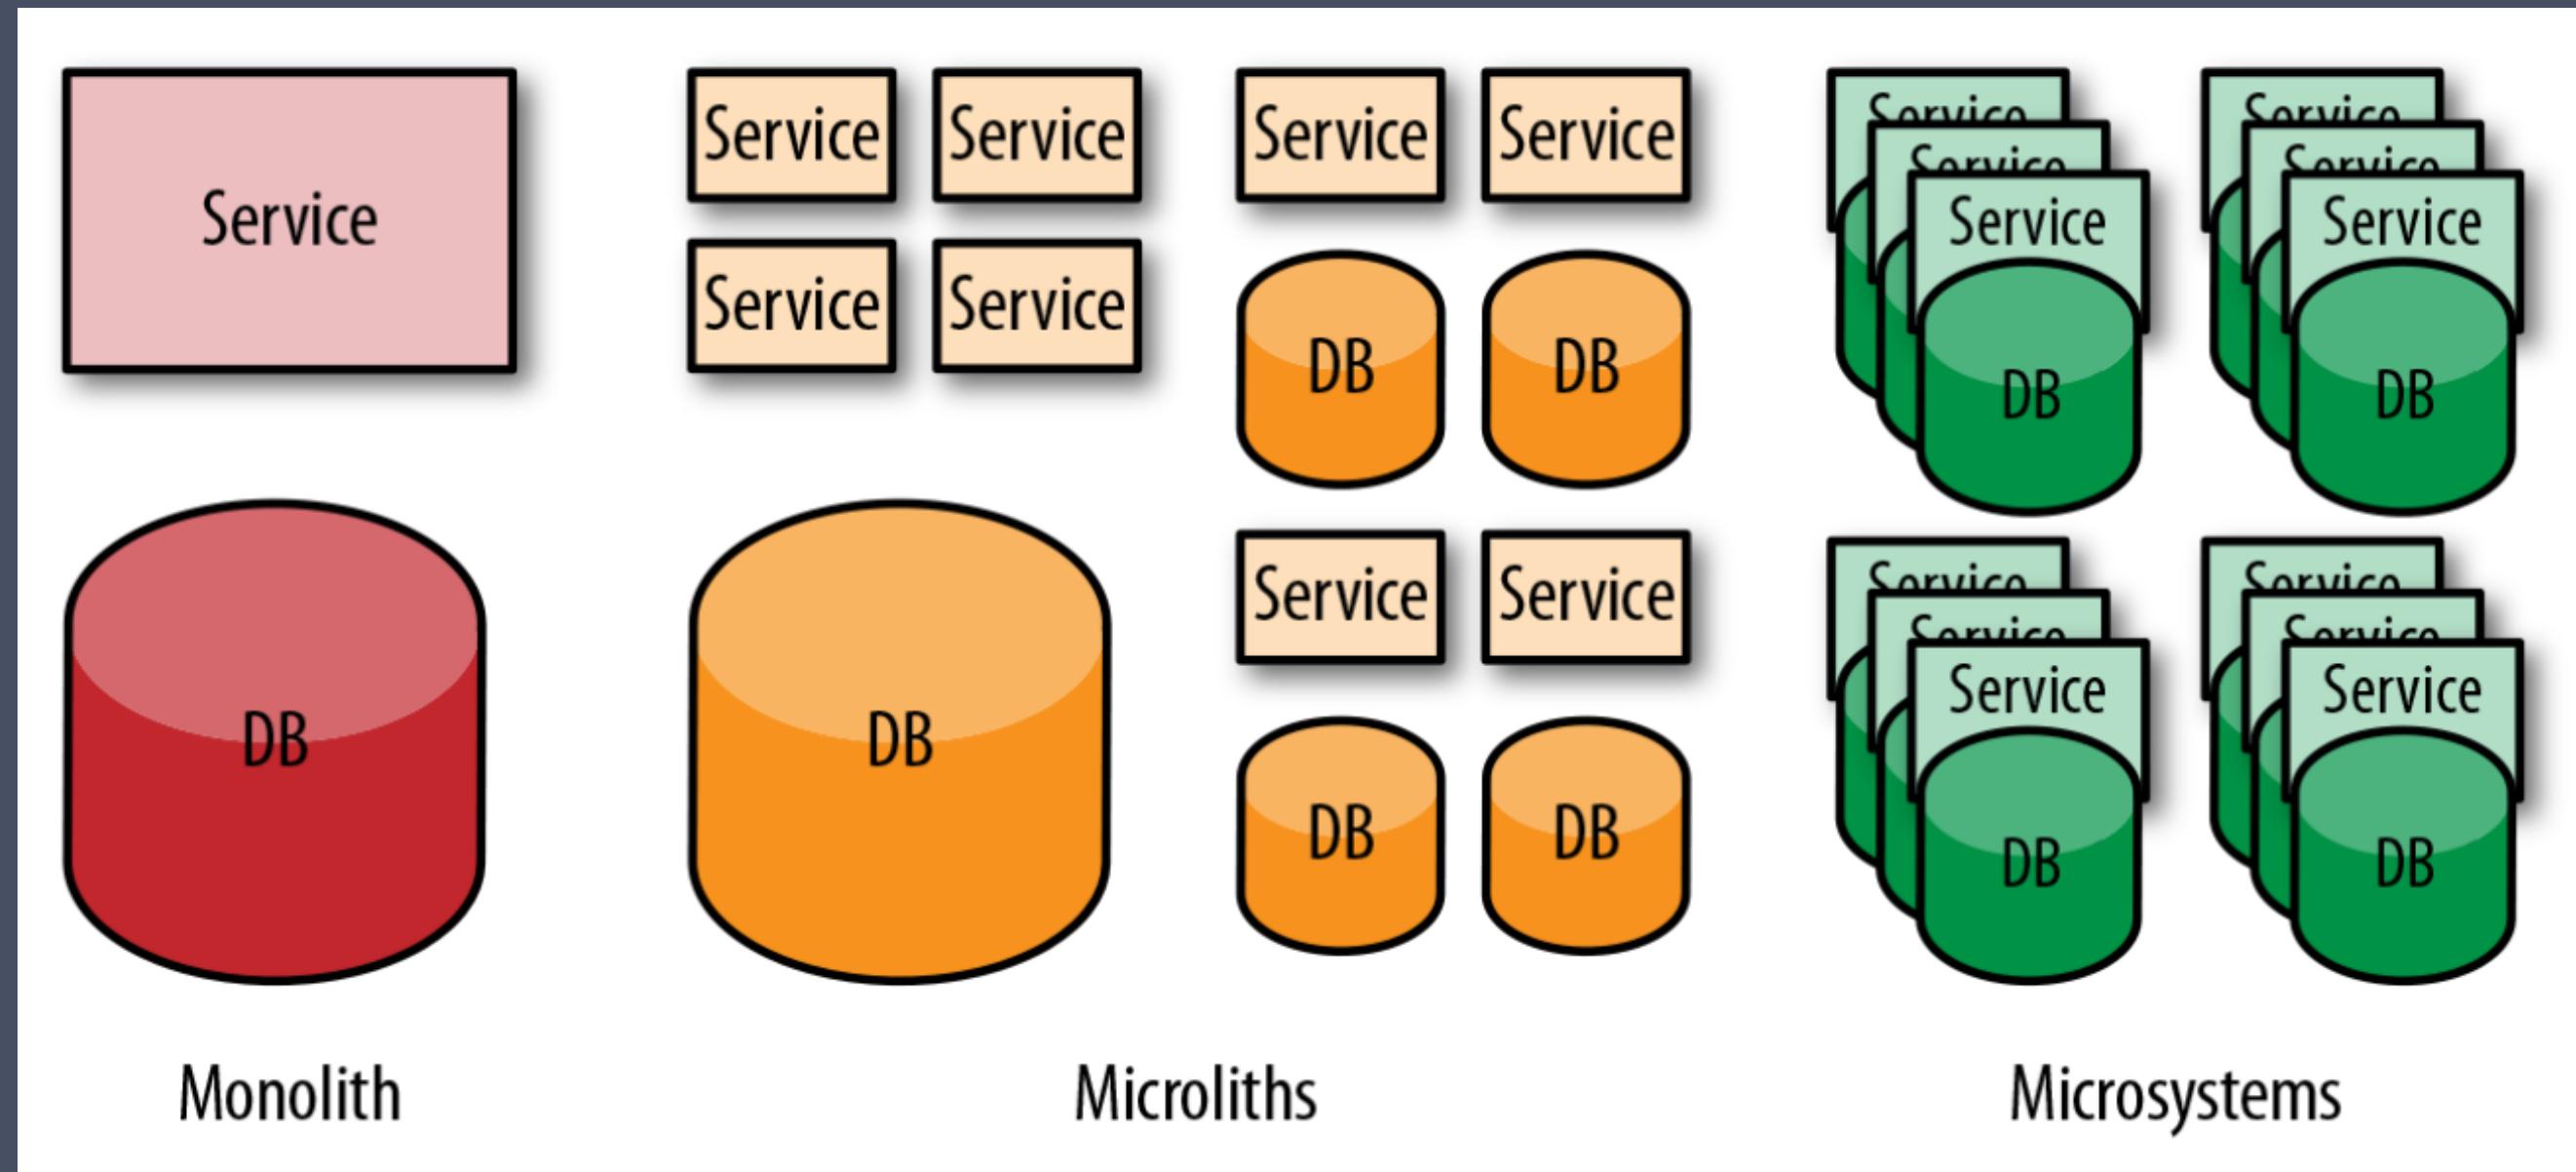


THINK FOR YOURSELF

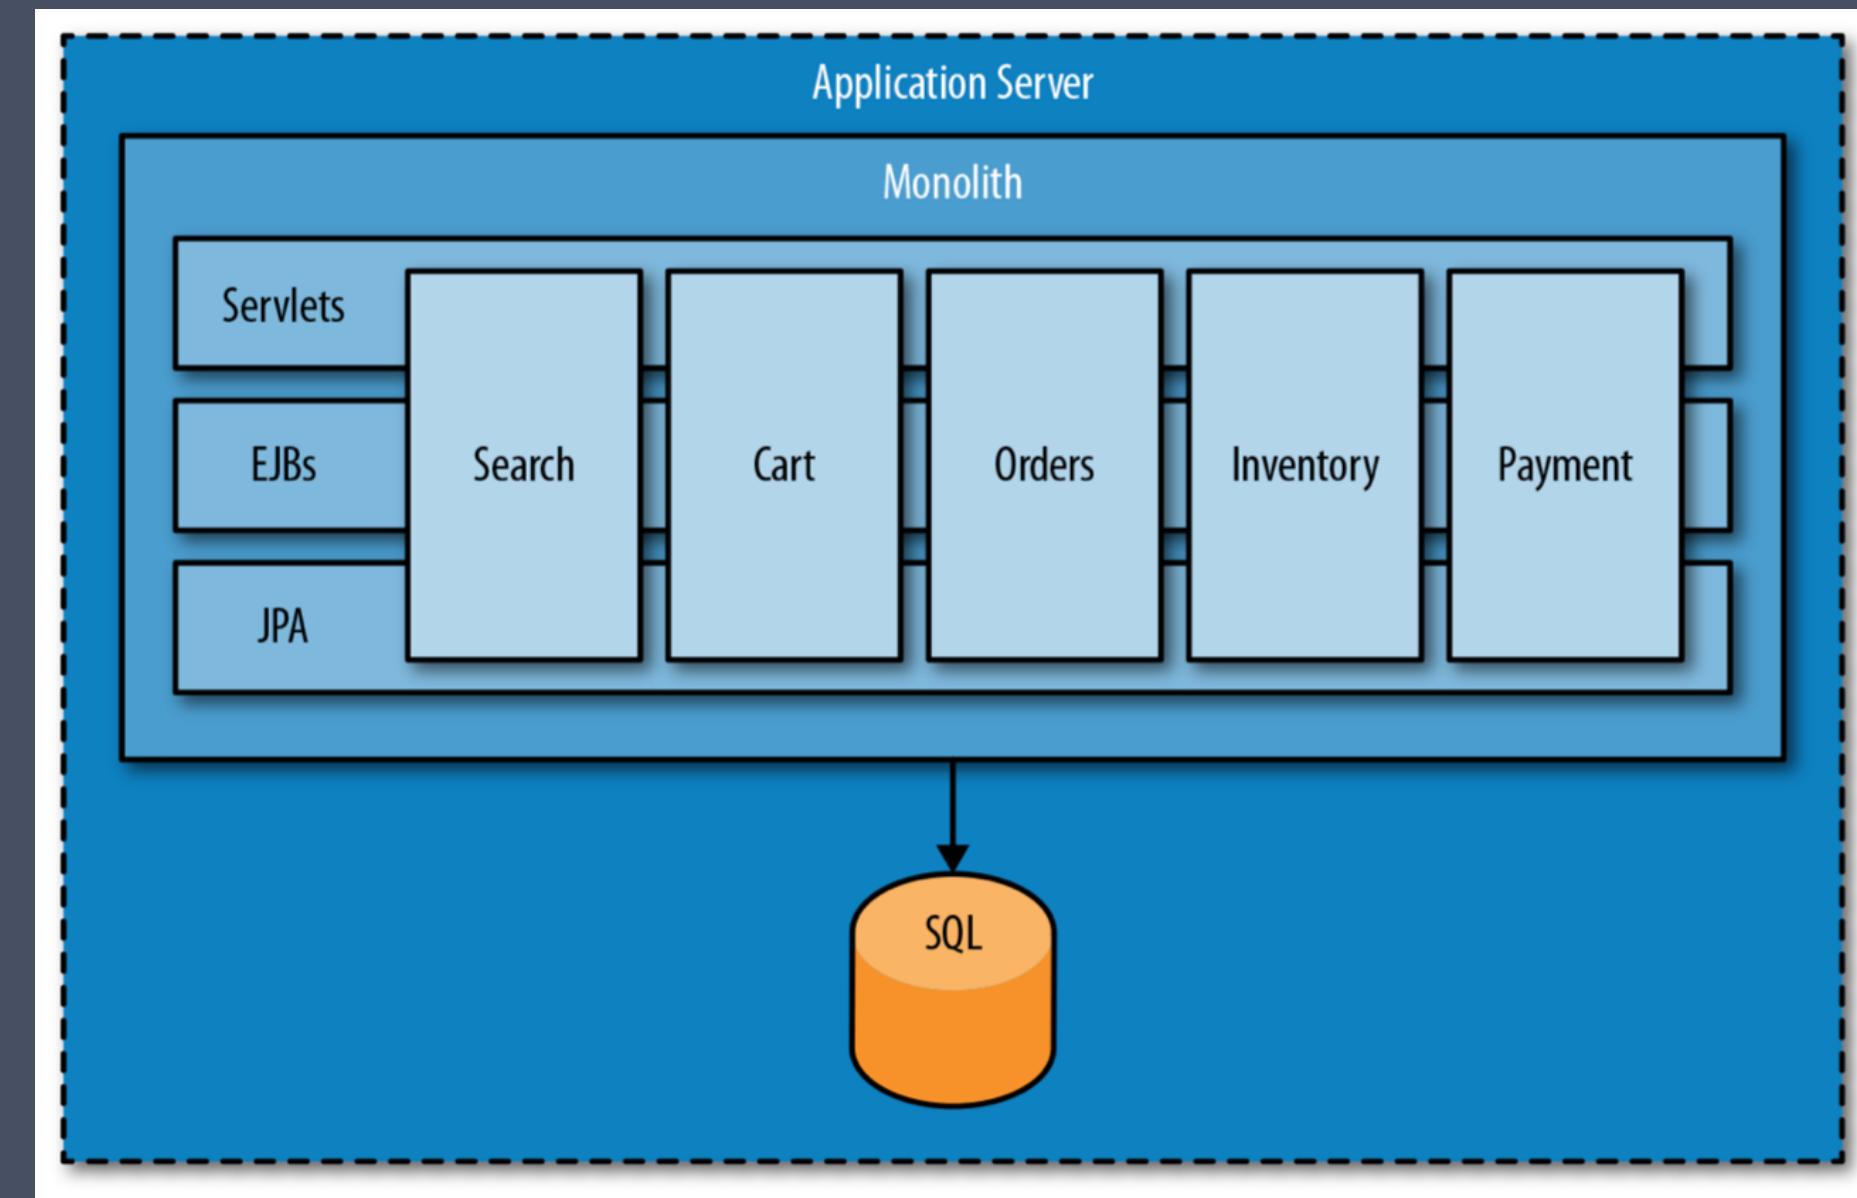
NO ONE WANTS
MICROSERVICES
IT'S A NECESSARY EVIL

ARCHITECTURAL CONTEXT OF MICROSERVICES: DISTRIBUTED SYSTEMS

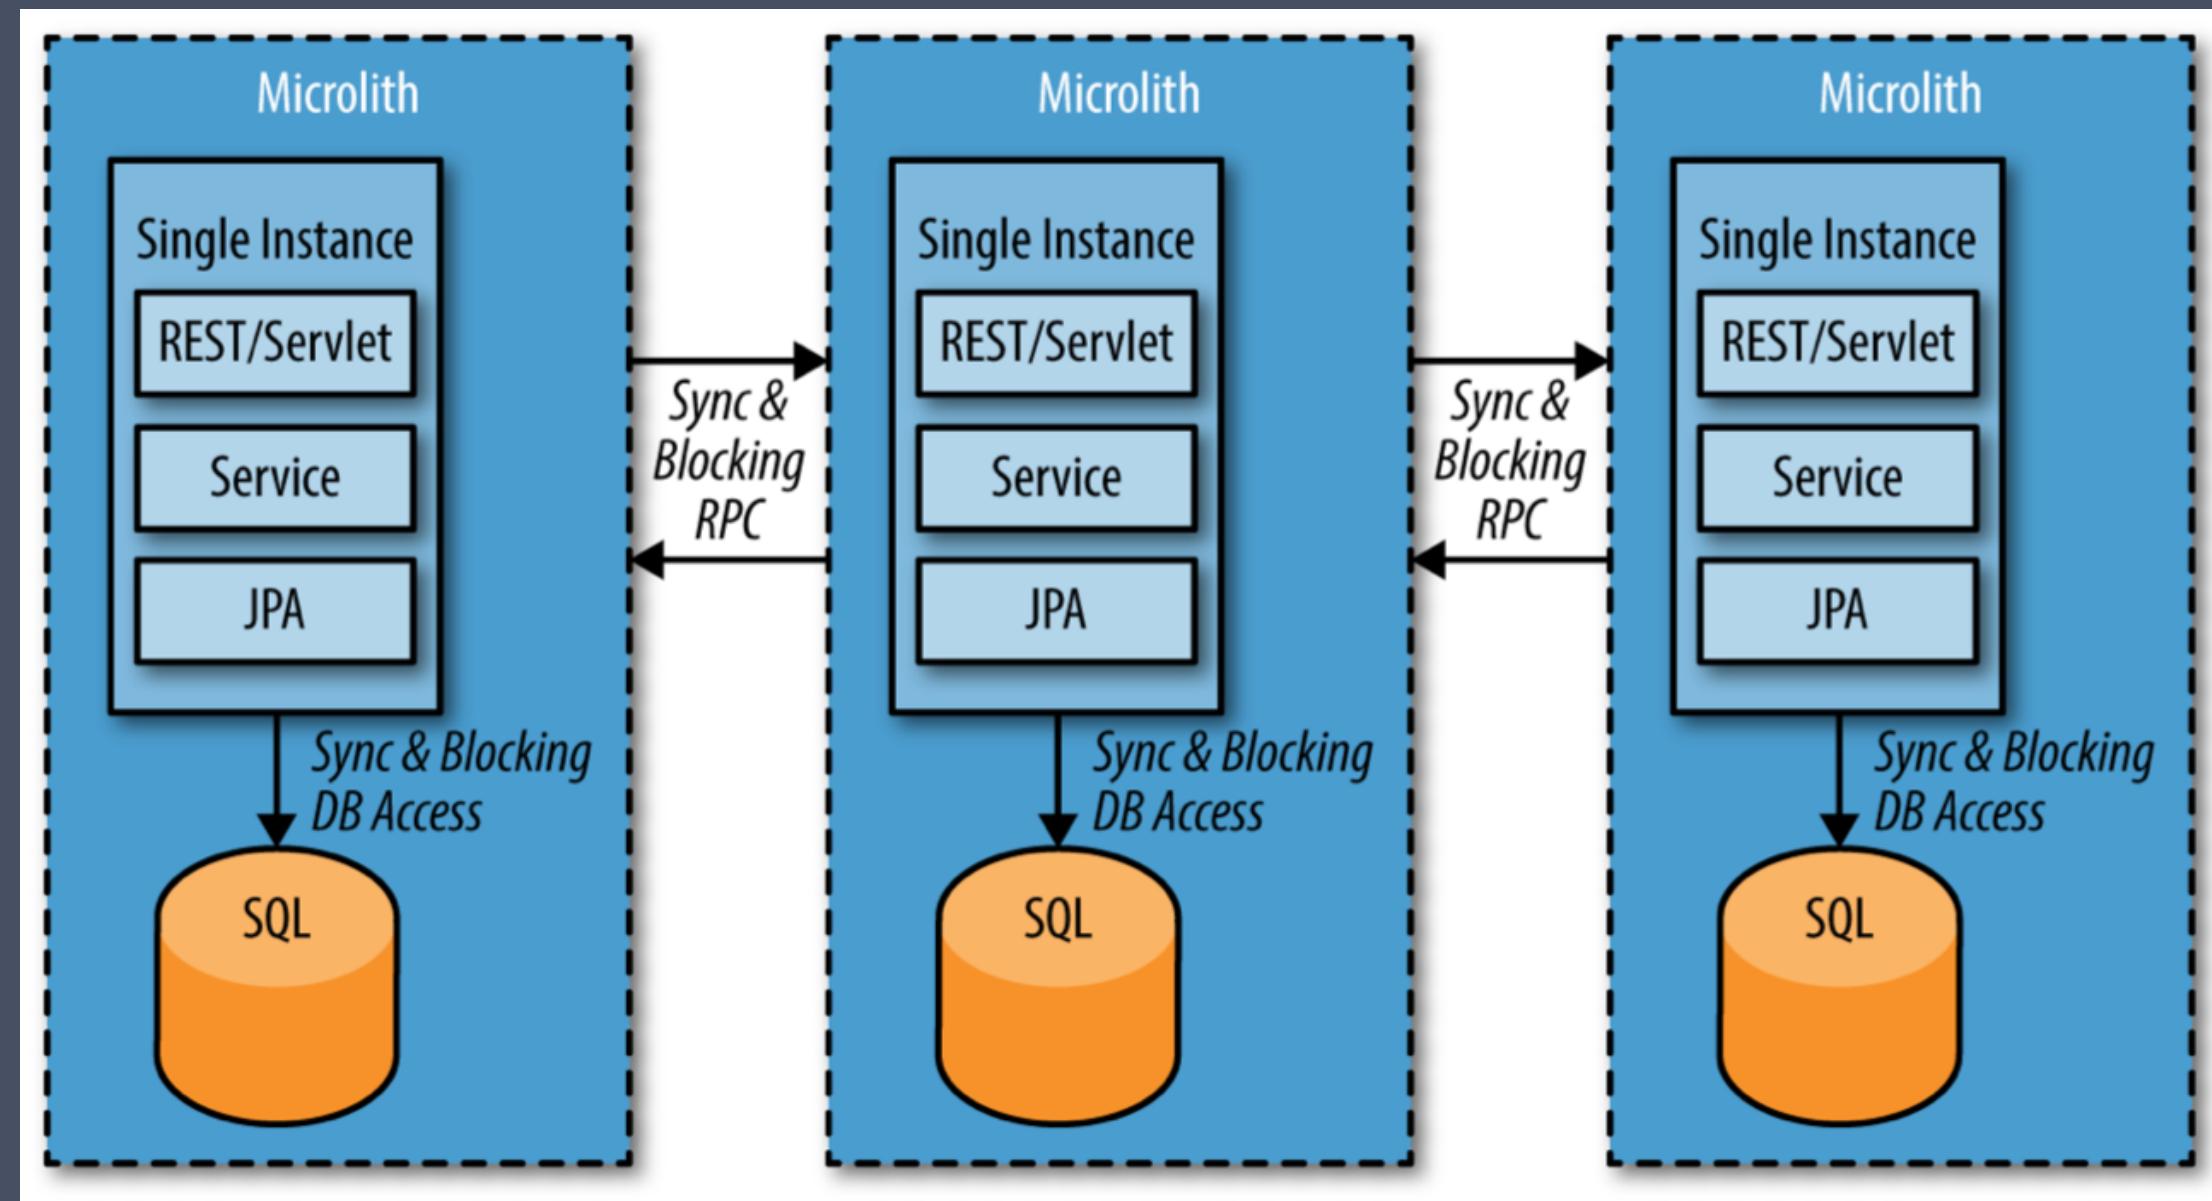
TODAY'S JOURNEY



LET'S SAY THAT WE WANT TO SLICE THIS MONOLITH UP



TOO MANY END UP WITH AN ARCHITECTURE LIKE THIS



MICROLITH: SINGLE INSTANCE MICROSERVICE

- NOT RESILIENT
- NOT ELASTIC



ONE ACTOR IS NO ACTOR.
ACTORS COME IN SYSTEMS.

- CARL HEWITT



MICROSERVICES
COME IN SYSTEMS

3 HELPFUL TOOLS

1. EVENTS-FIRST DDD
2. REACTIVE DESIGN
3. EVENT-BASED PERSISTENCE

PRACTICE

EVENTS-FIRST
DOMAIN-DRIVEN DESIGN

DON'T FOCUS ON THE THINGS

THE NOUNS
THE DOMAIN OBJECTS

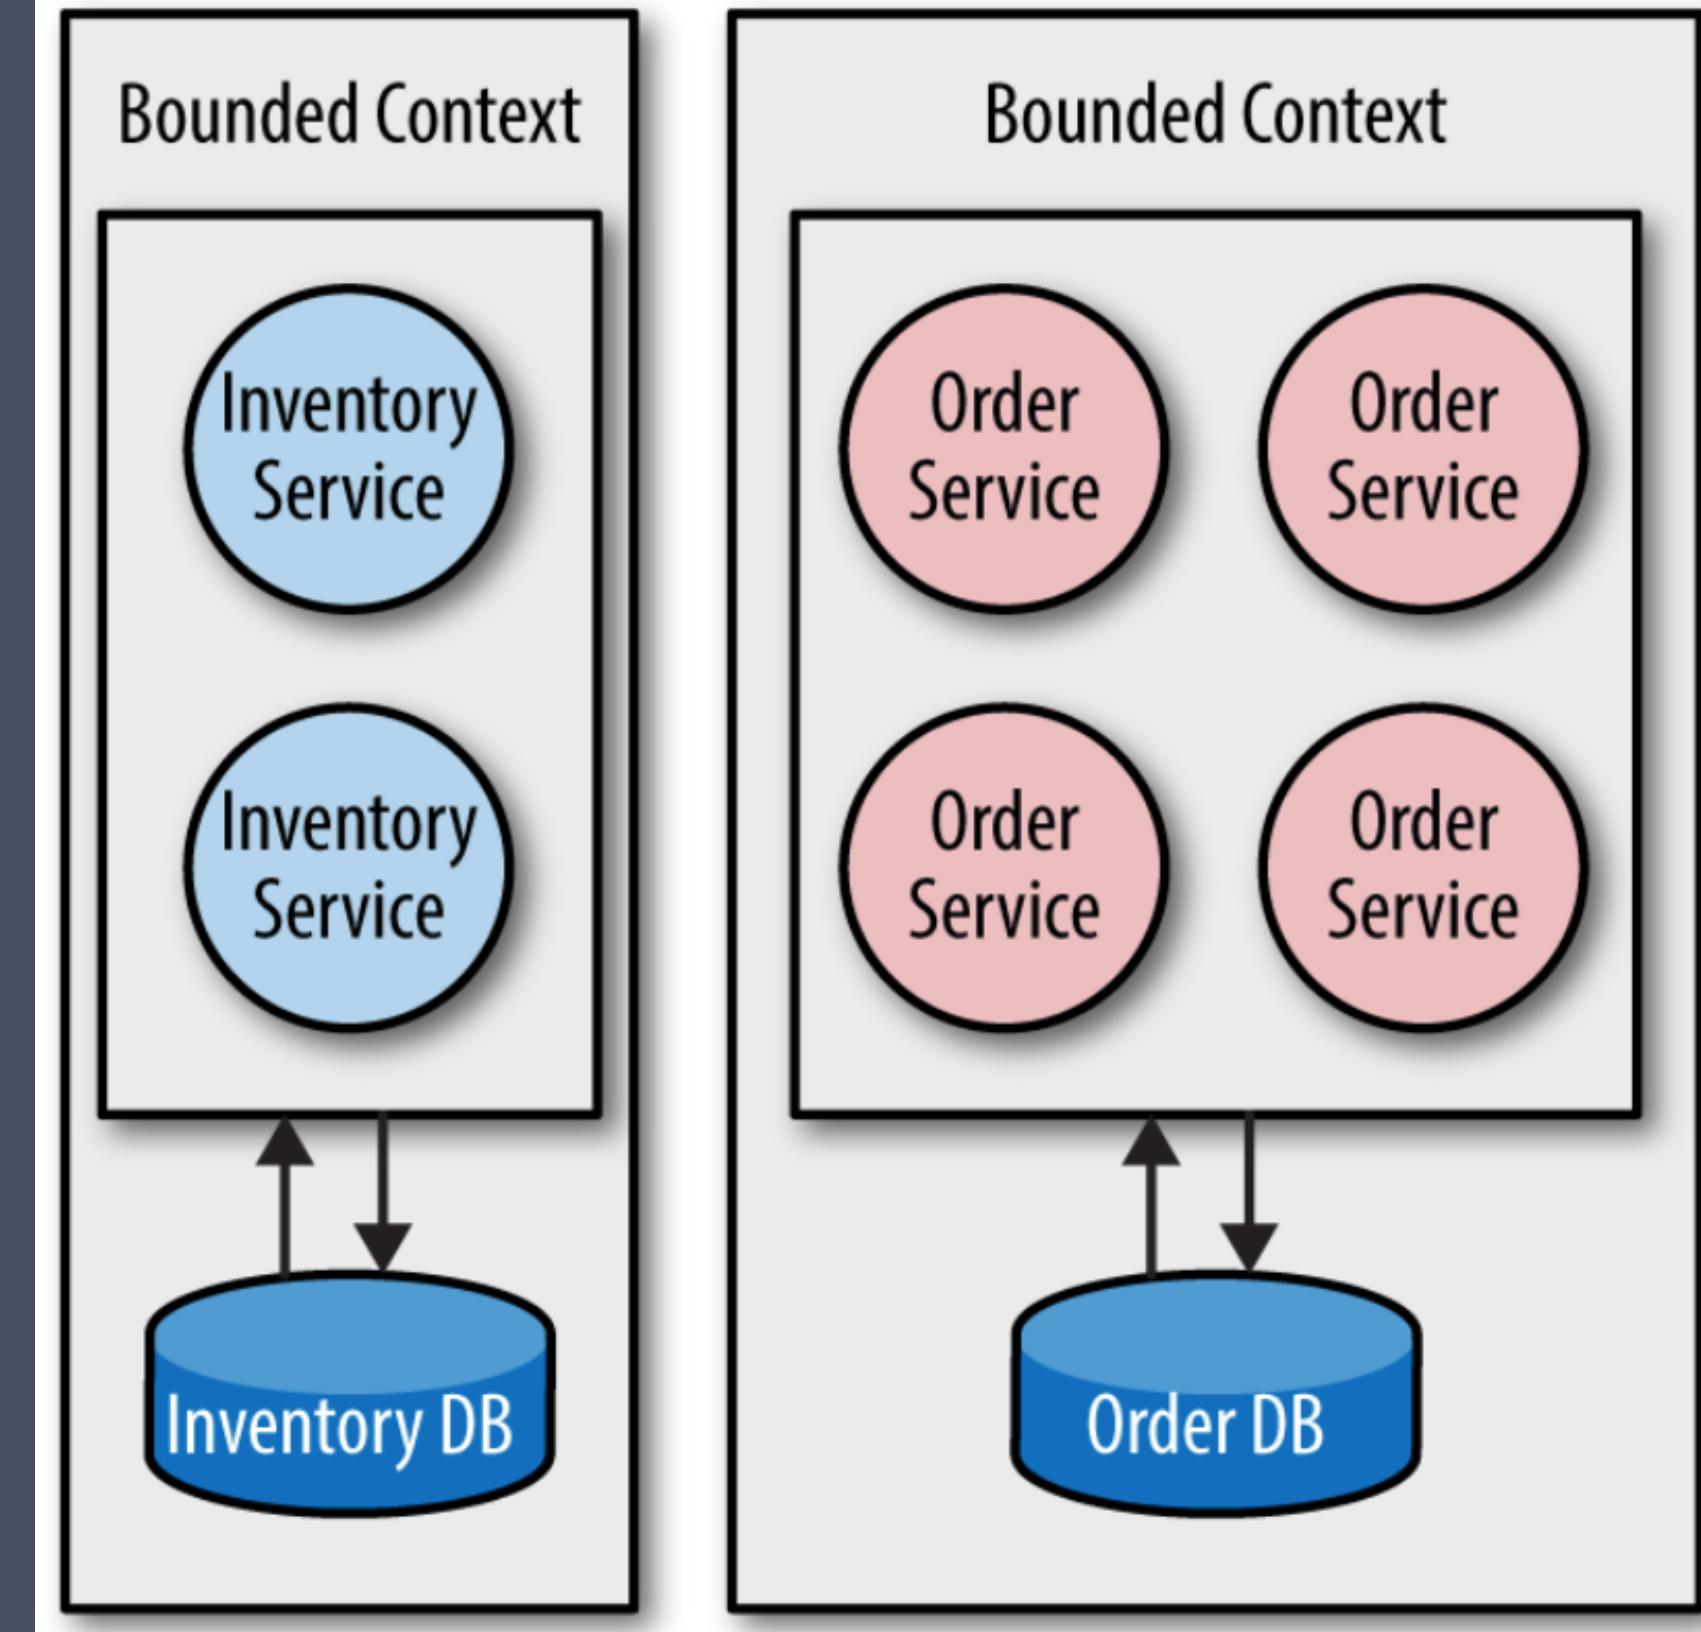
FOCUS ON WHAT HAPPENS

THE VERBS
THE EVENTS

WHEN YOU START MODELING EVENTS.
IT FORCES YOU TO THINK ABOUT THE
BEHAVIOR OF THE SYSTEM.
AS OPPOSED TO THINKING ABOUT
STRUCTURE INSIDE THE SYSTEM.

- GREG YOUNG

LET THE
EVENTS DEFINE
THE BOUNDED CONTEXT



EVENTS REPRESENT FACTS

TO CONDENSE FACT FROM
THE VAPOR OF NUANCE

- NEAL STEPHENSON. SNOW CRASH



WHAT ARE THE FACTS?



TRY OUT
EVENT STORMING

UNDERSTAND HOW FACTS ARE CAUSALLY RELATED HOW FACTS FLOW IN THE SYSTEM

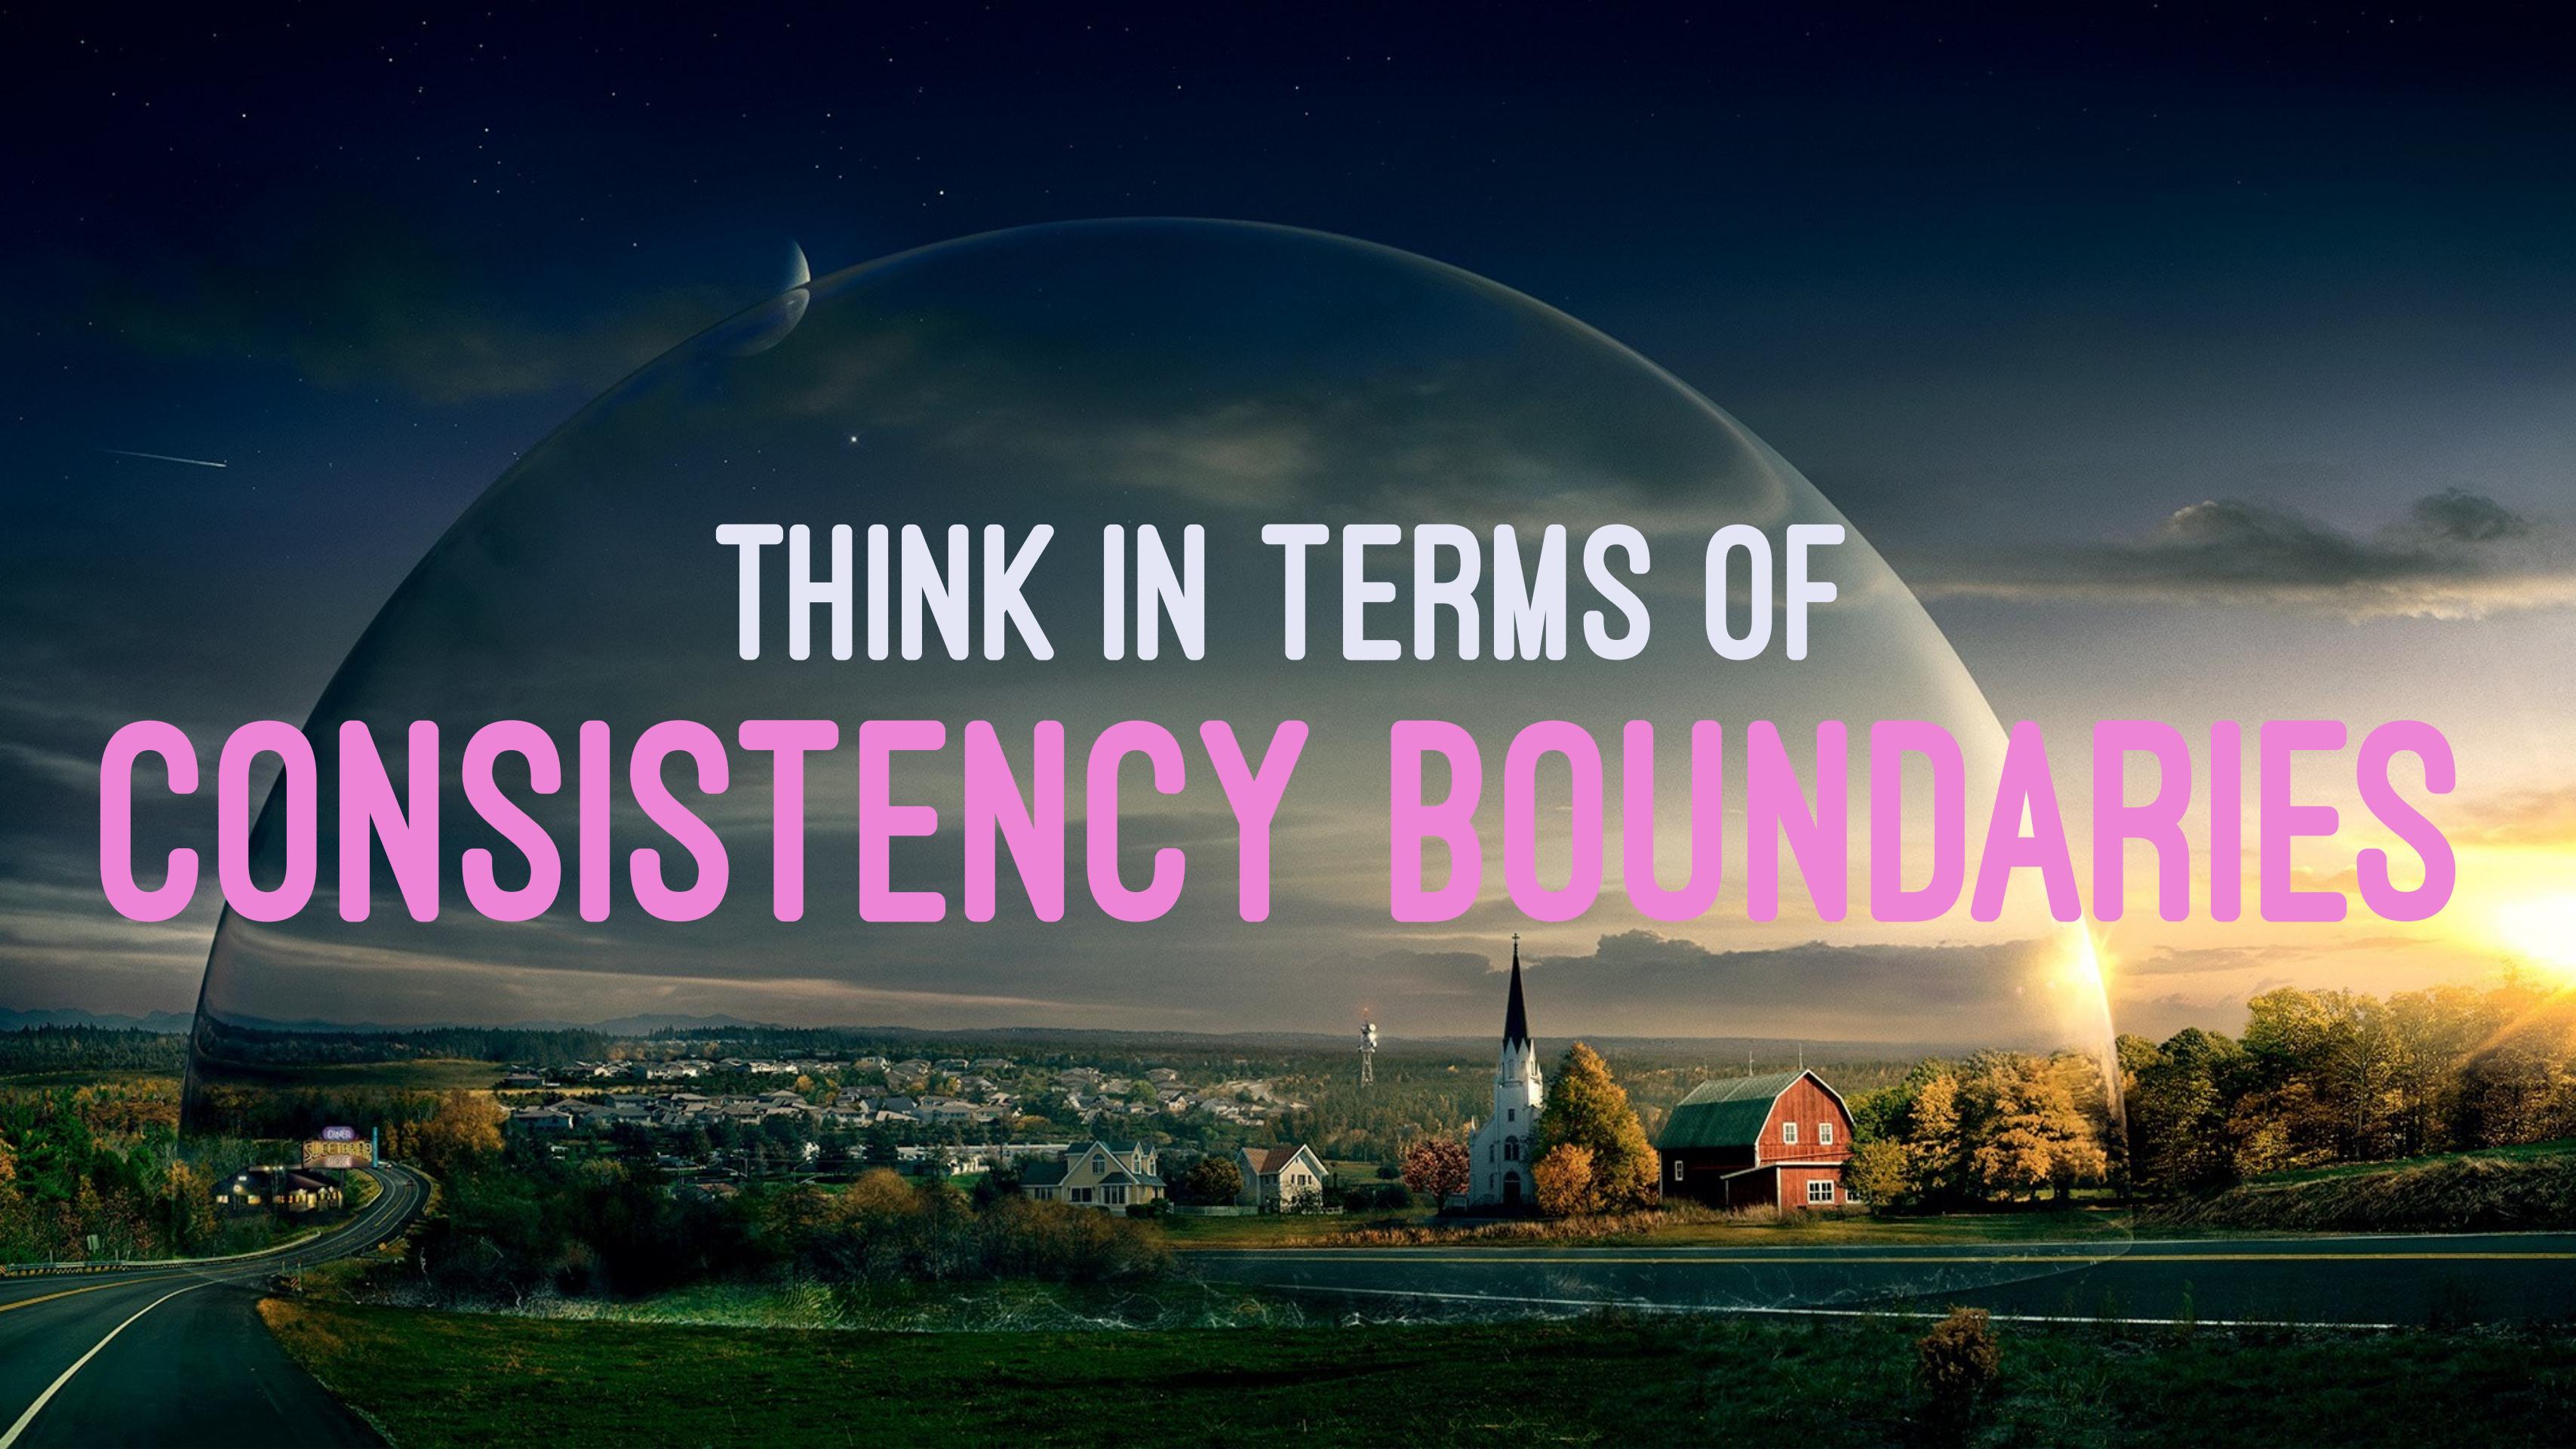
Peter Alvaro (@palvaro) posted a tweet on April 1, 2016, at 1:09 AM. The tweet reads: "I shall never tire of writing "causality is reachability in spacetime" on the blackboard". The post has 6 retweets and 38 likes. A blue "Following" button is visible next to the user's name.

RETWEETS 6 LIKES 38

1:09 AM - 1 Apr 2016

Following

...

A landscape photograph of a small town at sunset. In the foreground, there's a road curving through green fields. To the left, a gas station sign reads "DINER SWEET STOP". In the center, a white church with a tall steeple stands next to a red barn. The town extends into the background with more houses and trees. The sky is filled with warm, golden light from the setting sun, which is partially visible on the right. A large, semi-transparent dark sphere, resembling Earth or a planet, is positioned in the upper half of the image, casting a shadow over the town. The sphere has a bright ring around its equator and some internal cloud-like structures.

THINK IN TERMS OF
CONSISTENCY BOUNDARIES



WE NEED TO
**CONTAIN MUTABLE STATE &
PUBLISH FACTS**

AGGREGATE

→ UNIT OF CONSISTENCY
→ UNIT OF FAILURE

PRACTICE
REACTIVE DESIGN

REACTIVE PROGRAMMING VS REACTIVE SYSTEMS

**REACTIVE PROGRAMMING
CAN HELP US MAKE THE
INDIVIDUAL INSTANCE
HIGHLY PERFORMANT & EFFICIENT**

ASYNCHRONOUS



 **Viktor Klang**
@viktorklang

"To get consistent & fast response in a real time system, you've to work asynchronously [...] never wait on something happening." [@hintjens](#)

RETWEETS LIKES
6 **6**

12:44 PM - 21 Sep 2016

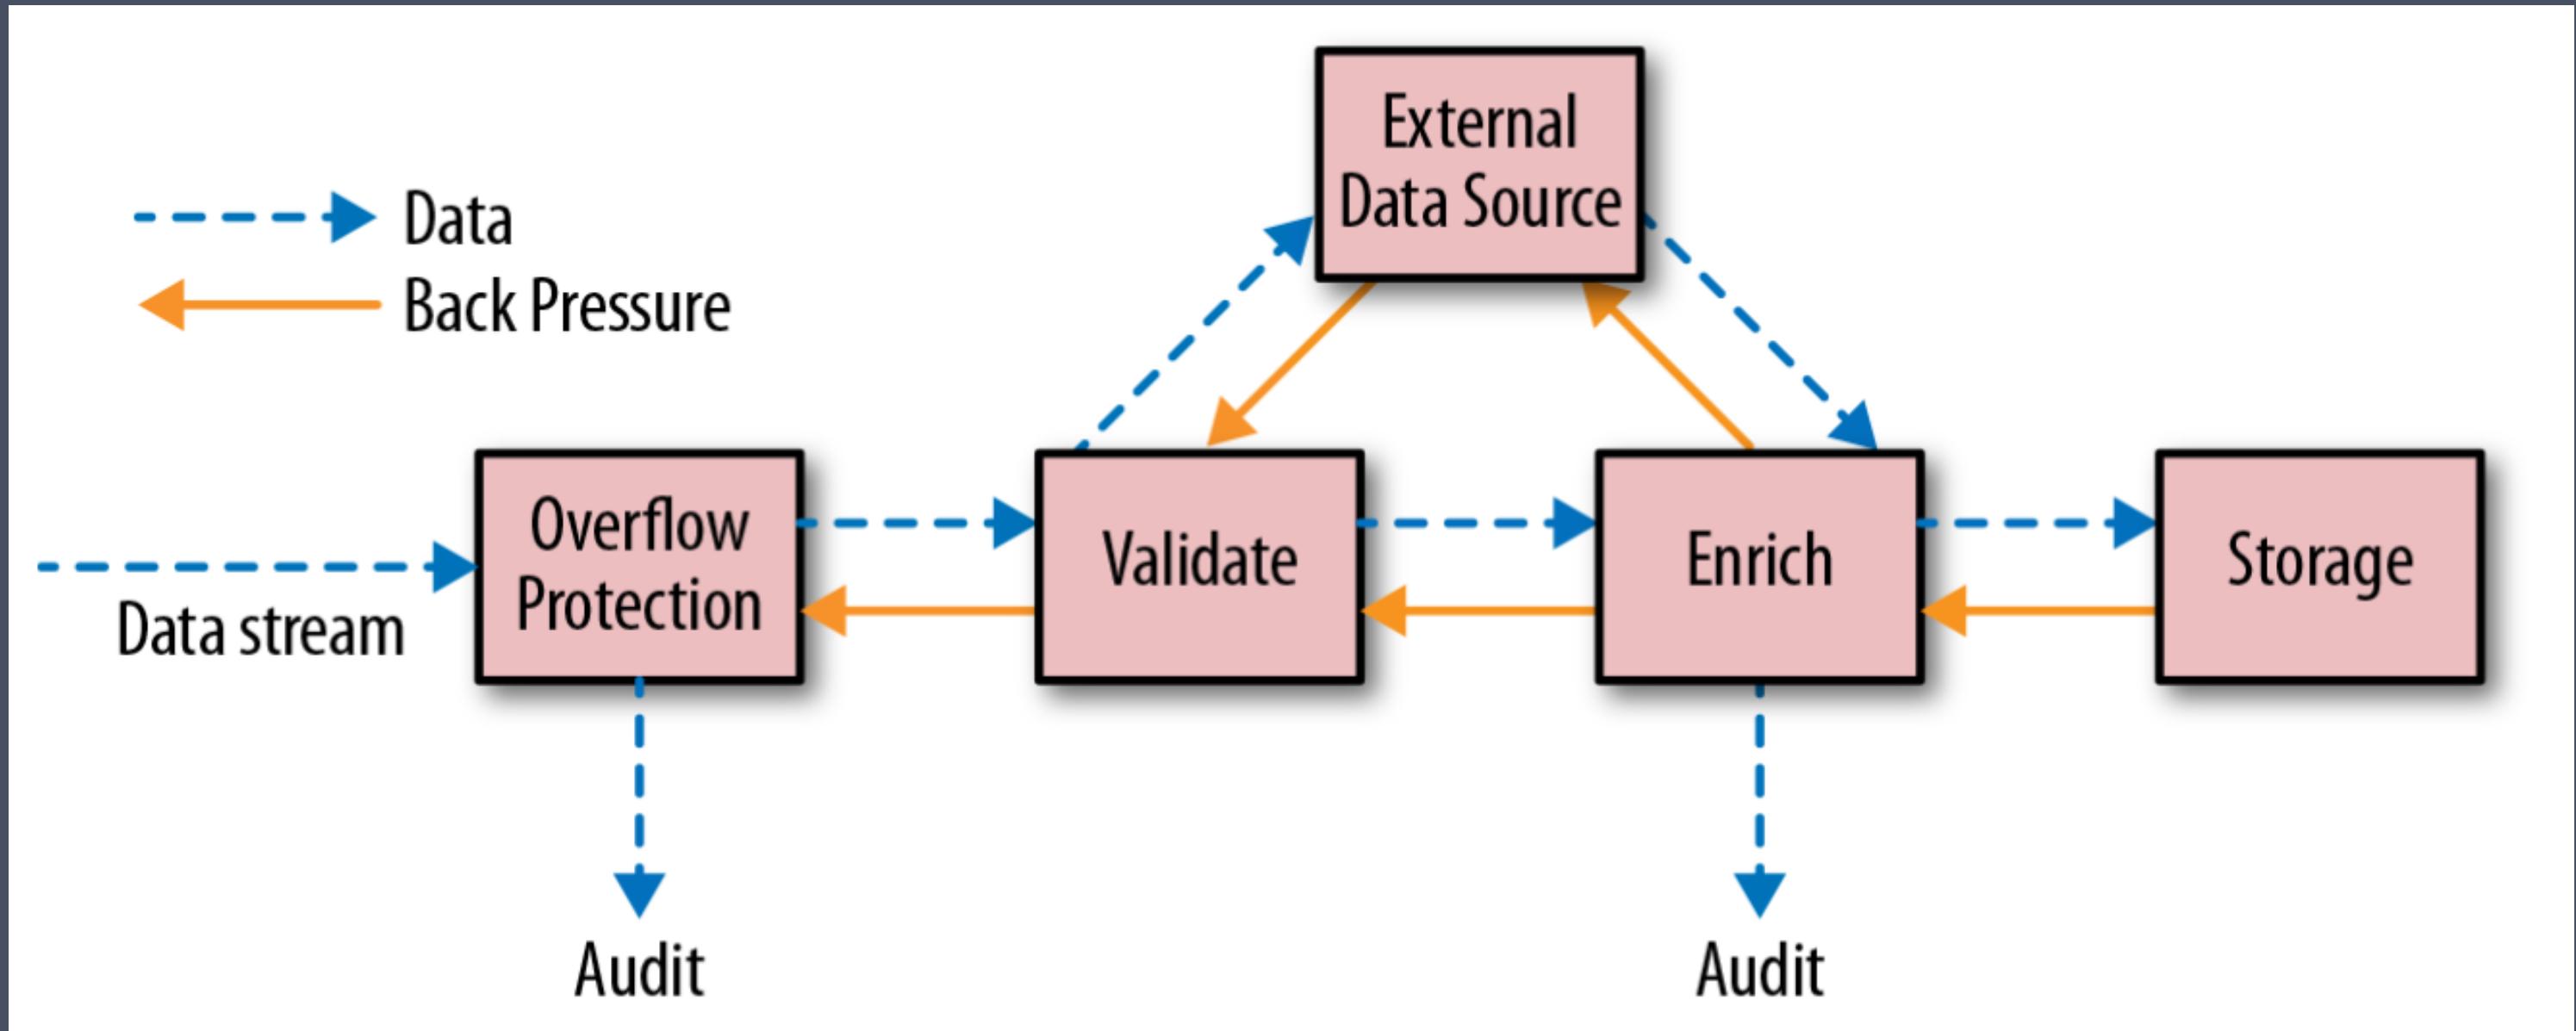
ASYNCHRONOUS & NON-BLOCKING

- MORE EFFICIENT USE OF RESOURCES
- MINIMIZES CONTENTION ON SHARED RESOURCES

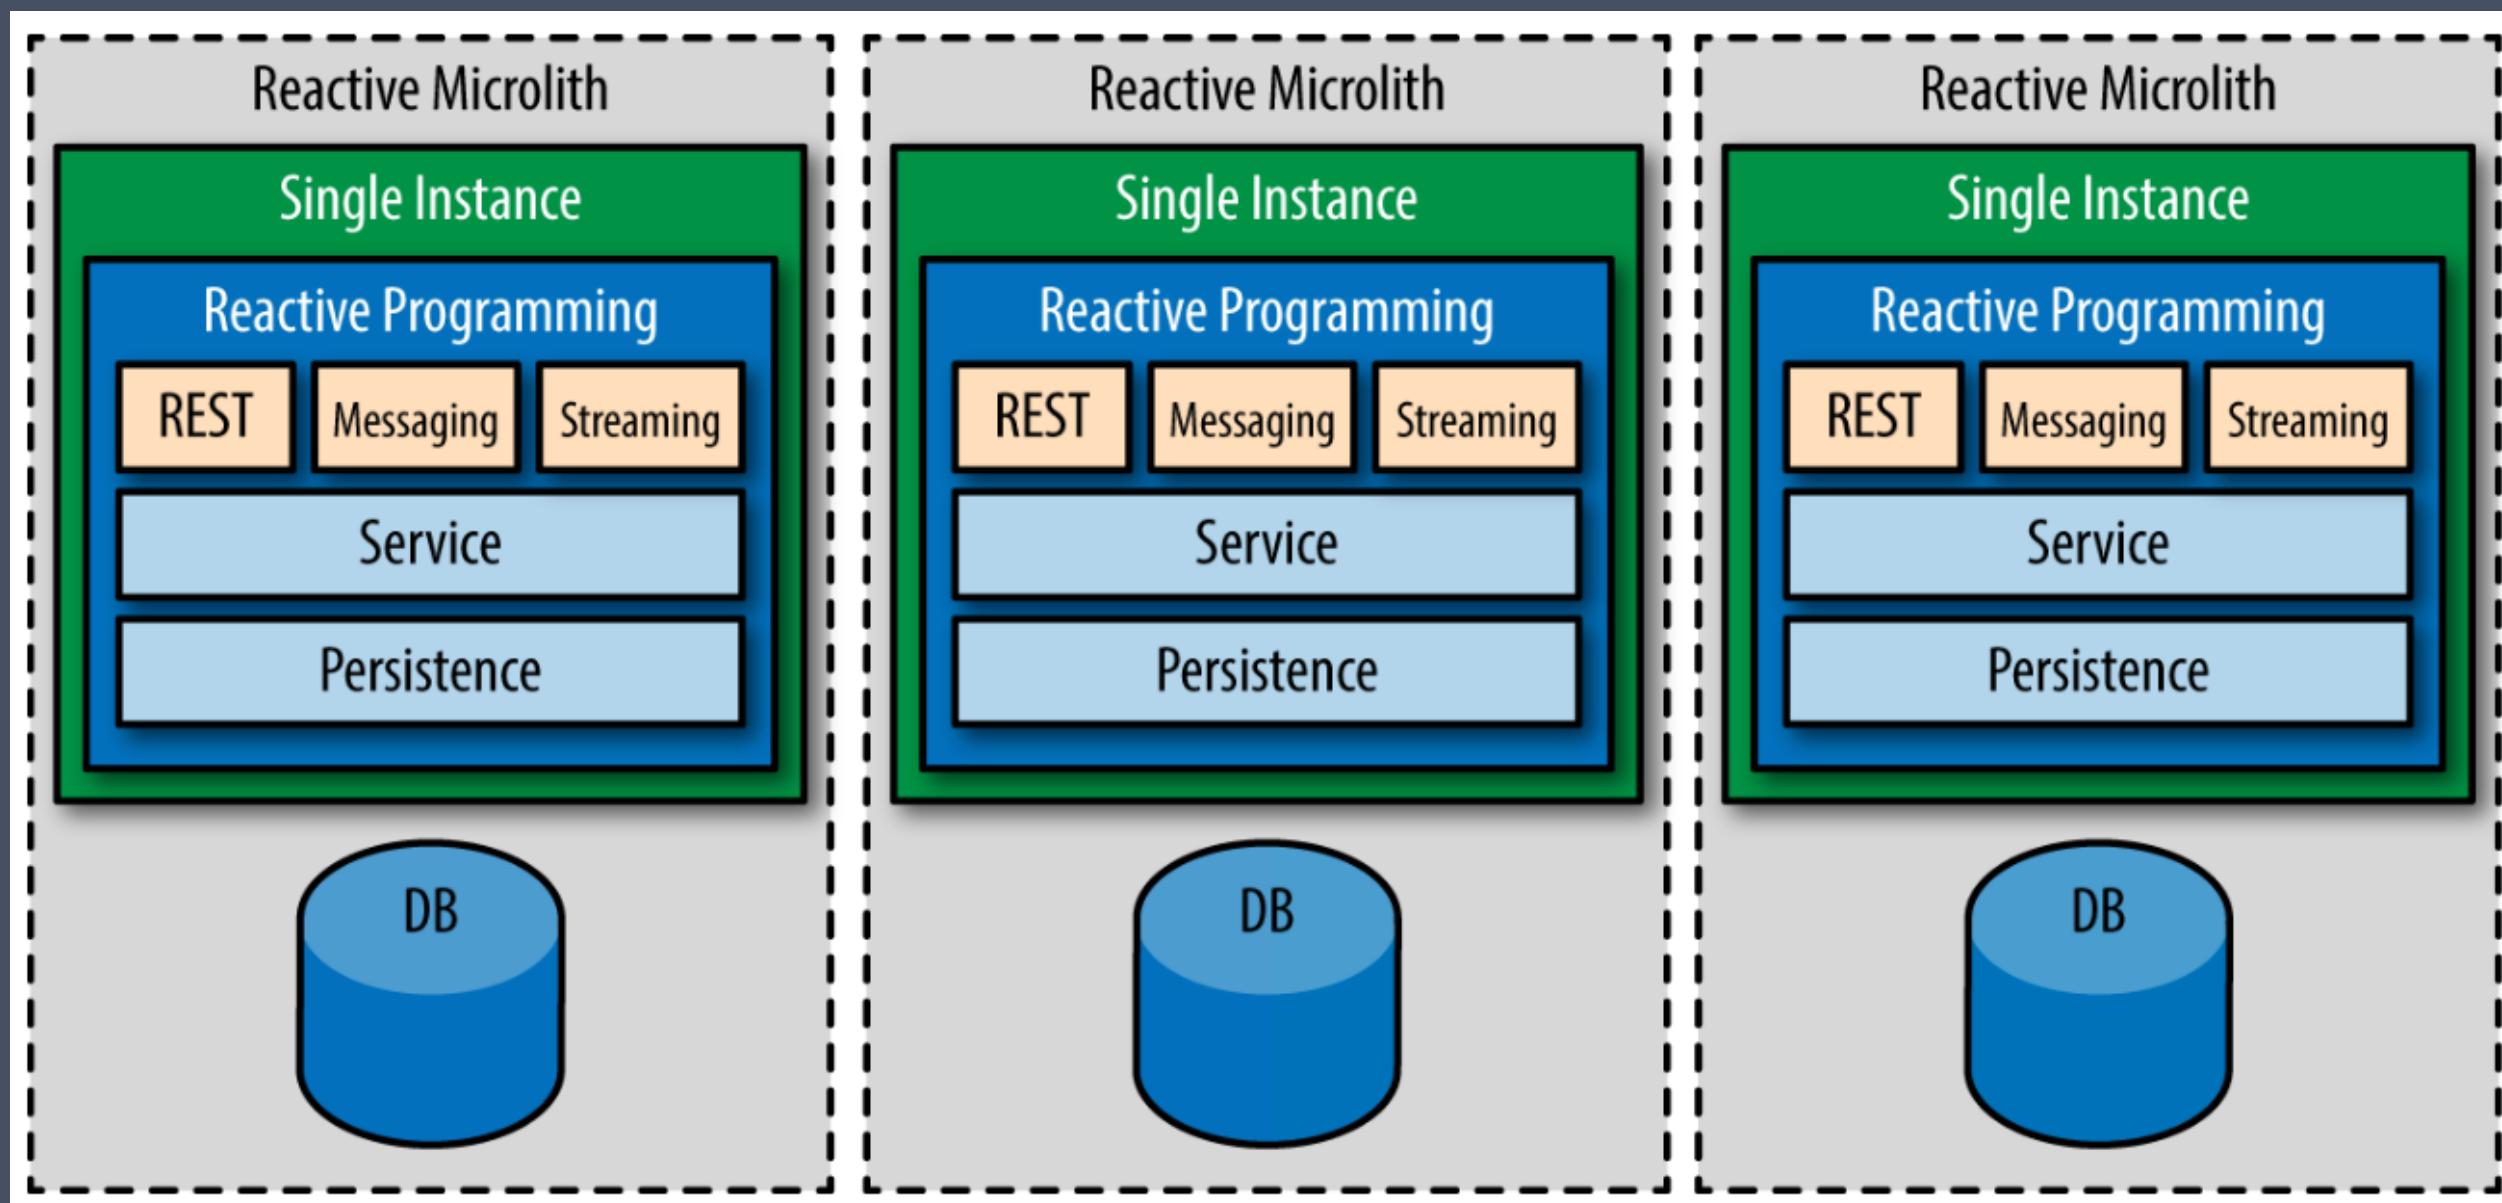
**ALWAYS APPLY BACKPRESSURE
A FAST SYSTEM
SHOULD NOT OVERLOAD
A SLOW SYSTEM**



BACKPRESSURE

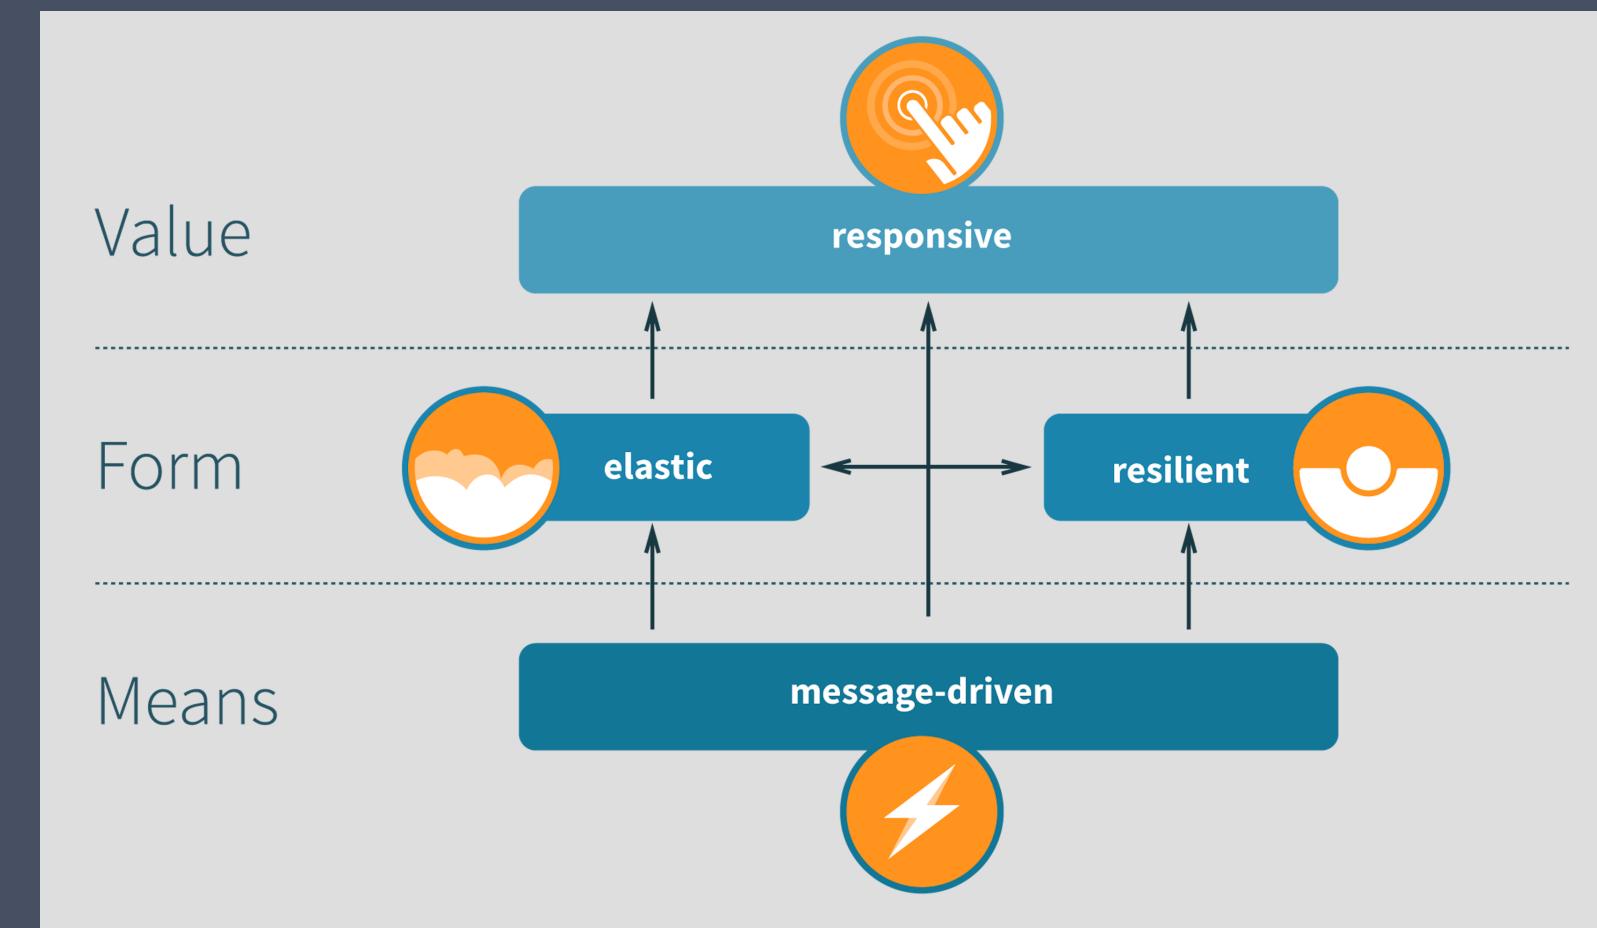


LET'S APPLY REACTIVE PROGRAMMING TO OUR MICROLITHS



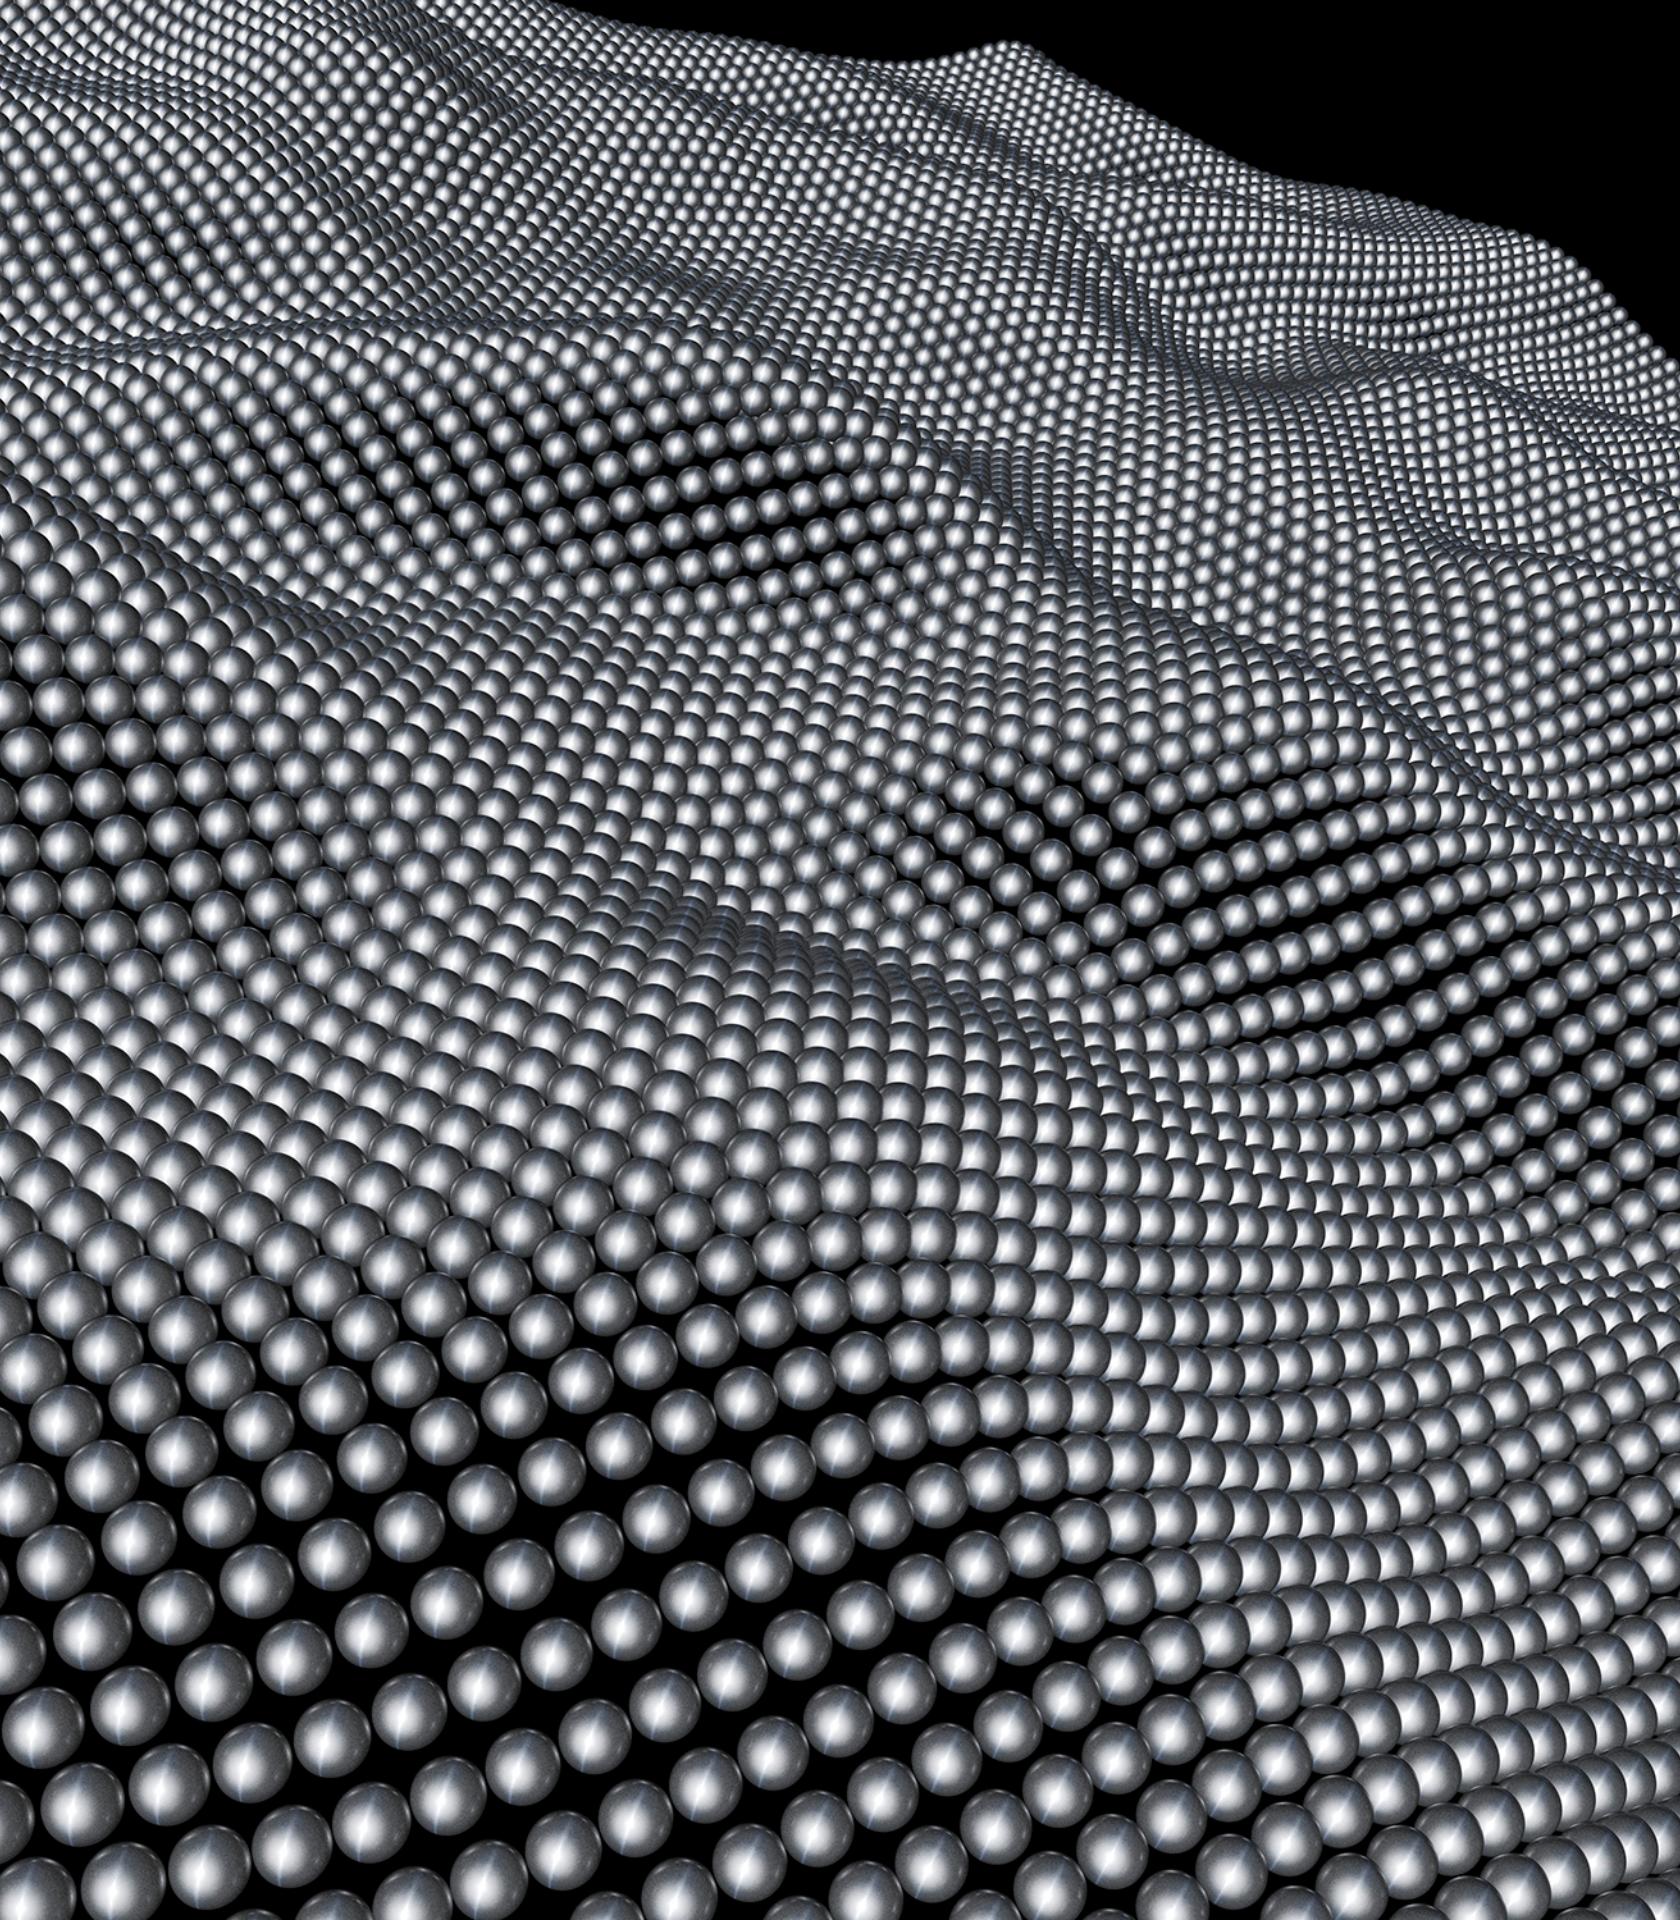
WE'RE GETTING THERE, BUT WE STILL HAVE A
SINGLE INSTANCE MICROSERVICE
→ NOT SCALABLE
→ NOT RESILIENT

REACTIVE SYSTEMS CAN HELP US BUILD DISTRIBUTED SYSTEMS THAT ARE ELASTIC & RESILIENT



REACTIVE SYSTEMS ARE BASED ON
ASYNCHRONOUS
MESSAGE-PASSING

ALLOWS DECOUPLING IN
SPACE
AND
TIME



ALLOWS FOR LOCATION TRANSPARENCY

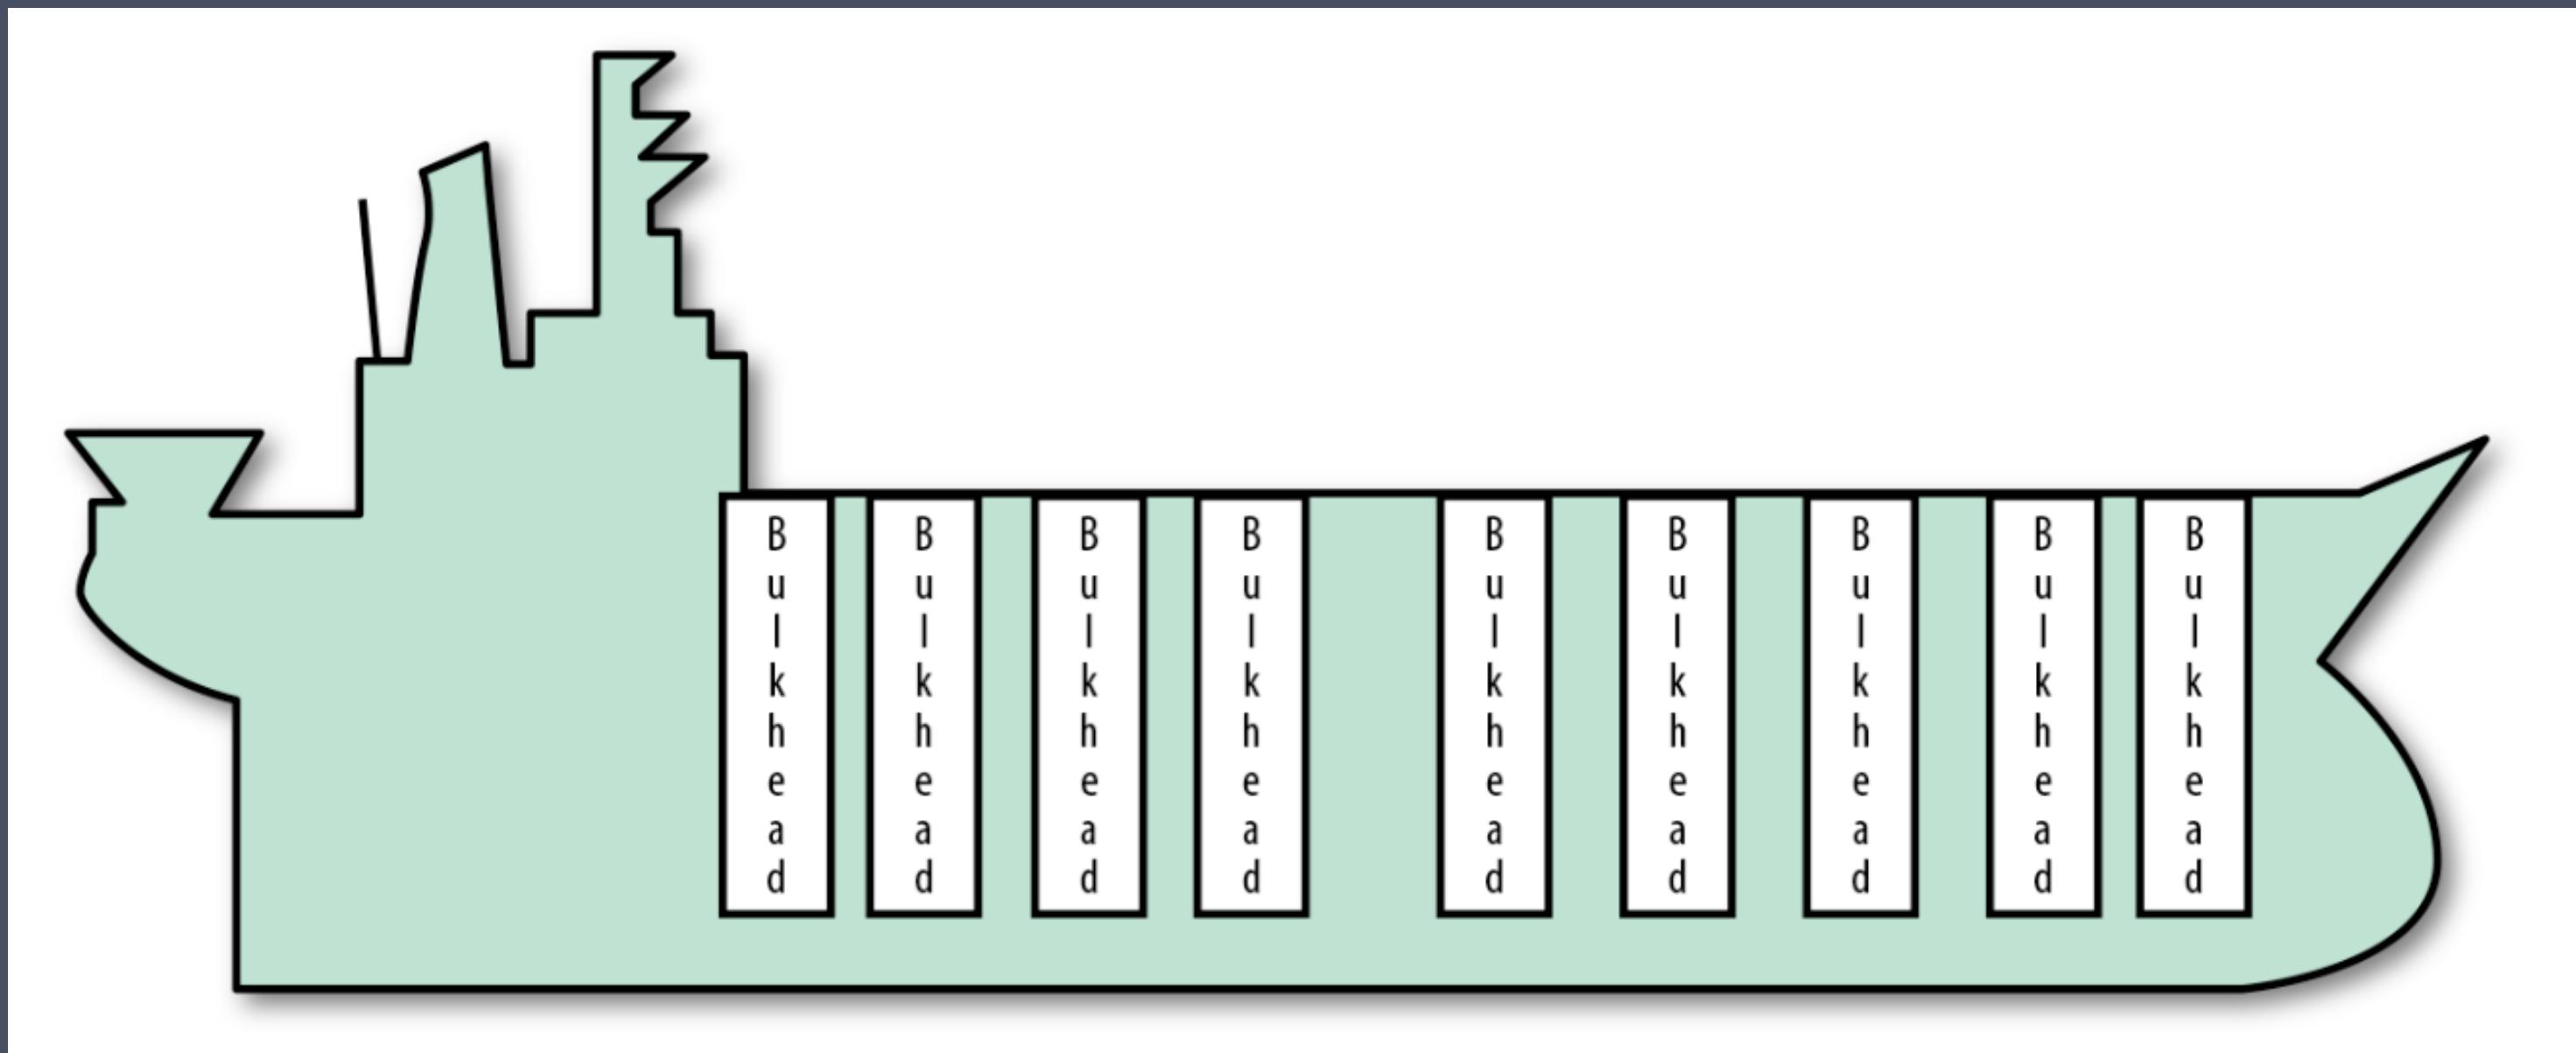
ONE COMMUNICATION ABSTRACTION ACROSS ALL DIMENSIONS OF SCALE

CORE \Rightarrow SOCKET \Rightarrow CPU \Rightarrow

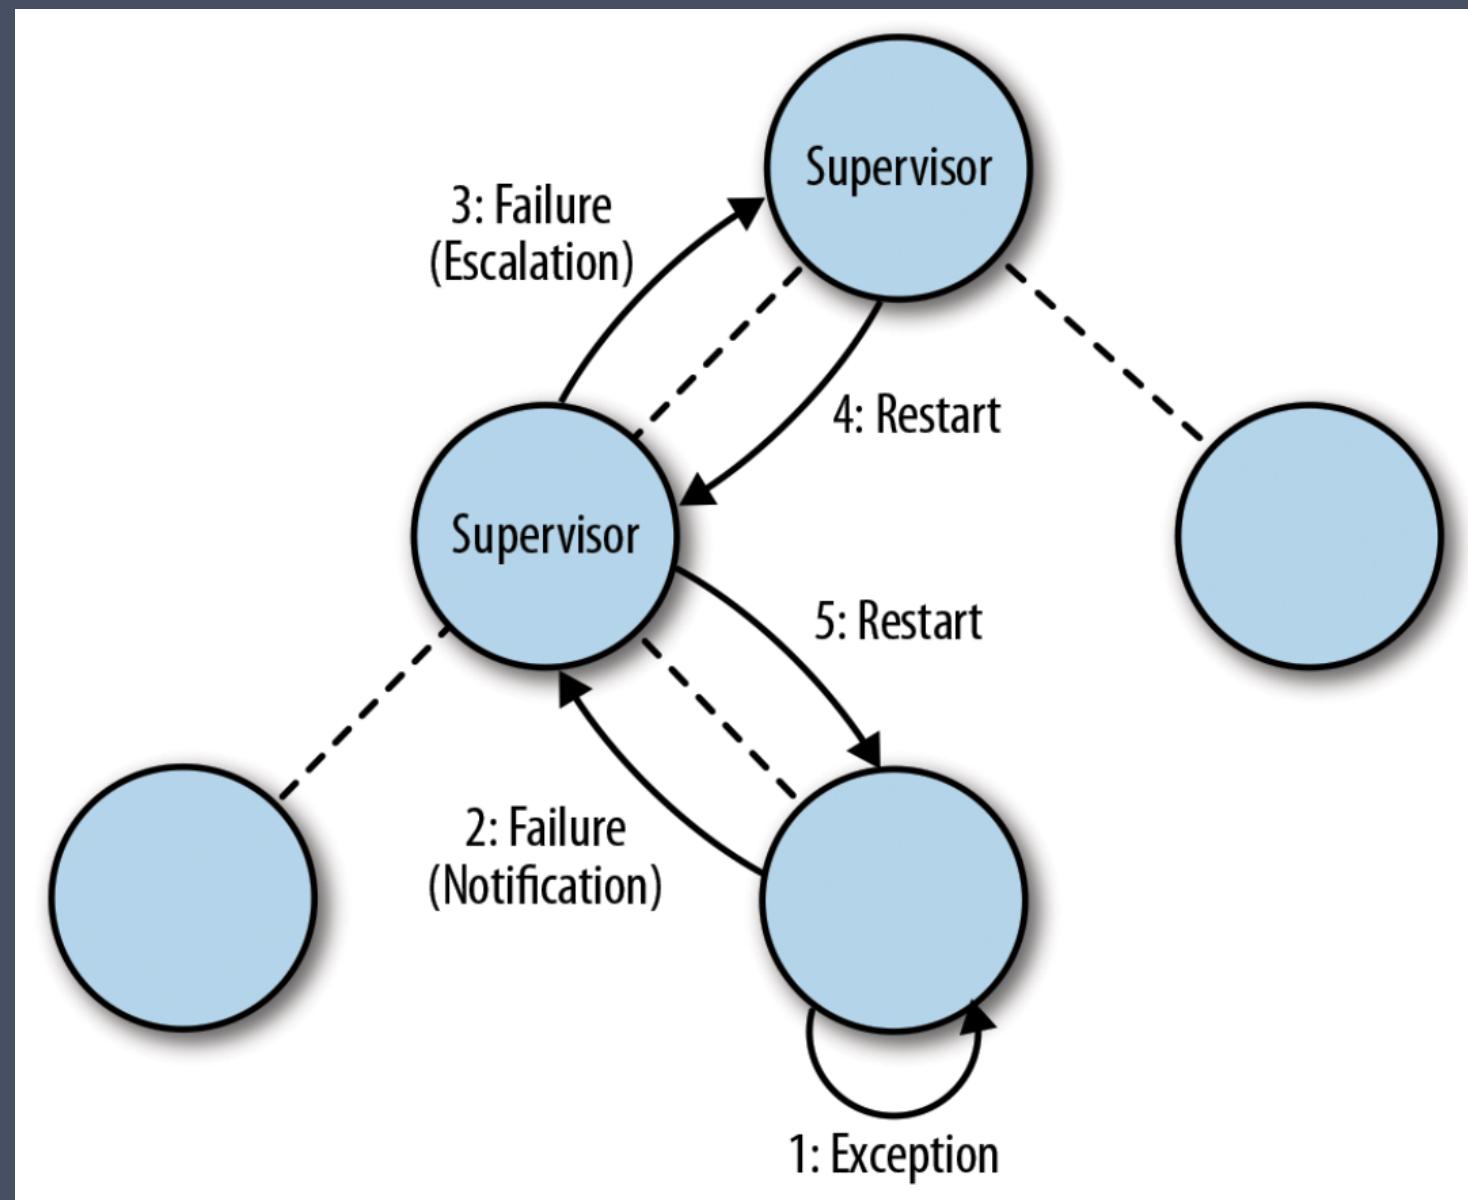
CONTAINER \Rightarrow SERVER \Rightarrow RACK \Rightarrow

DATA CENTER \Rightarrow SYSTEM

SELF-HEALING THROUGH BULKHEADING



SELF-HEALING THROUGH SUPERVISION

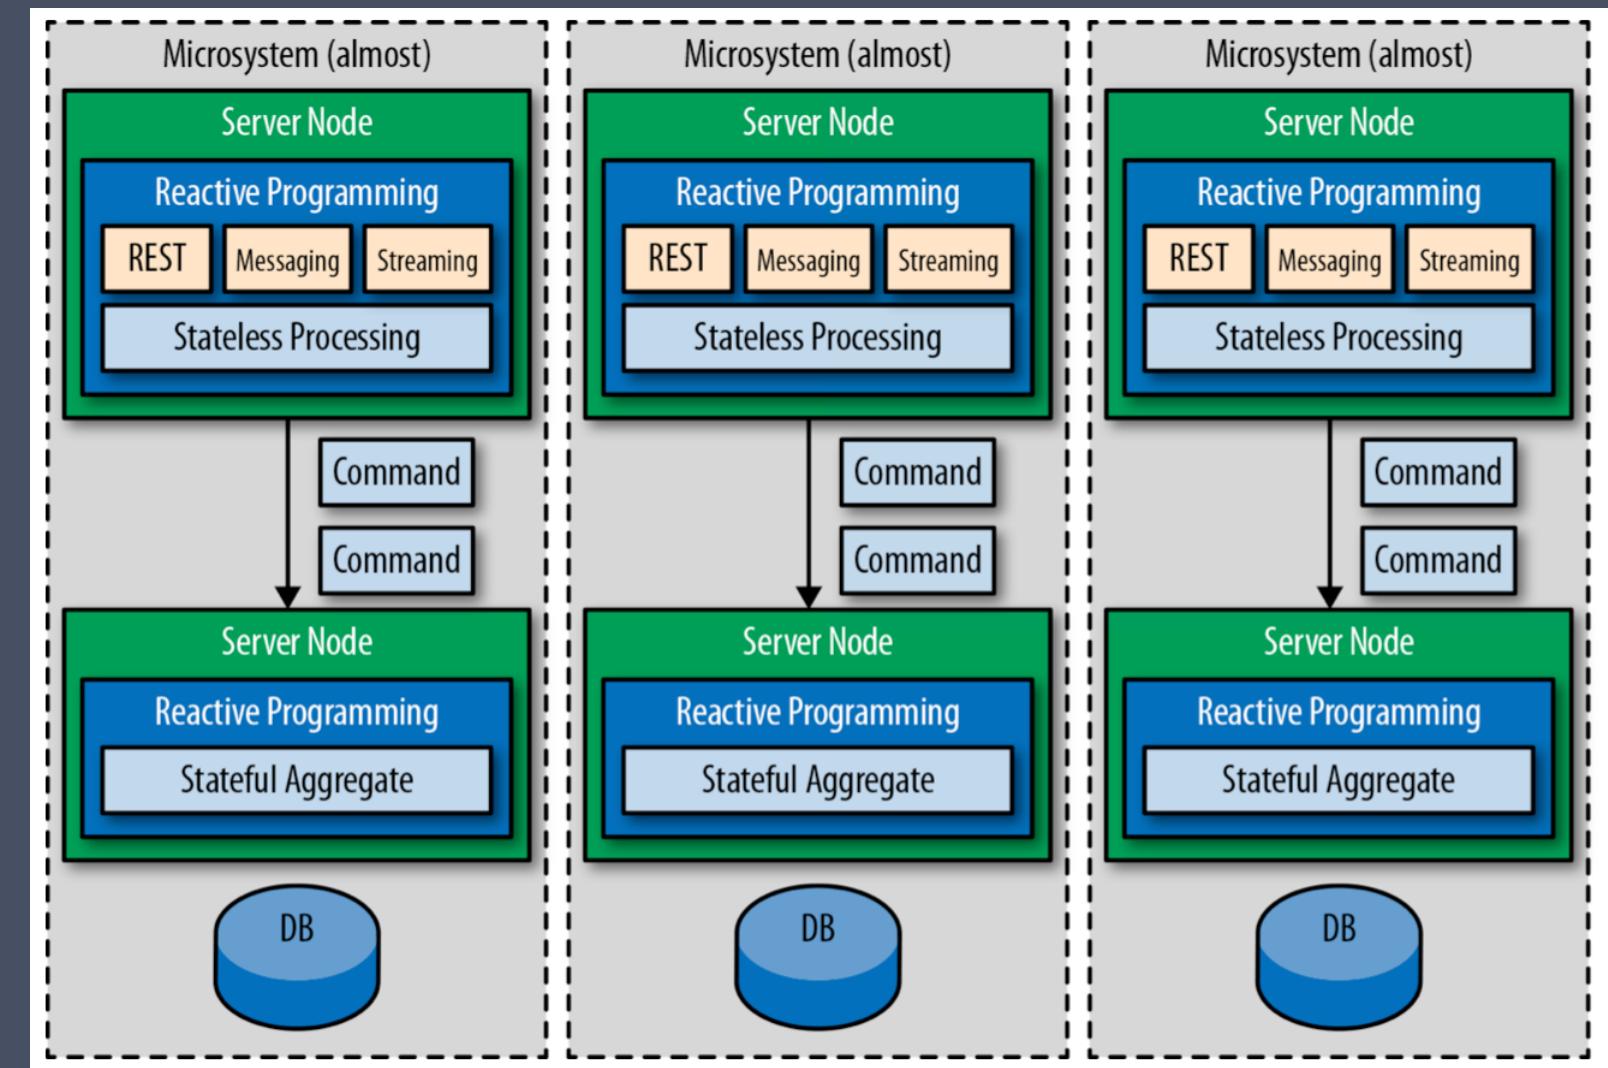


MICROSERVICES
COME AS SYSTEMS

EACH MICROSERVICE
NEEDS BE DESIGNED AS
A DISTRIBUTED SYSTEM
A MICROSYSTEM

WE NEED TO MOVE
FROM MICROLITHS
TO MICROSYSTEMS

SEPARATE THE
STATELESS BEHAVIOR
FROM THE
STATEFUL ENTITY
TO SCALE THEM INDIVIDUALLY



SCALING (STATELESS) BEHAVIOR

IS FAIR

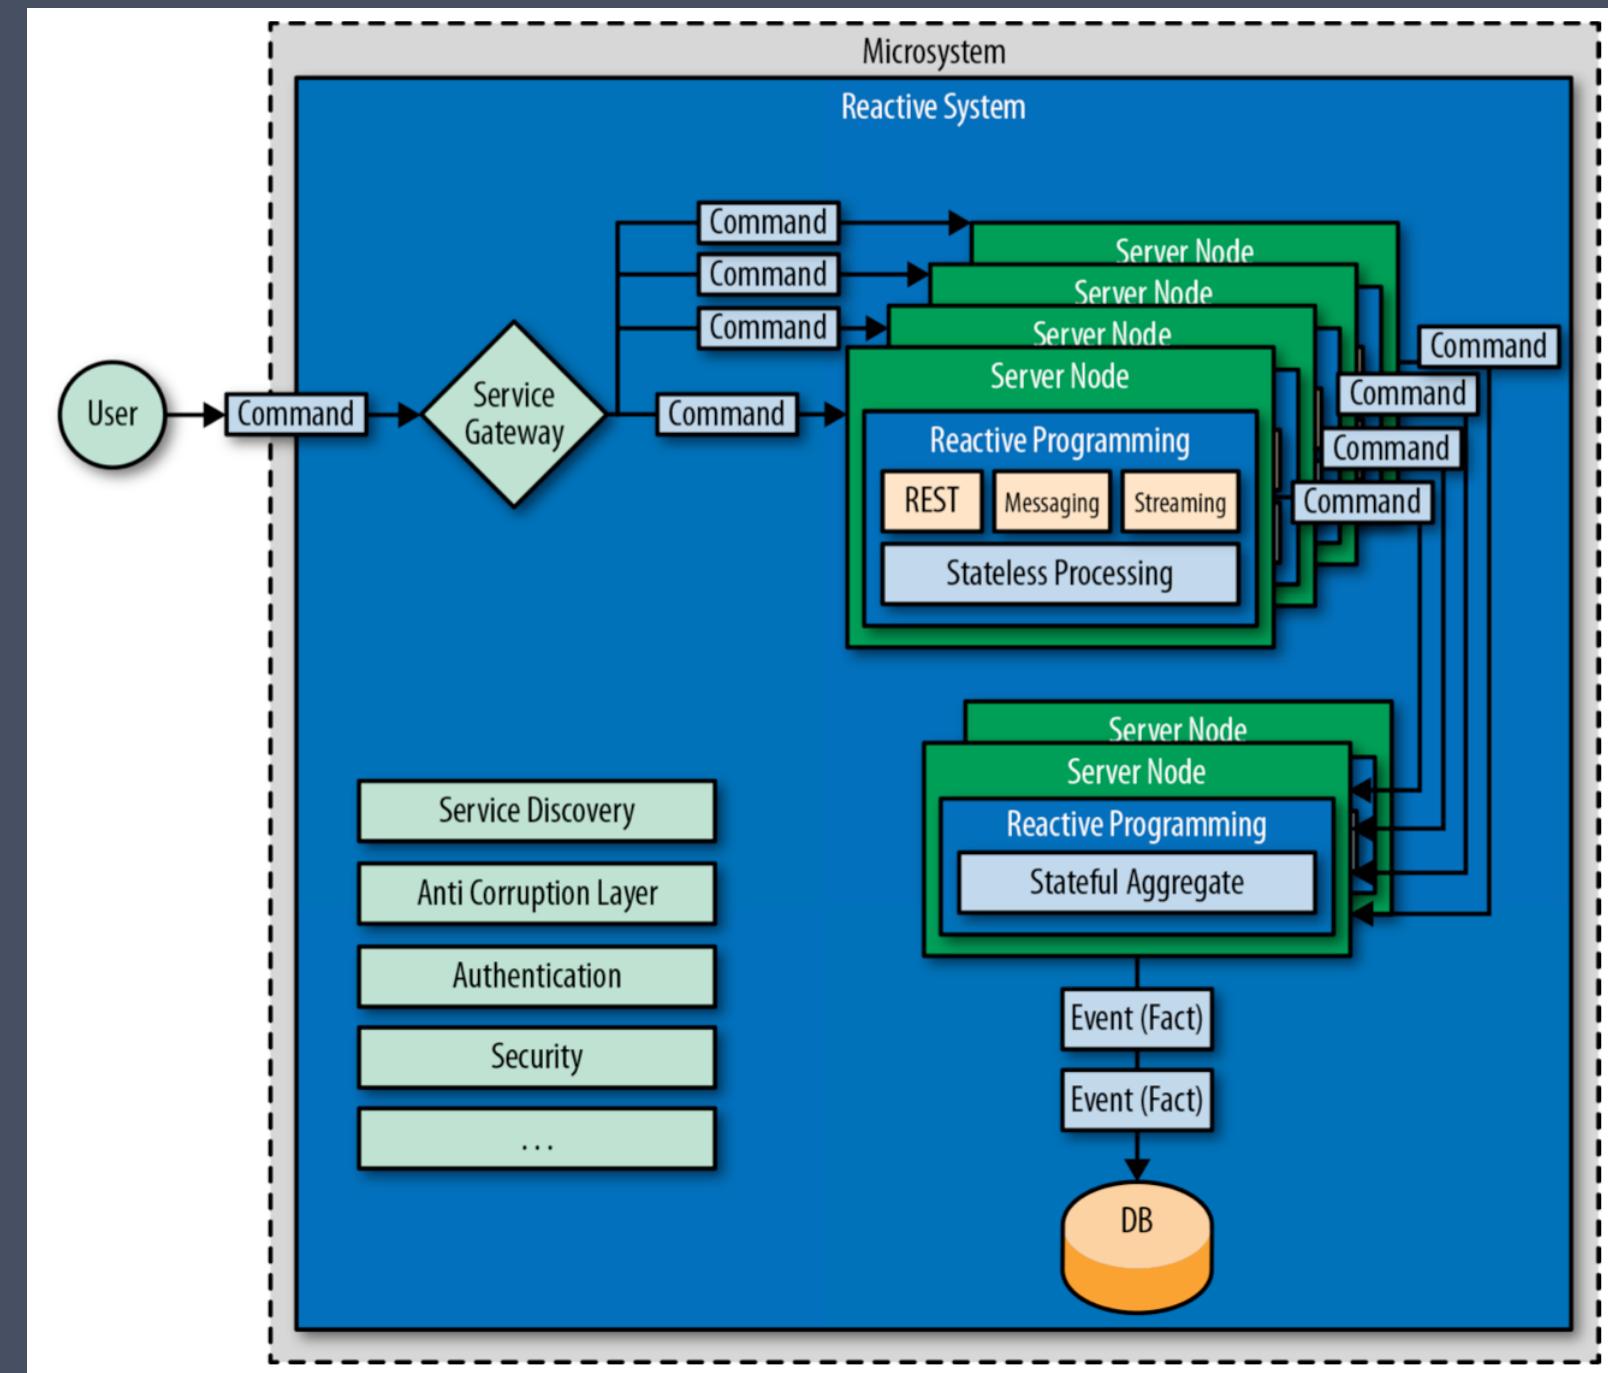
SCALING (STATEFUL) ENTITIES

IS HARD



THERE IS NO SUCH THING AS A
"STATELESS" ARCHITECTURE
IT'S JUST SOMEONE ELSE'S PROBLEM

REACTIVE MICROSYSTEM



PRACTICE

EVENT-BASED
PERSISTENCE

UPDATE-IN-PLACE STRIKES SYSTEMS
DESIGNERS AS A CARDINAL SIN:
IT VIOLATES TRADITIONAL ACCOUNTING
PRACTICES THAT HAVE BEEN OBSERVED
FOR HUNDREDS OF YEARS.

- JIM GRAY

**THE TRUTH IS THE LOG.
THE DATABASE IS A CACHE OF A
SUBSET OF THE LOG.**

- PAT HELLAND

EVENT LOGGING FOR SCALABLE PERSISTENCE

CRUD IS DEAD



EVENT SOURCING

A CURE FOR THE CARDINAL SIN

MODELING EVENTS FORCES YOU TO
HAVE A TEMPORAL FOCUS ON WHAT'S
GOING ON IN THE SYSTEM.
TIME BECOMES A CRUCIAL FACTOR OF
THE SYSTEM.

- GREG YOUNG

**THE
HUE
LOG**

**IS A DATABASE OF THE PAST
NOT JUST A DATABASE OF THE PRESENT**

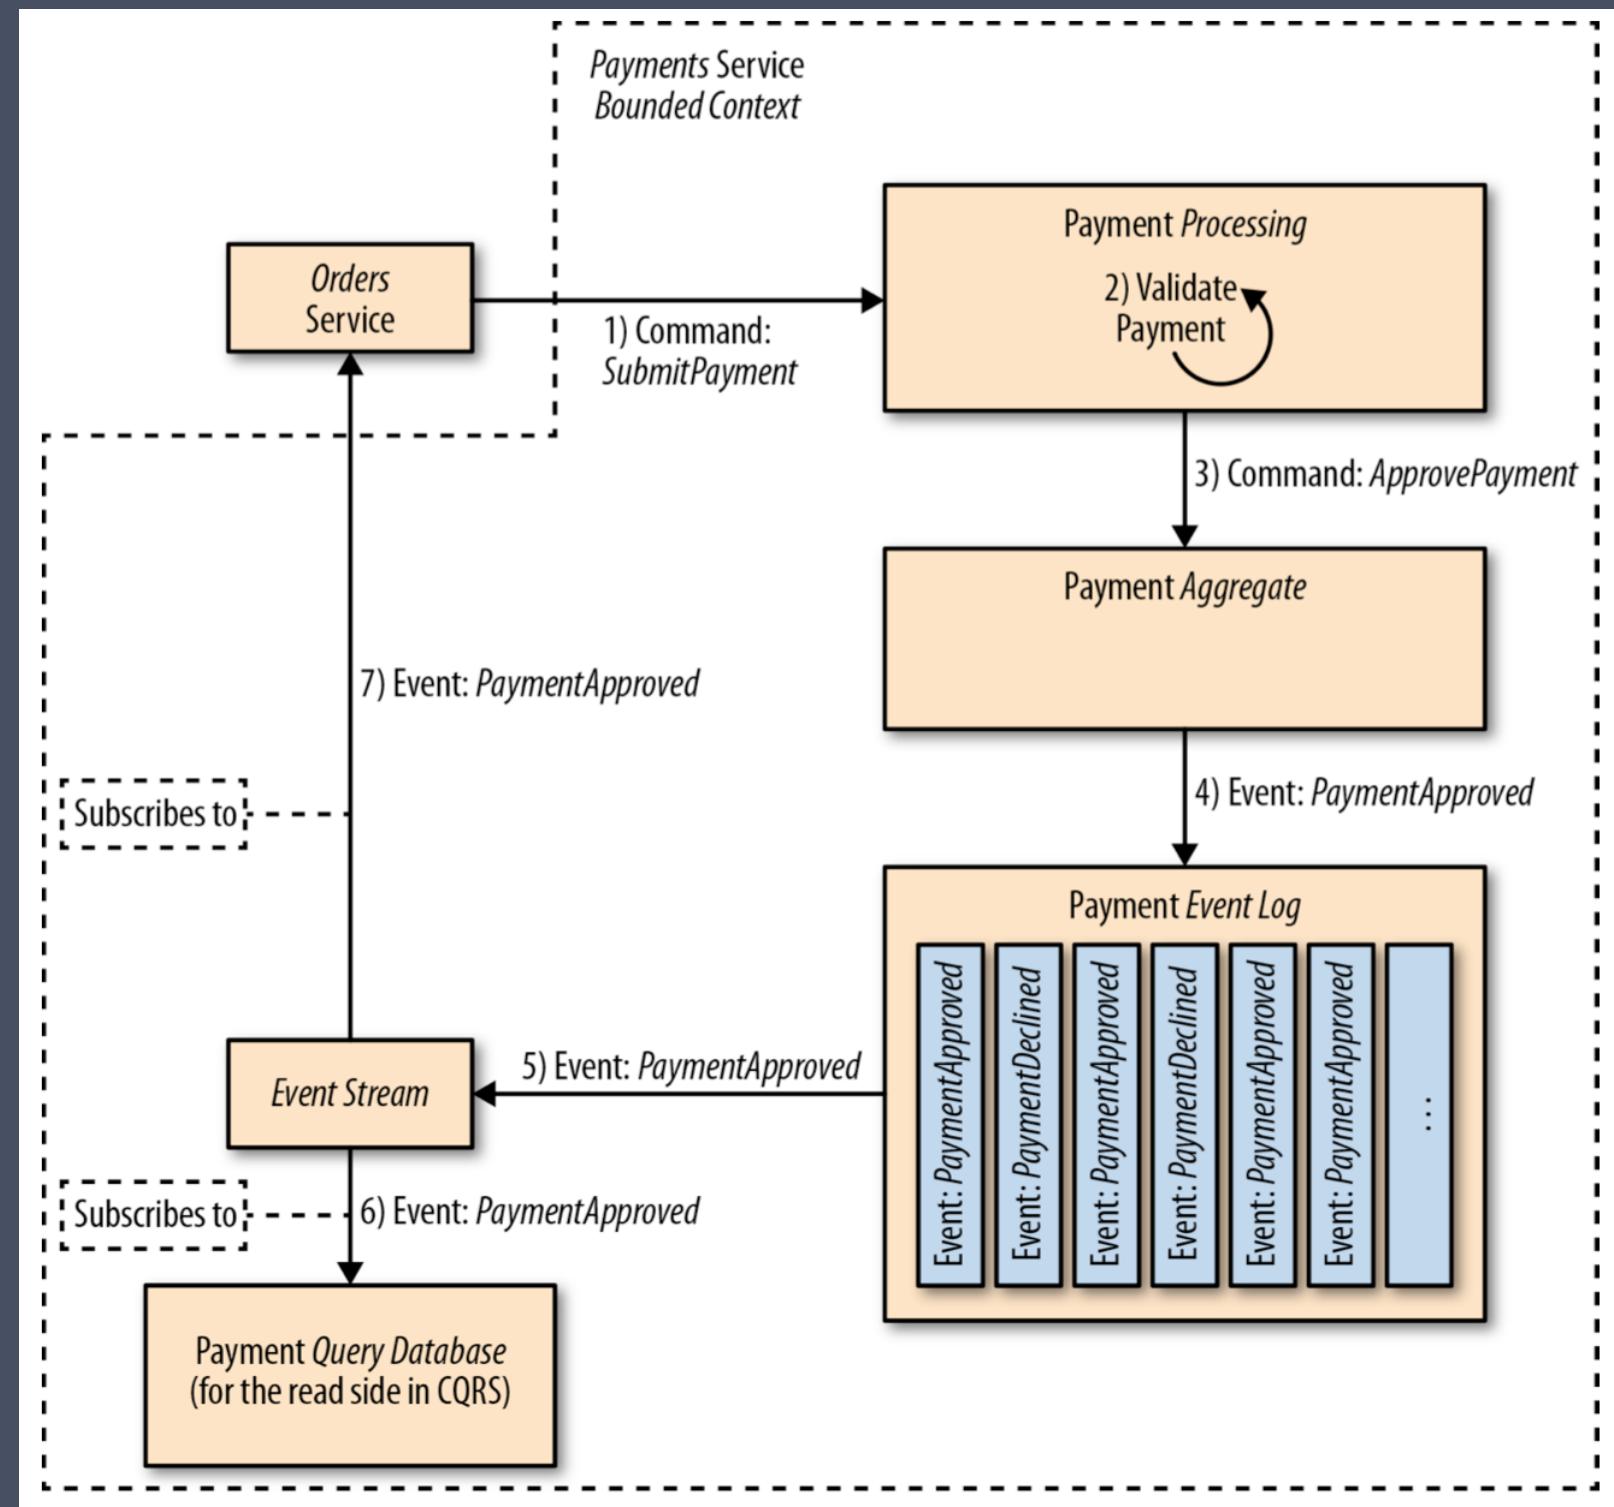
EVENT LOGGING AVOIDS THE INFAMOUS
**OBJECT-RELATIONAL
IMPEDENCE MISMATCH**

UNTANGLE THE
READ & WRITE
MODELS WITH
CQRS

ADVANTAGES OF USING CQRS:

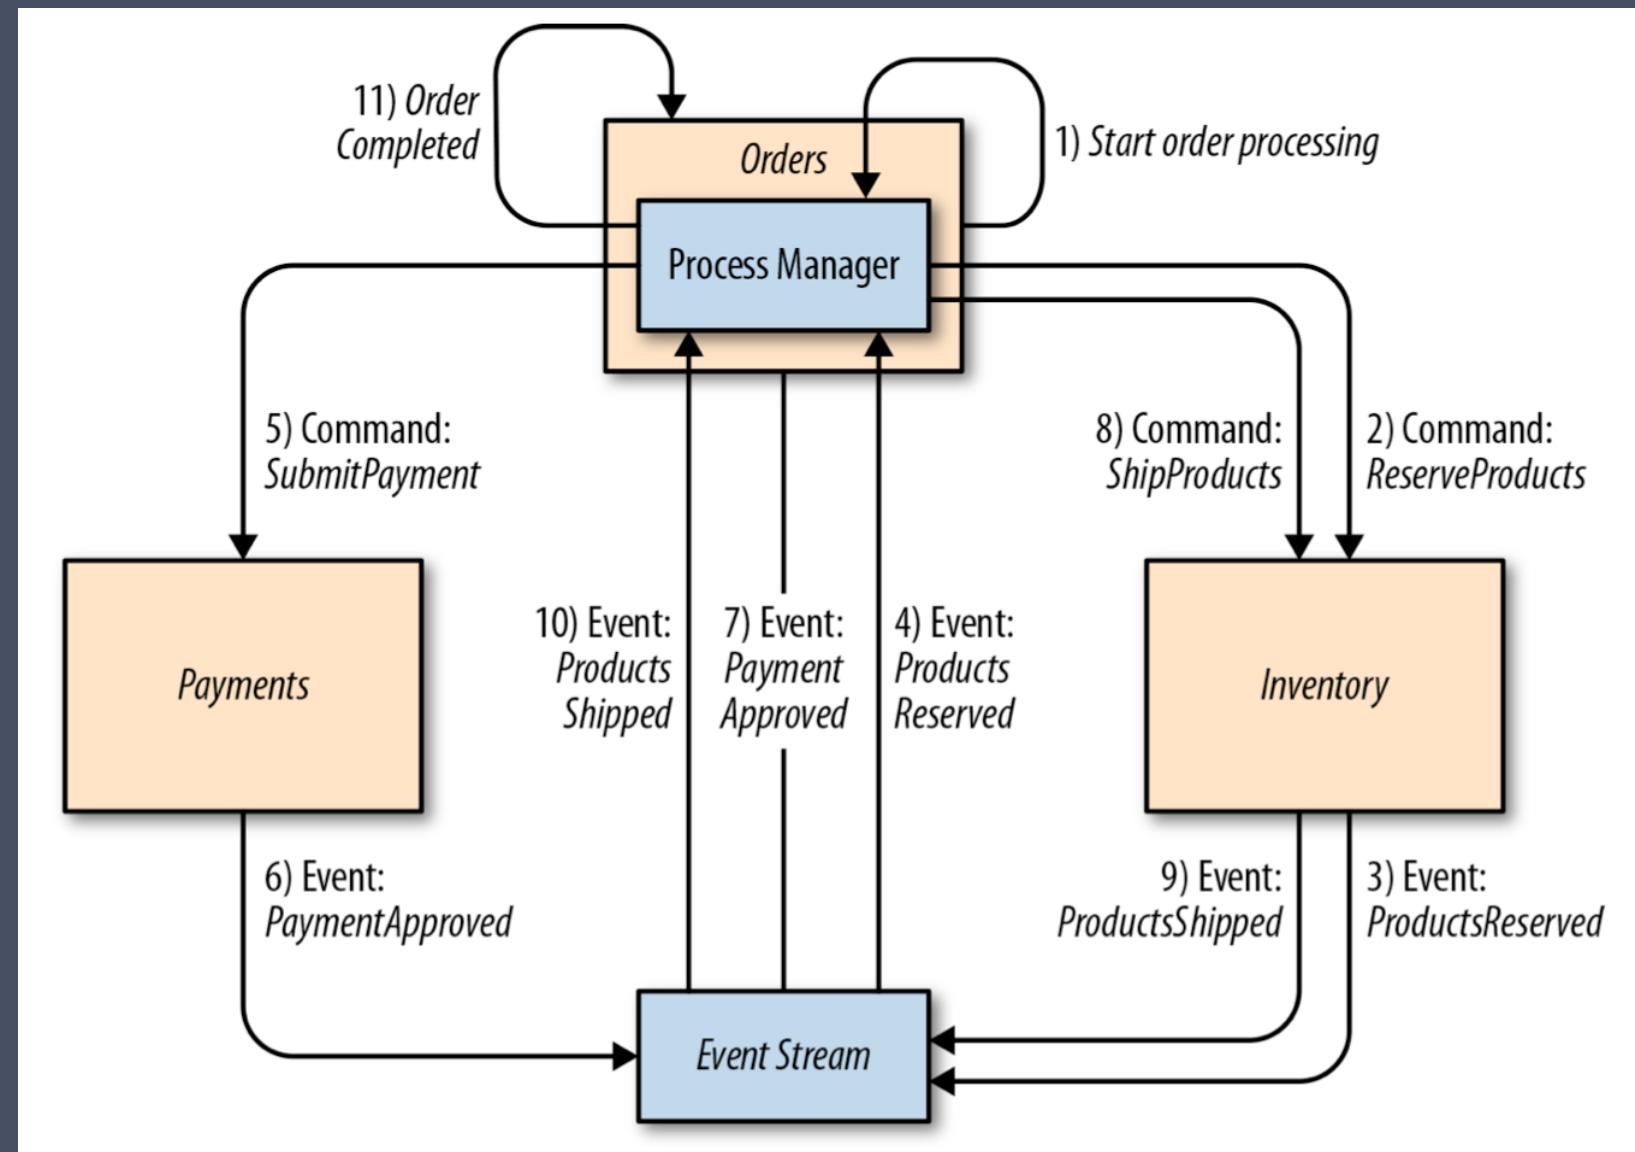
1. RESILIENCE
2. SCALABILITY
3. POLYGLOT
PERSISTENCE

USE EVENT SOURCING AND EVENT STREAMING



HOW CAN WE
COORDINATE WORK
ACROSS AGGREGATES?

EVENT-DRIVEN WORKFLOW



BUT WHAT ABOUT
TRANSACTIONS?

TWO-PHASE COMMIT IS THE
ANTI-AVAILABILITY PROTOCOL.

- PAT HELLAND

IN GENERAL. APPLICATION DEVELOPERS SIMPLY DO NOT IMPLEMENT LARGE SCALABLE APPLICATIONS ASSUMING DISTRIBUTED TRANSACTIONS.

– PAT HELLAND

Life beyond Distributed Transactions: an Apostate's Opinion

Position Paper

Pat Helland

Amazon.Com
705 Fifth Ave South
Seattle, WA 98104
USA

PHelland@Amazon.com

The positions expressed in this paper are personal opinions and do not in any way reflect the positions of my employer Amazon.com.

ABSTRACT

Many decades of work have been invested in the area of distributed transactions including protocols such as 2PC, Paxos, and various approaches to quorum. These protocols provide the application programmer a façade of global serializability. Personally, I have invested a non-trivial portion of my career as a strong advocate for the implementation and use of platforms providing guarantees of global serializability.

My experience over the last decade has led me to liken these platforms to the Maginot Line¹. In general, application developers simply do not implement large scalable applications assuming distributed transactions. When they attempt to use distributed transactions, the projects founder because the performance costs and fragility make them impractical. Natural selection kicks in...

¹ The Maginot Line was a huge fortress that ran the length of the Franco-German border and was constructed at great expense between World War I and World War II. It successfully kept the German army from directly crossing the border between France and Germany. It was quickly bypassed by the Germans in 1940 who invaded through Belgium.

This article is published under a Creative Commons License Agreement (<http://creativecommons.org/licenses/by/2.5/>). You may copy, distribute, display, and perform the work, make derivative works and make commercial use of the work, but you must attribute the work to the author and CIDR 2007. 3rd Biennial Conference on Innovative DataSystems Research (CIDR) January 7-10, Asilomar, California USA.

Instead, applications are built using different techniques which do not provide the same transactional guarantees but still meet the needs of their businesses.

This paper explores and names some of the practical approaches used in the implementations of large-scale mission-critical applications in a world which rejects distributed transactions. We discuss the management of fine-grained pieces of application data which may be repartitioned over time as the application grows. We also discuss the design patterns used in sending messages between these repartitionable pieces of data.

The reason for starting this discussion is to raise awareness of new patterns for two reasons. First, it is my belief that this awareness can ease the challenges of people hand-crafting very large scalable applications. Second, by observing the patterns, hopefully the industry can work towards the creation of platforms that make it easier to build these very large applications.

1. INTRODUCTION

Let's examine some goals for this paper, some assumptions that I am making for this discussion, and then some opinions derived from the assumptions. While I am keenly interested in high availability, this paper will ignore that issue and focus on scalability alone. In particular, we focus on the implications that fall out of assuming we cannot have large-scale distributed transactions.

Goals

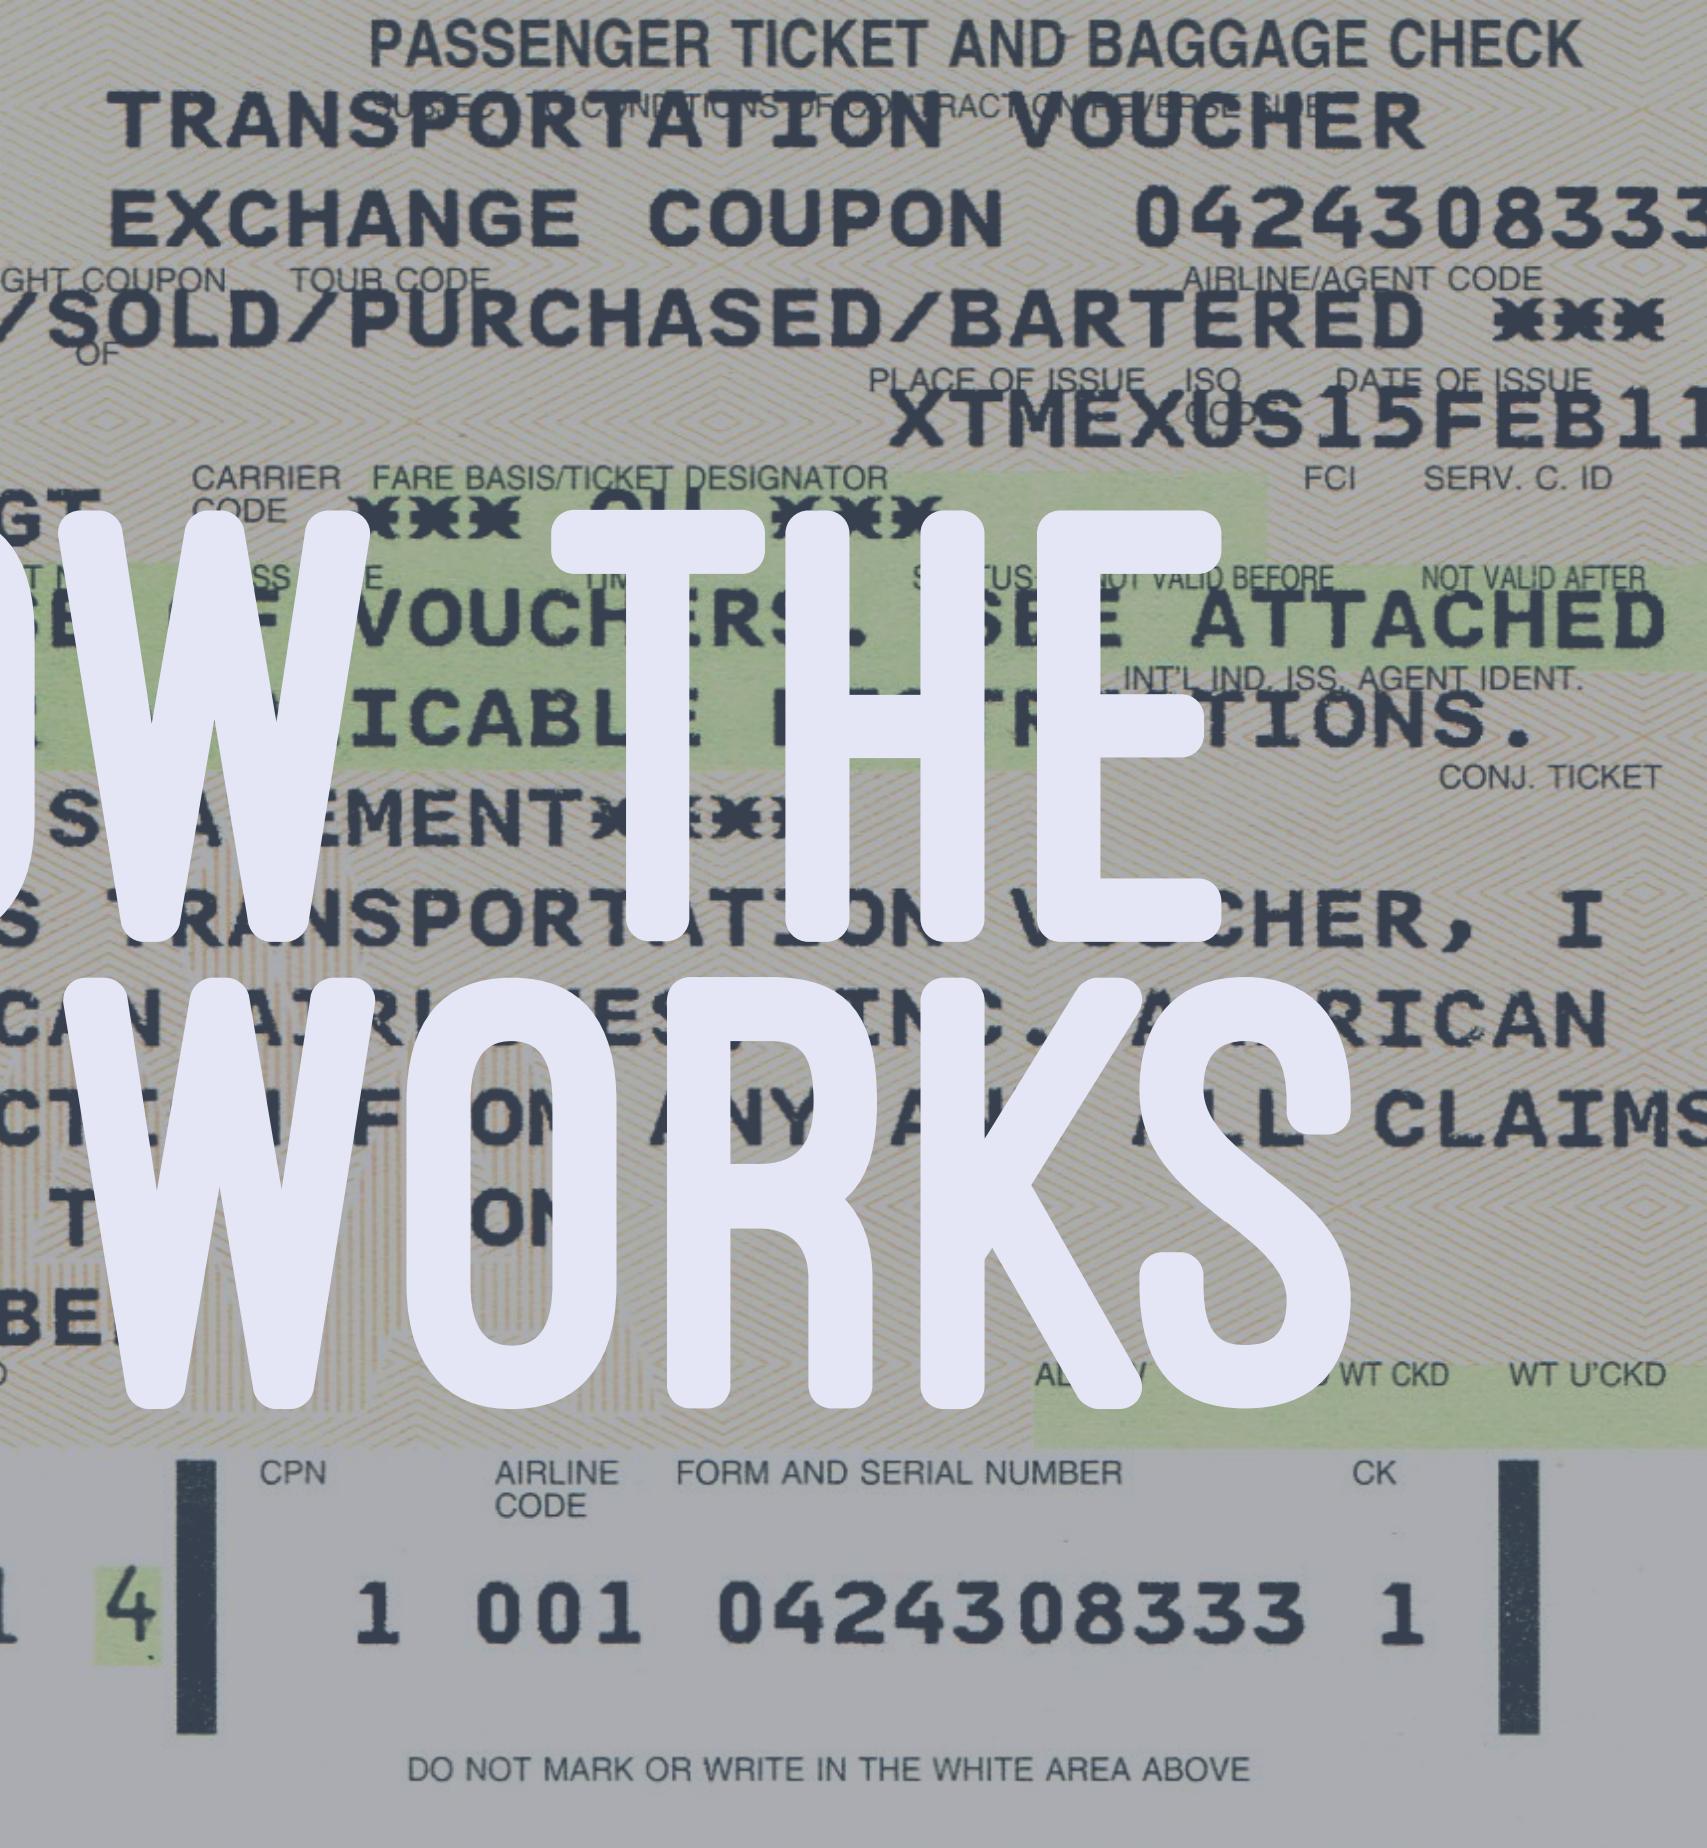
This paper has three broad goals:

- Discuss Scalable Applications

Many of the requirements for the design of scalable systems are understood implicitly by many application designers who build large systems.

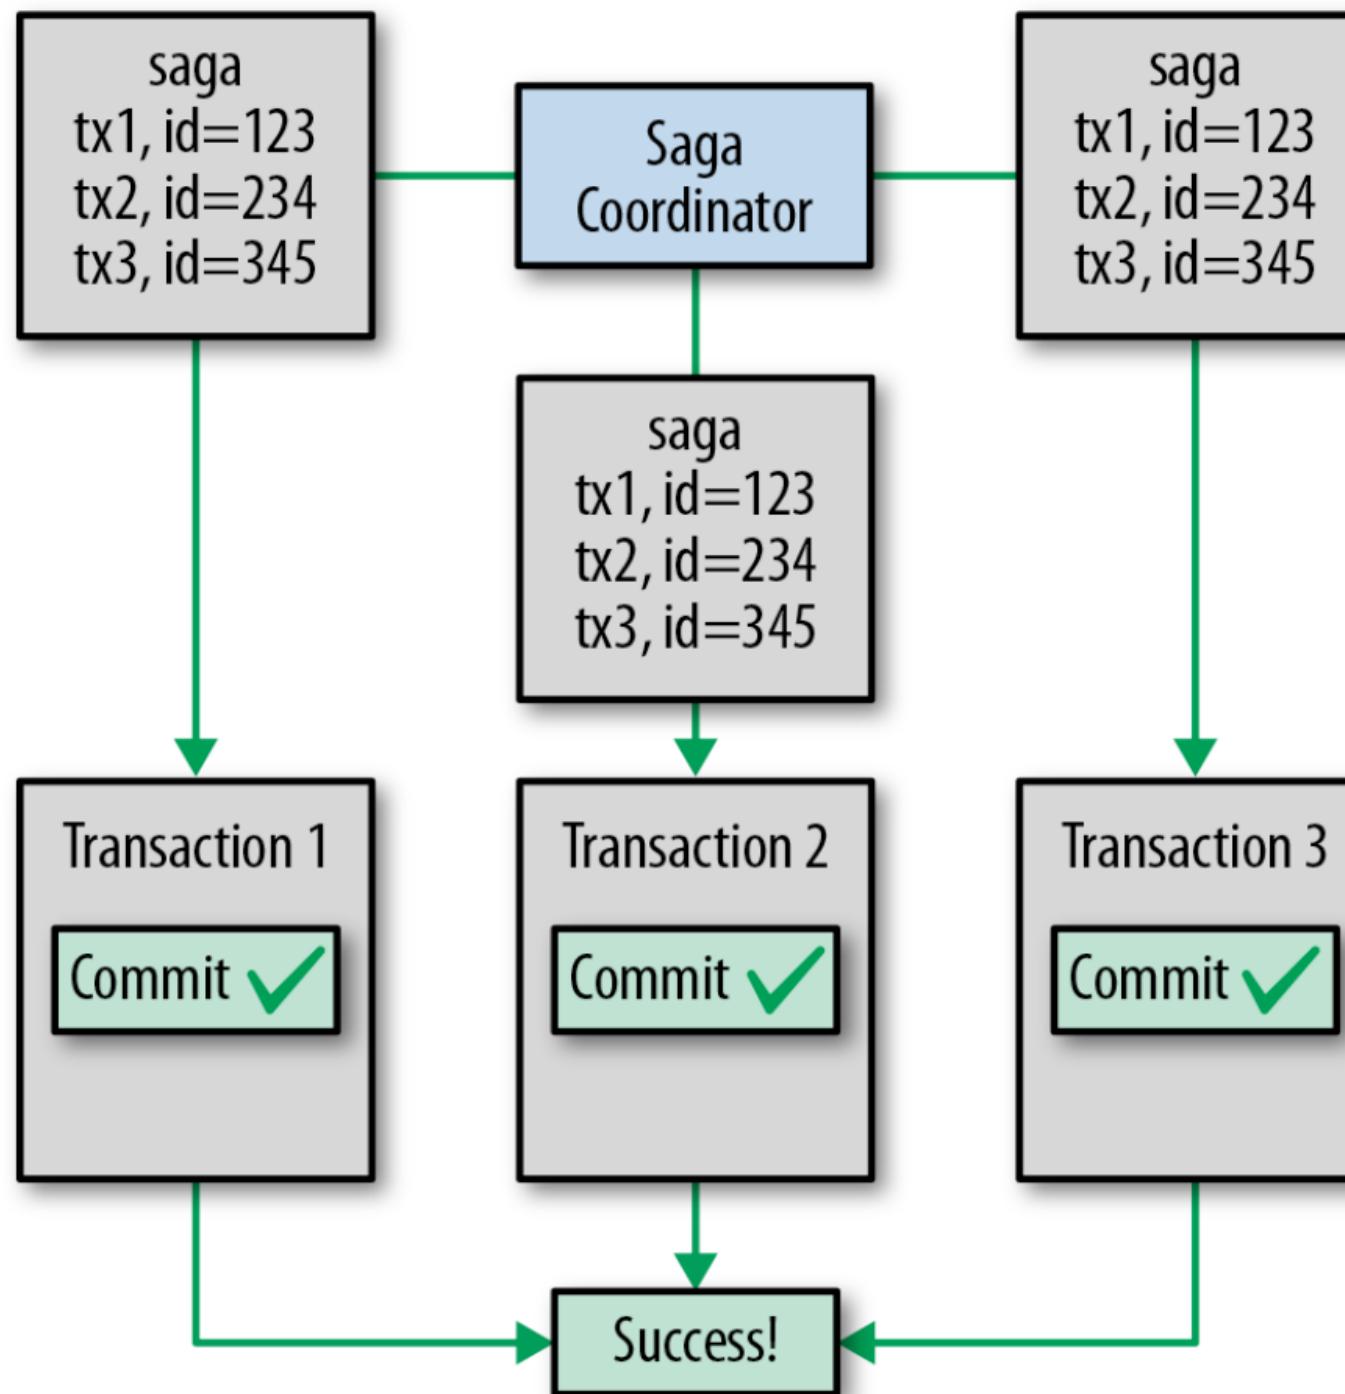
USE A PROTOCOL OF
GUESS.
APOLOGIZE.
COMPENSATE.

IT'S HOW THE
WORLD WORKS

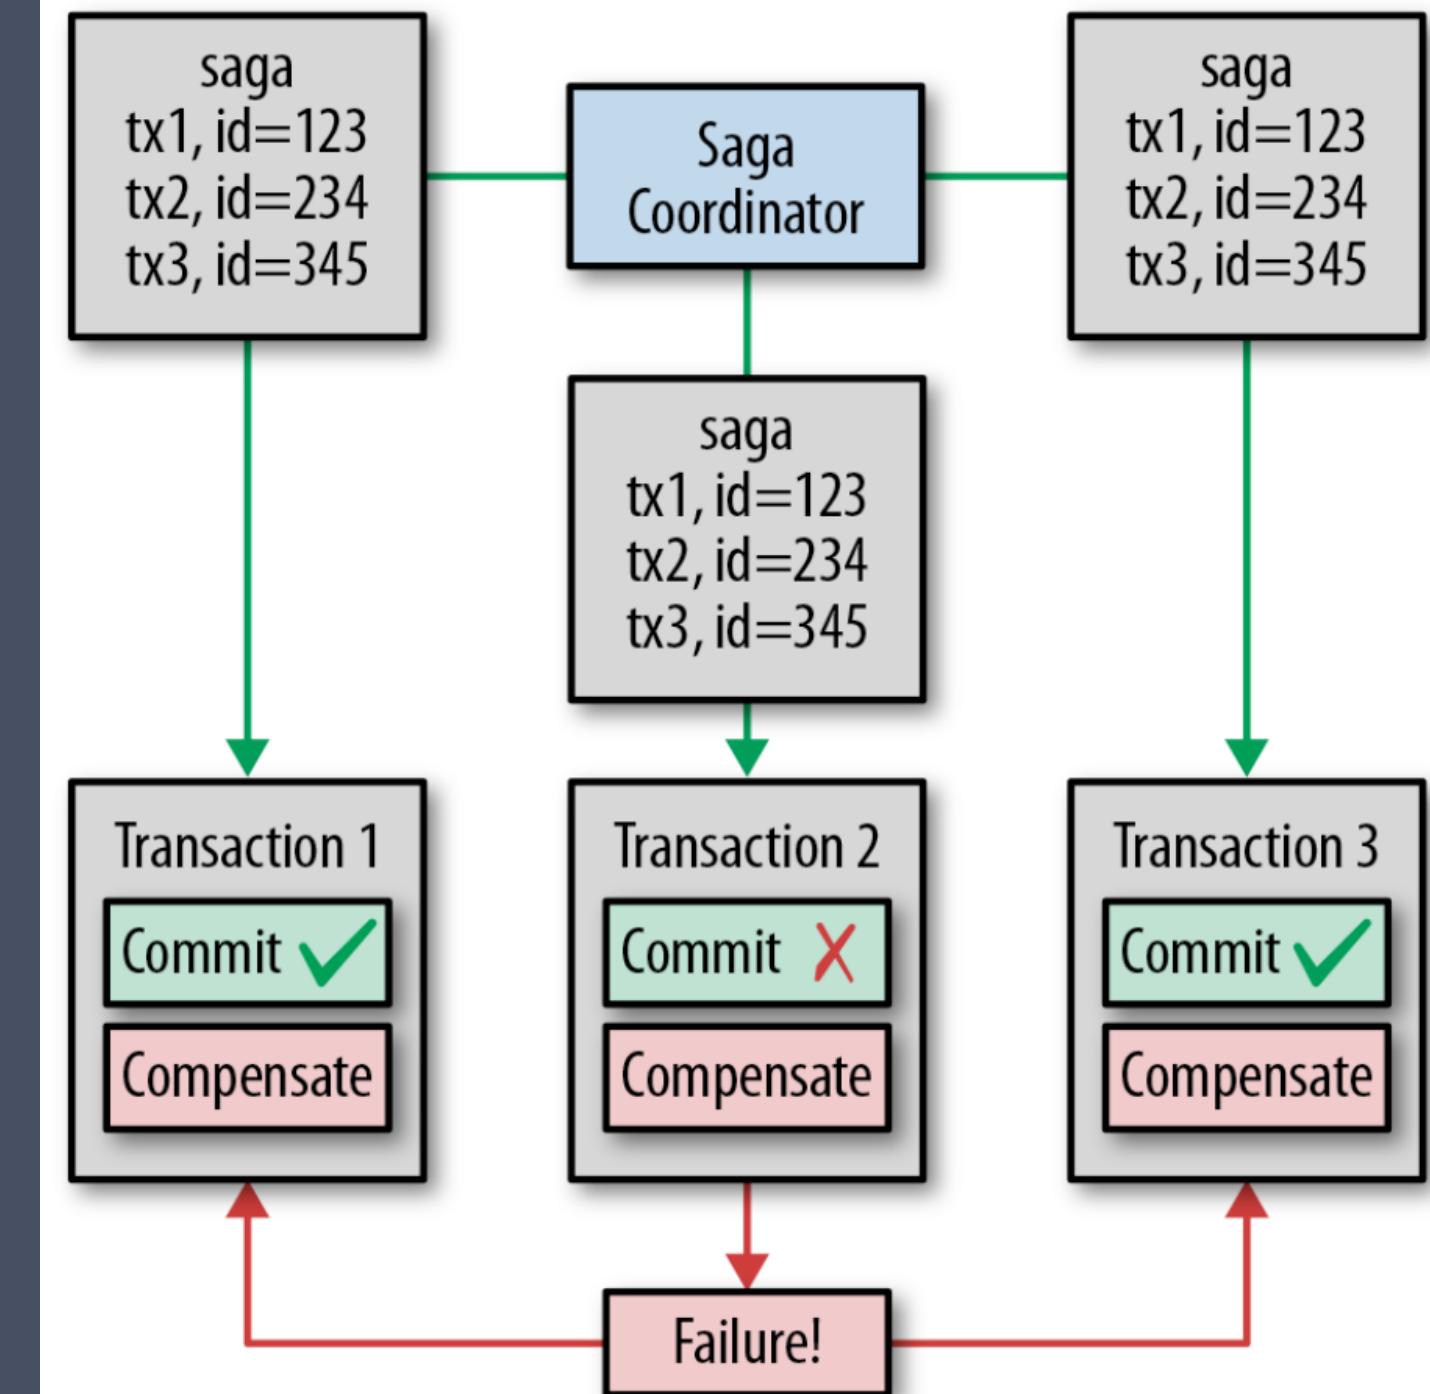


**USE SAGAS
NOT DISTRIBUTED TRANSACTIONS**

The saga pattern is all about failure management.



If failure is detected in one or more transactions of the saga...



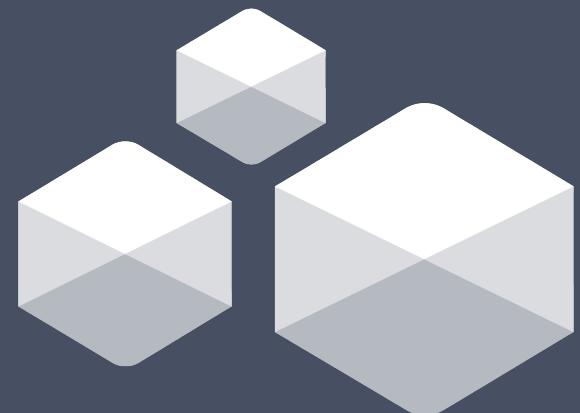
A saga knows all of the steps involved in a long running transaction.

...all transactions can be undone using compensating actions.

IN SUMMARY

1. DON'T BUILD MICROLITHS
2. MICROSERVICES COME IN (DISTRIBUTED) SYSTEMS
3. MICROSERVICES COME AS (MICRO)SYSTEMS
4. EMBRACE THE REACTIVE PRINCIPLES
5. EMBRACE EVENT-FIRST DDD & PERSISTENCE
6. PROFIT!

TRY THE
LAGOM
MICROSERVICES FRAMEWORK
POWERED BY AKKA & PLAY
LAGOMFRAMEWORK.COM

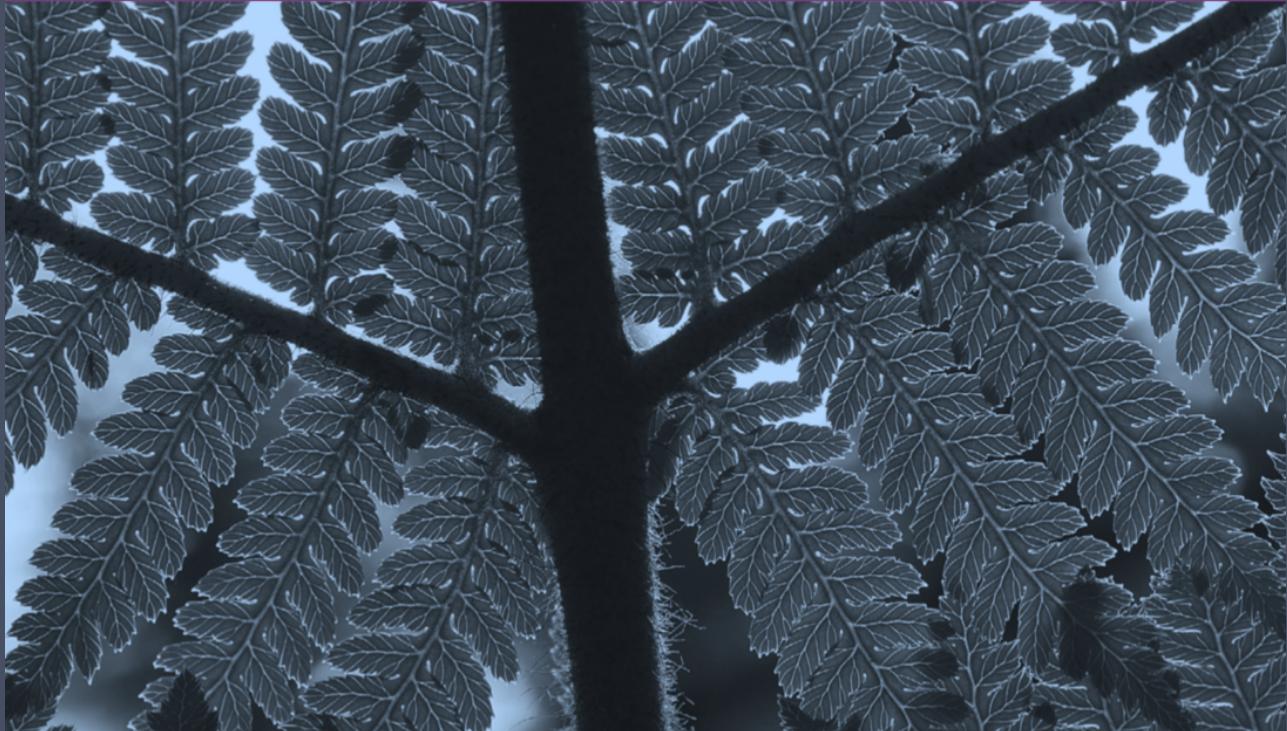


lagom

LEARN MORE
DOWNLOAD MY BOOK FOR FREE AT:
[BIT.LY/REACTIVE-MICROSYSTEMS](http://bit.ly/reactive-microsystems)

Reactive Microsystems

The Evolution of Microservices at Scale



Jonas Bonér