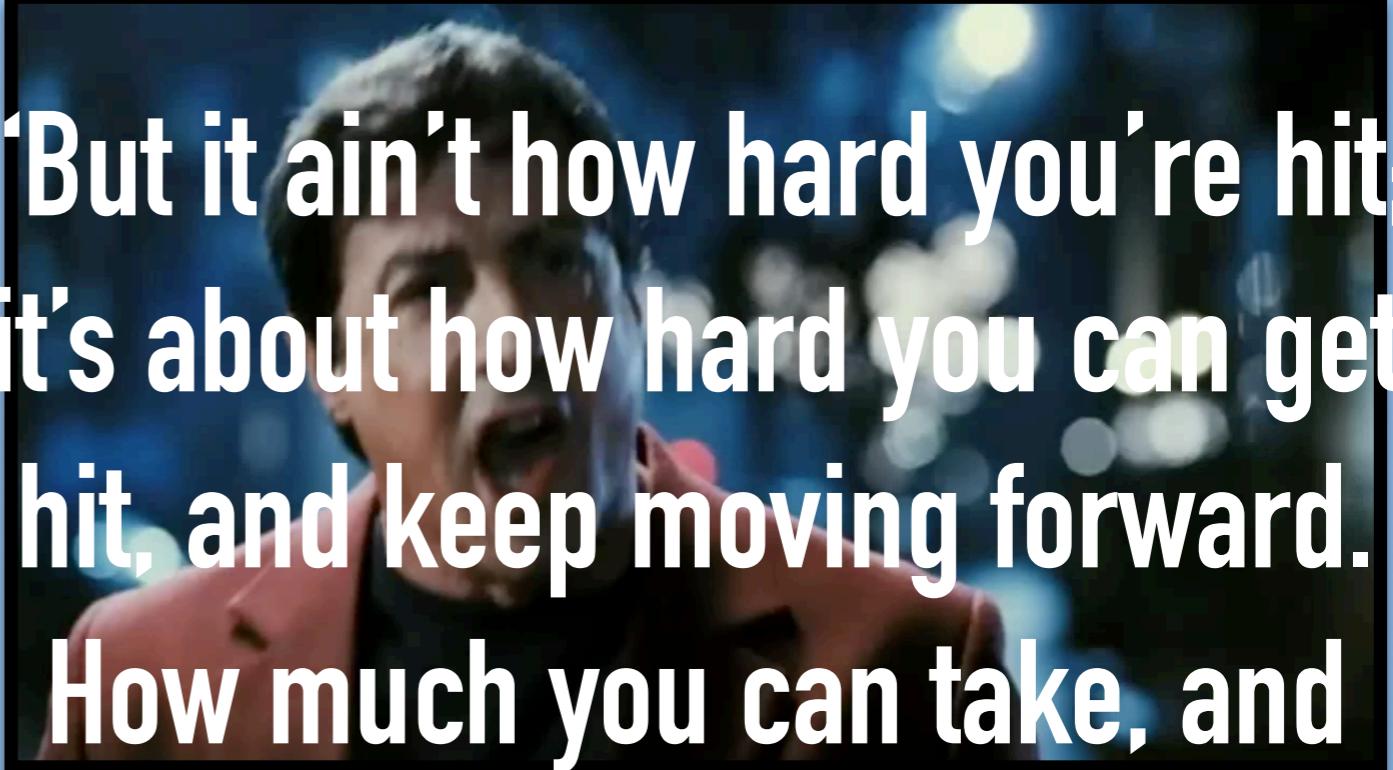


Without Resilience Nothing Else Matters

JONAS BONÉR
CTO Lightbend
@jboner

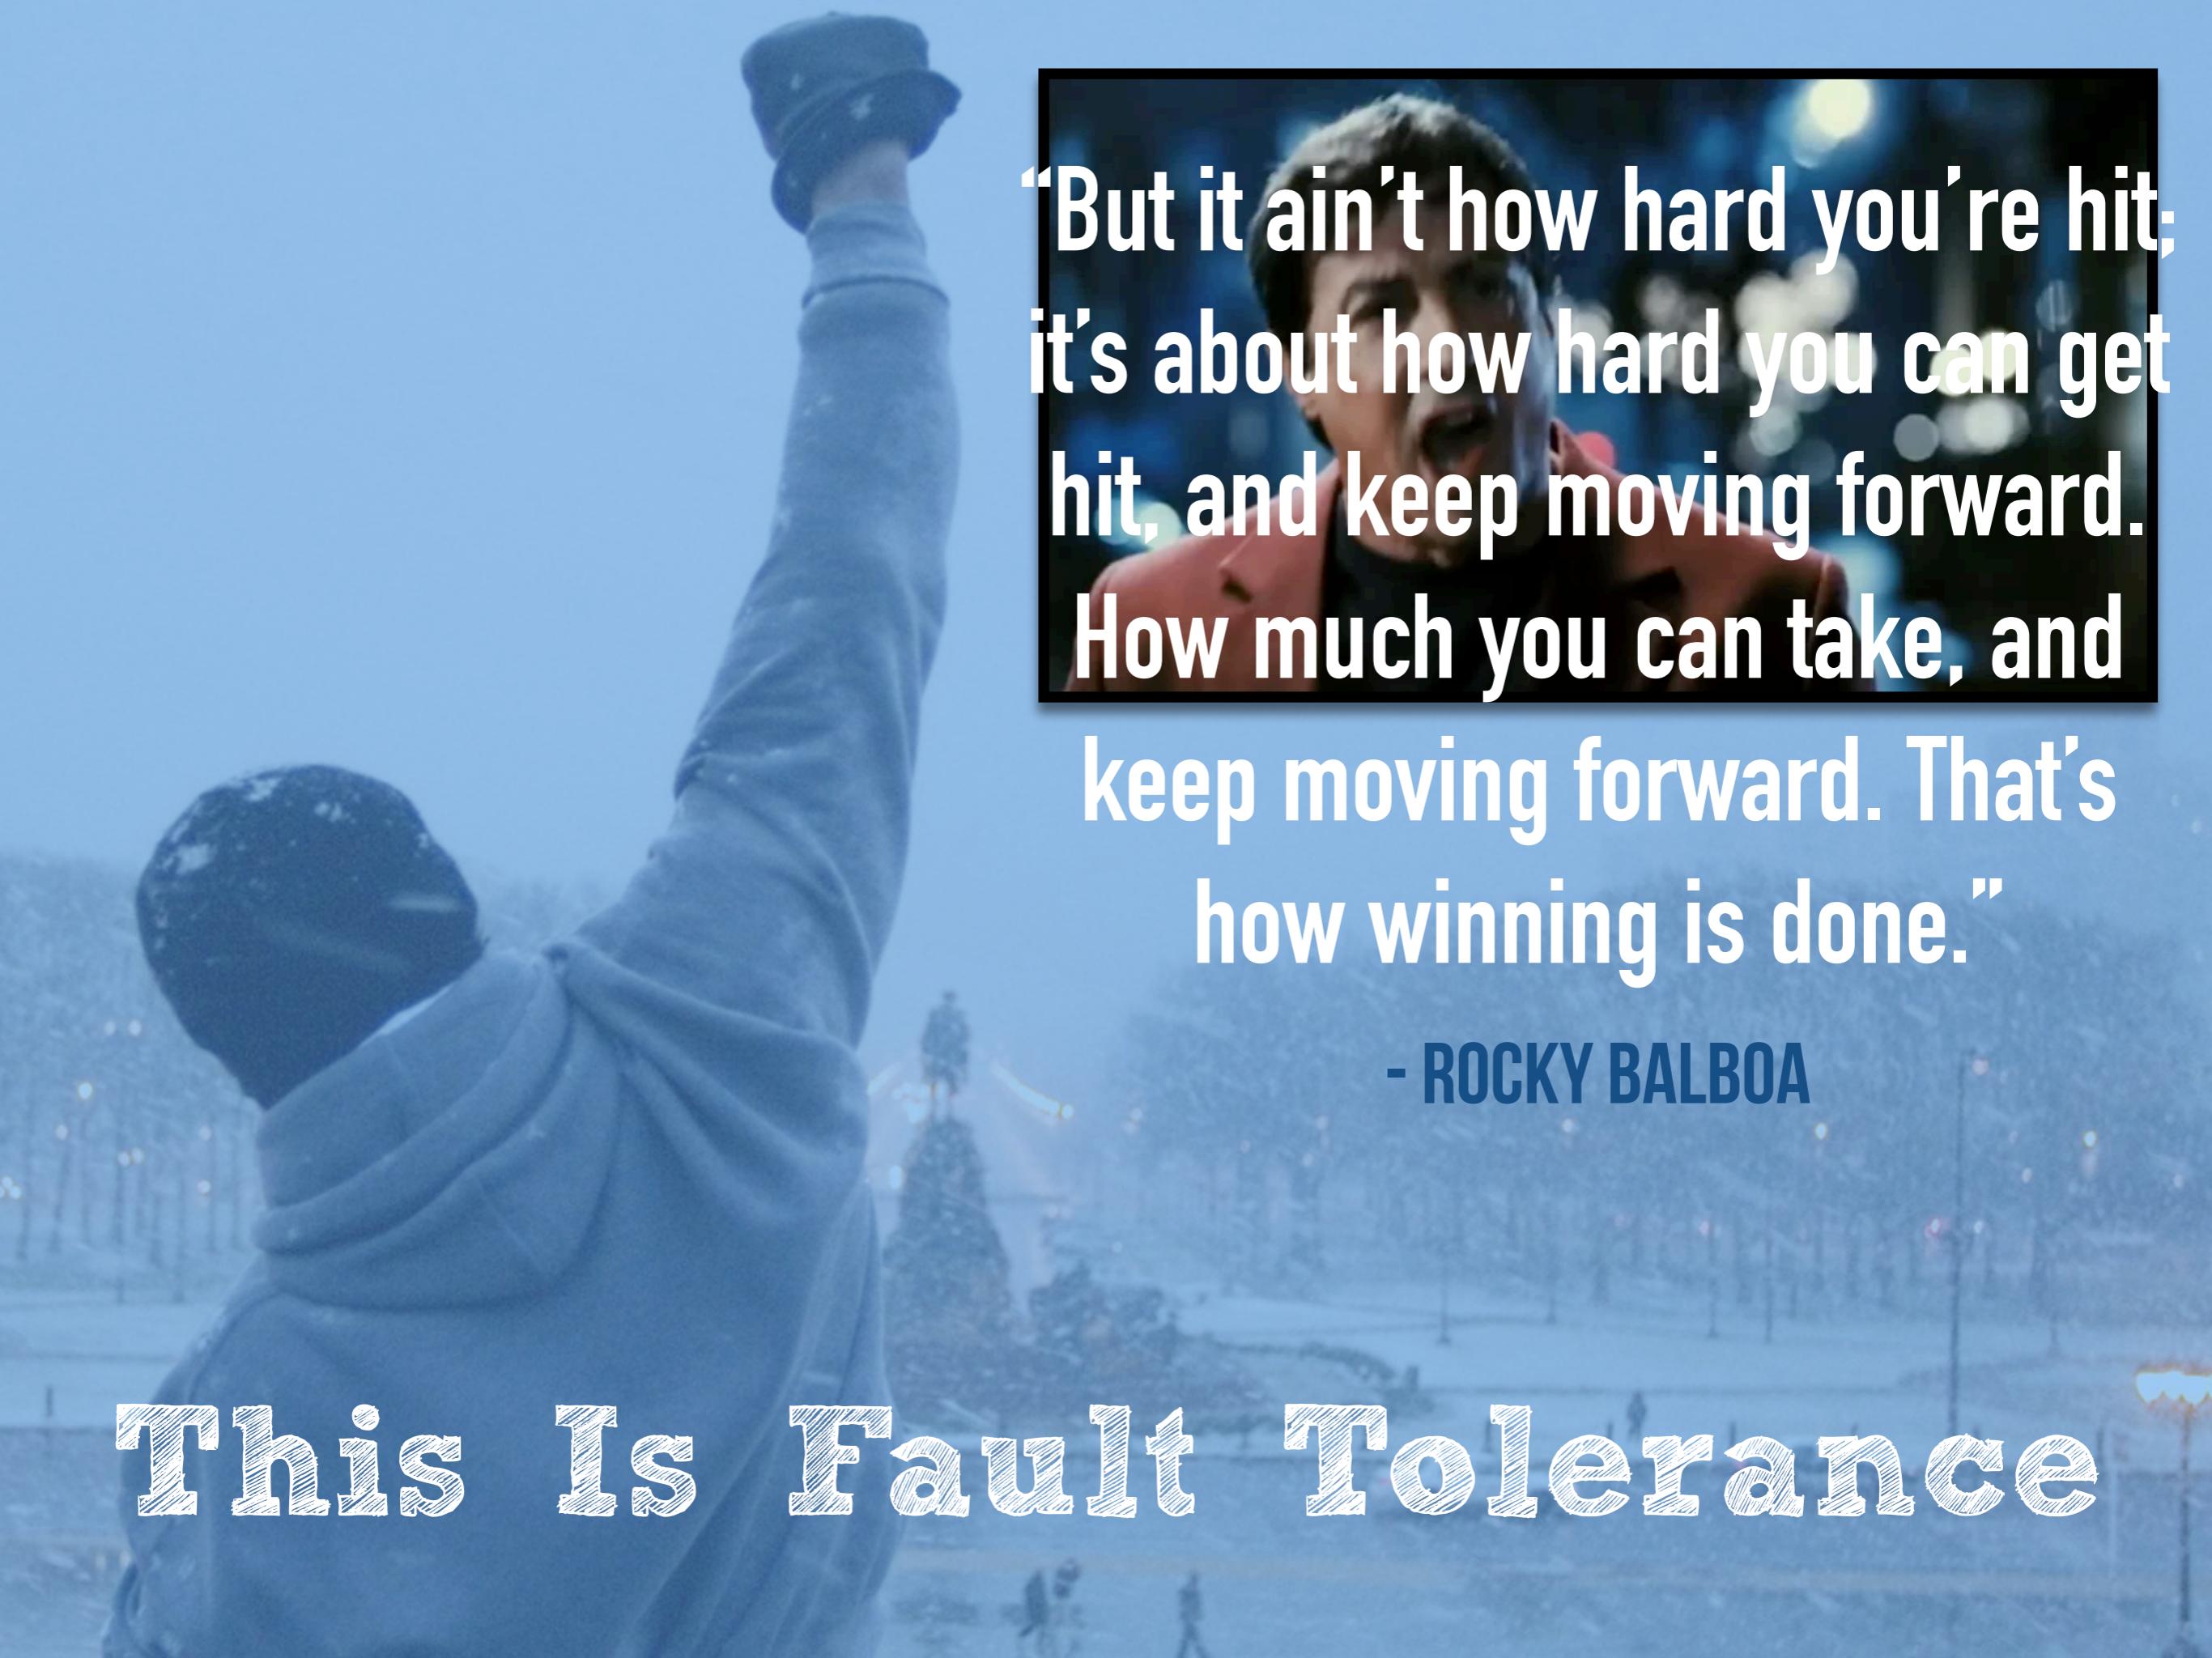




"But it ain't how hard you're hit;
it's about how hard you can get
hit, and keep moving forward.
How much you can take, and

keep moving forward. That's
how winning is done."

- ROCKY BALBOA



This Is Fault Tolerance

**Resilience
Is Beyond
Fault
Tolerance**

Resilience

“The ability of a substance or object to spring back into shape.
The capacity to recover quickly from difficulties.”

-MERRIAM WEBSTER

Antifragility

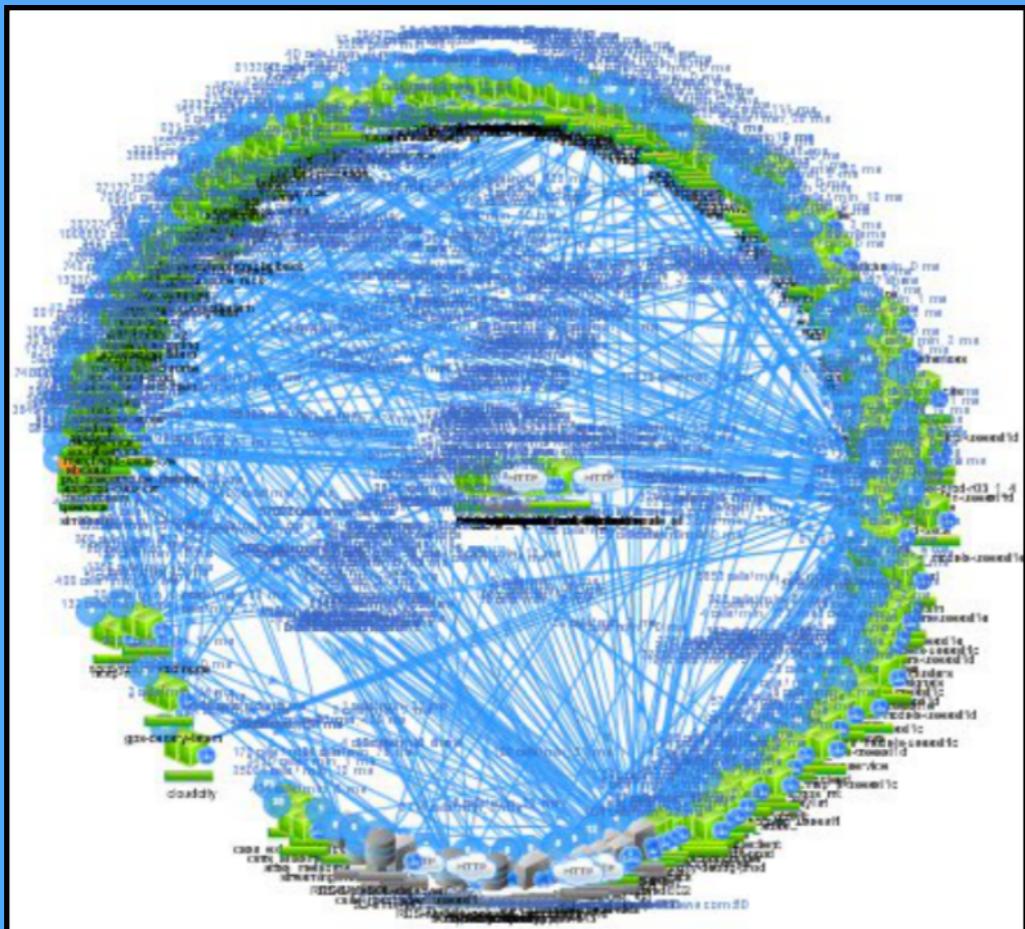
“Antifragility is beyond resilience and robustness. The resilient resists shock and stays the same; the antifragile gets better.”

- NASSEM NICHOLAS TALEB

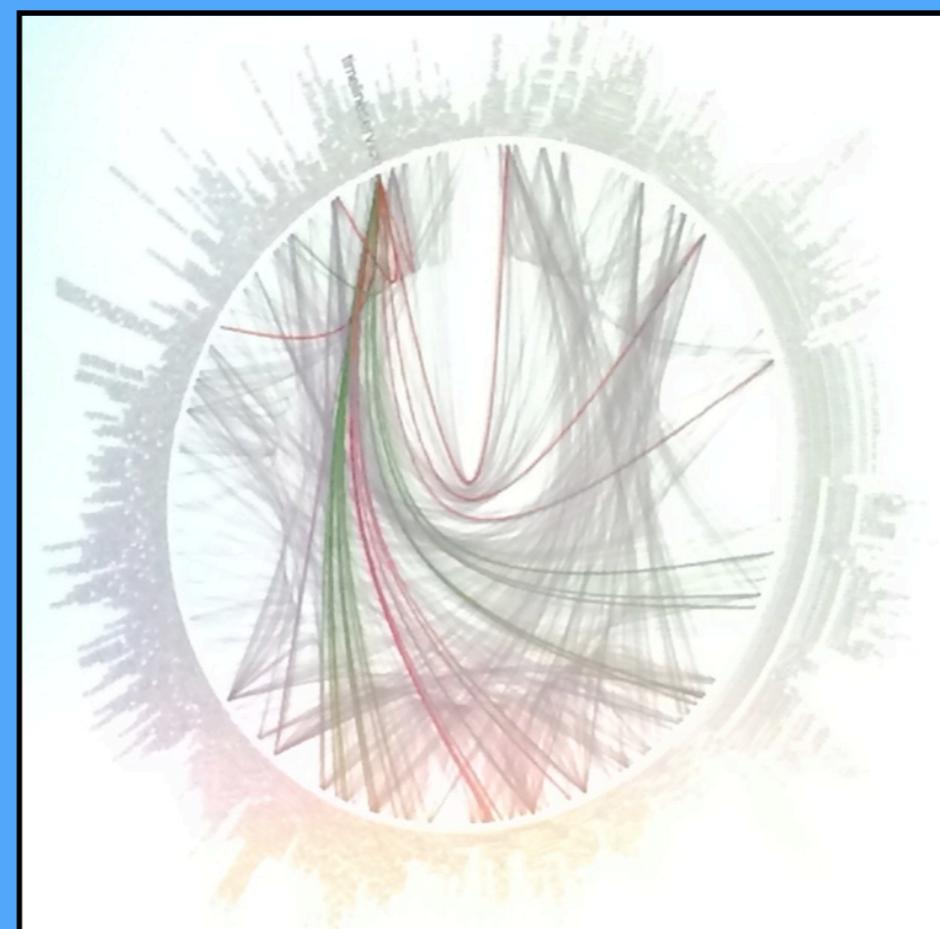
“We can model and understand in isolation.
But, when released into competitive nominally
regulated societies, their connections proliferate,
their interactions and interdependencies multiply,
their complexities mushroom.
And we are caught short.”

- SIDNEY DEKKER

Software Systems Today Are Incredibly Complex



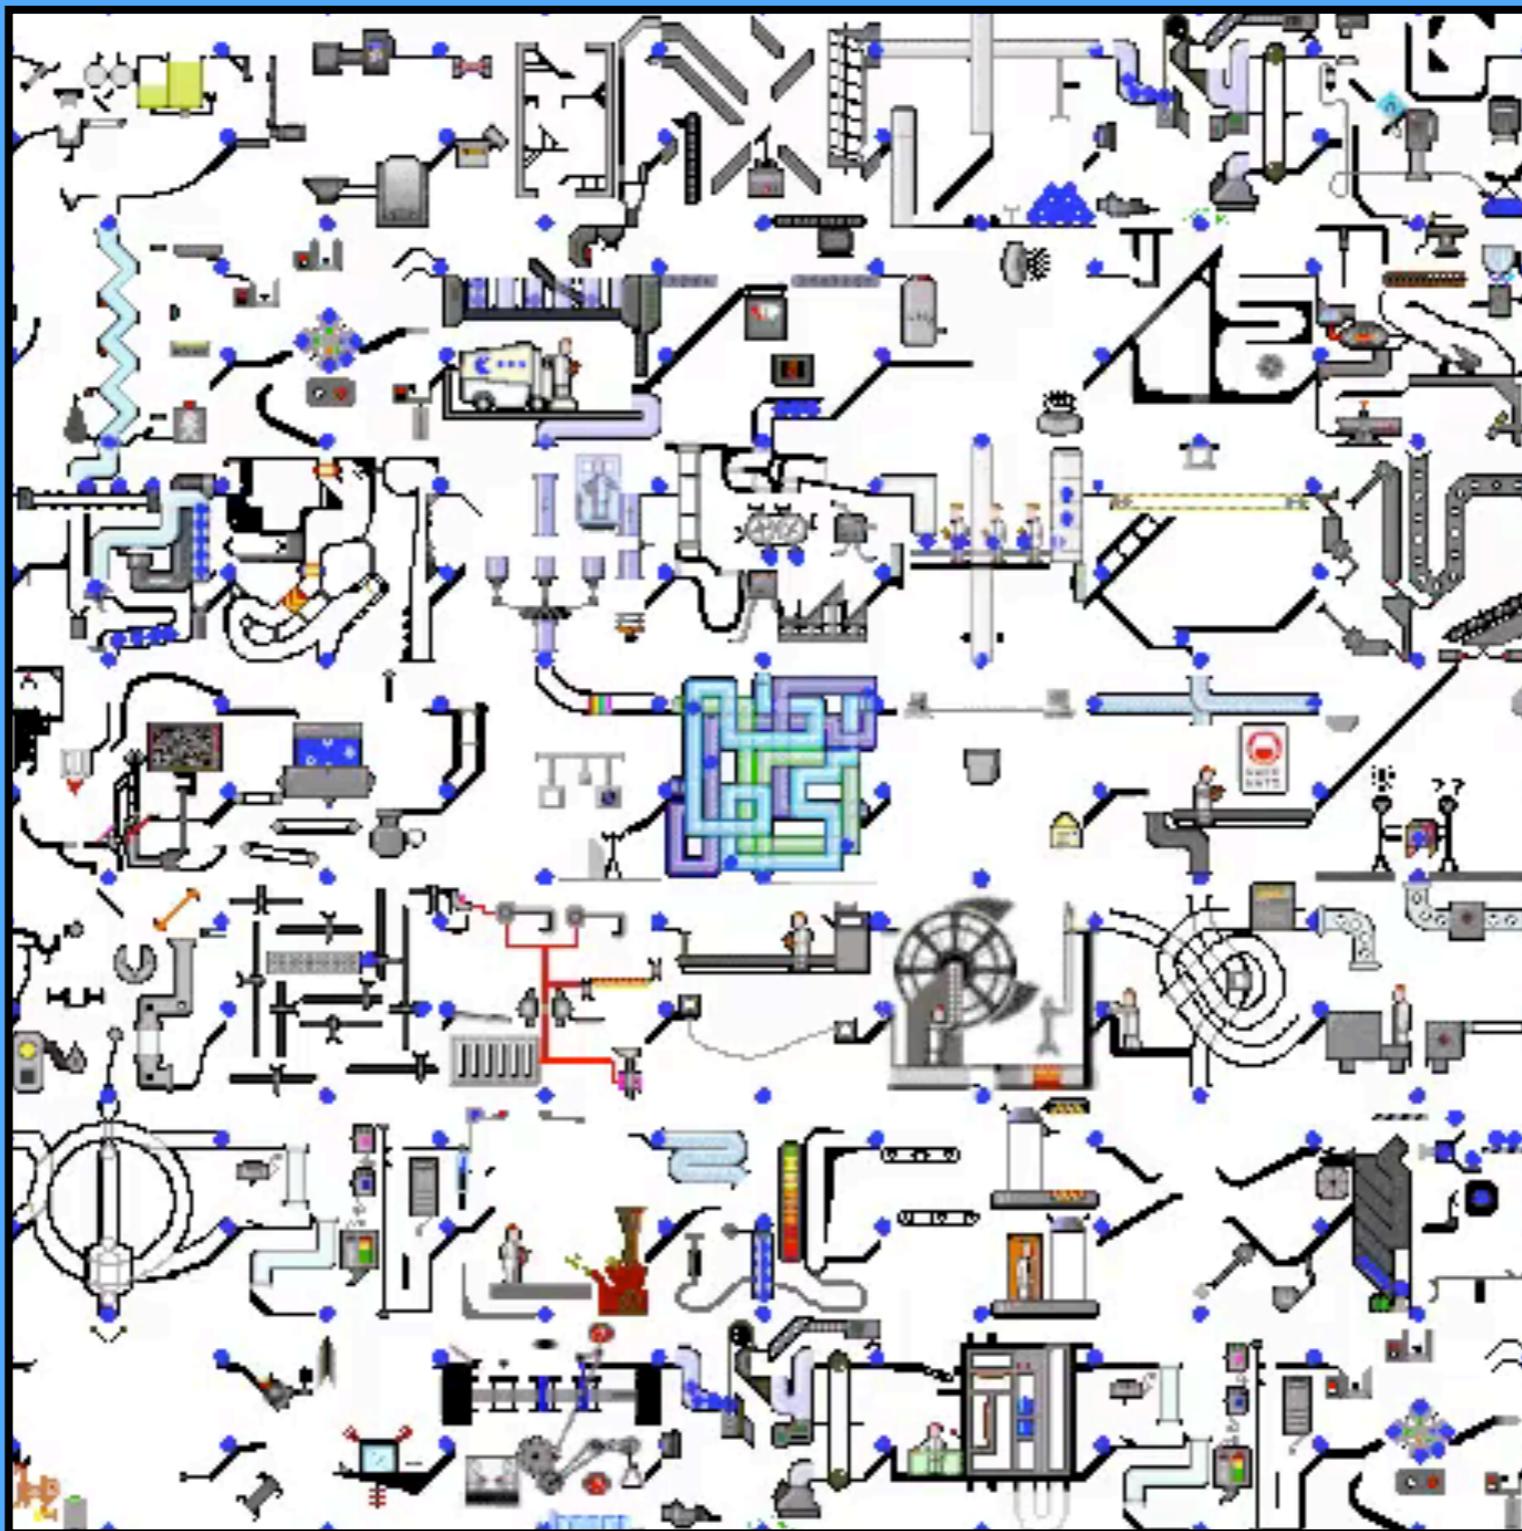
NETFLIX



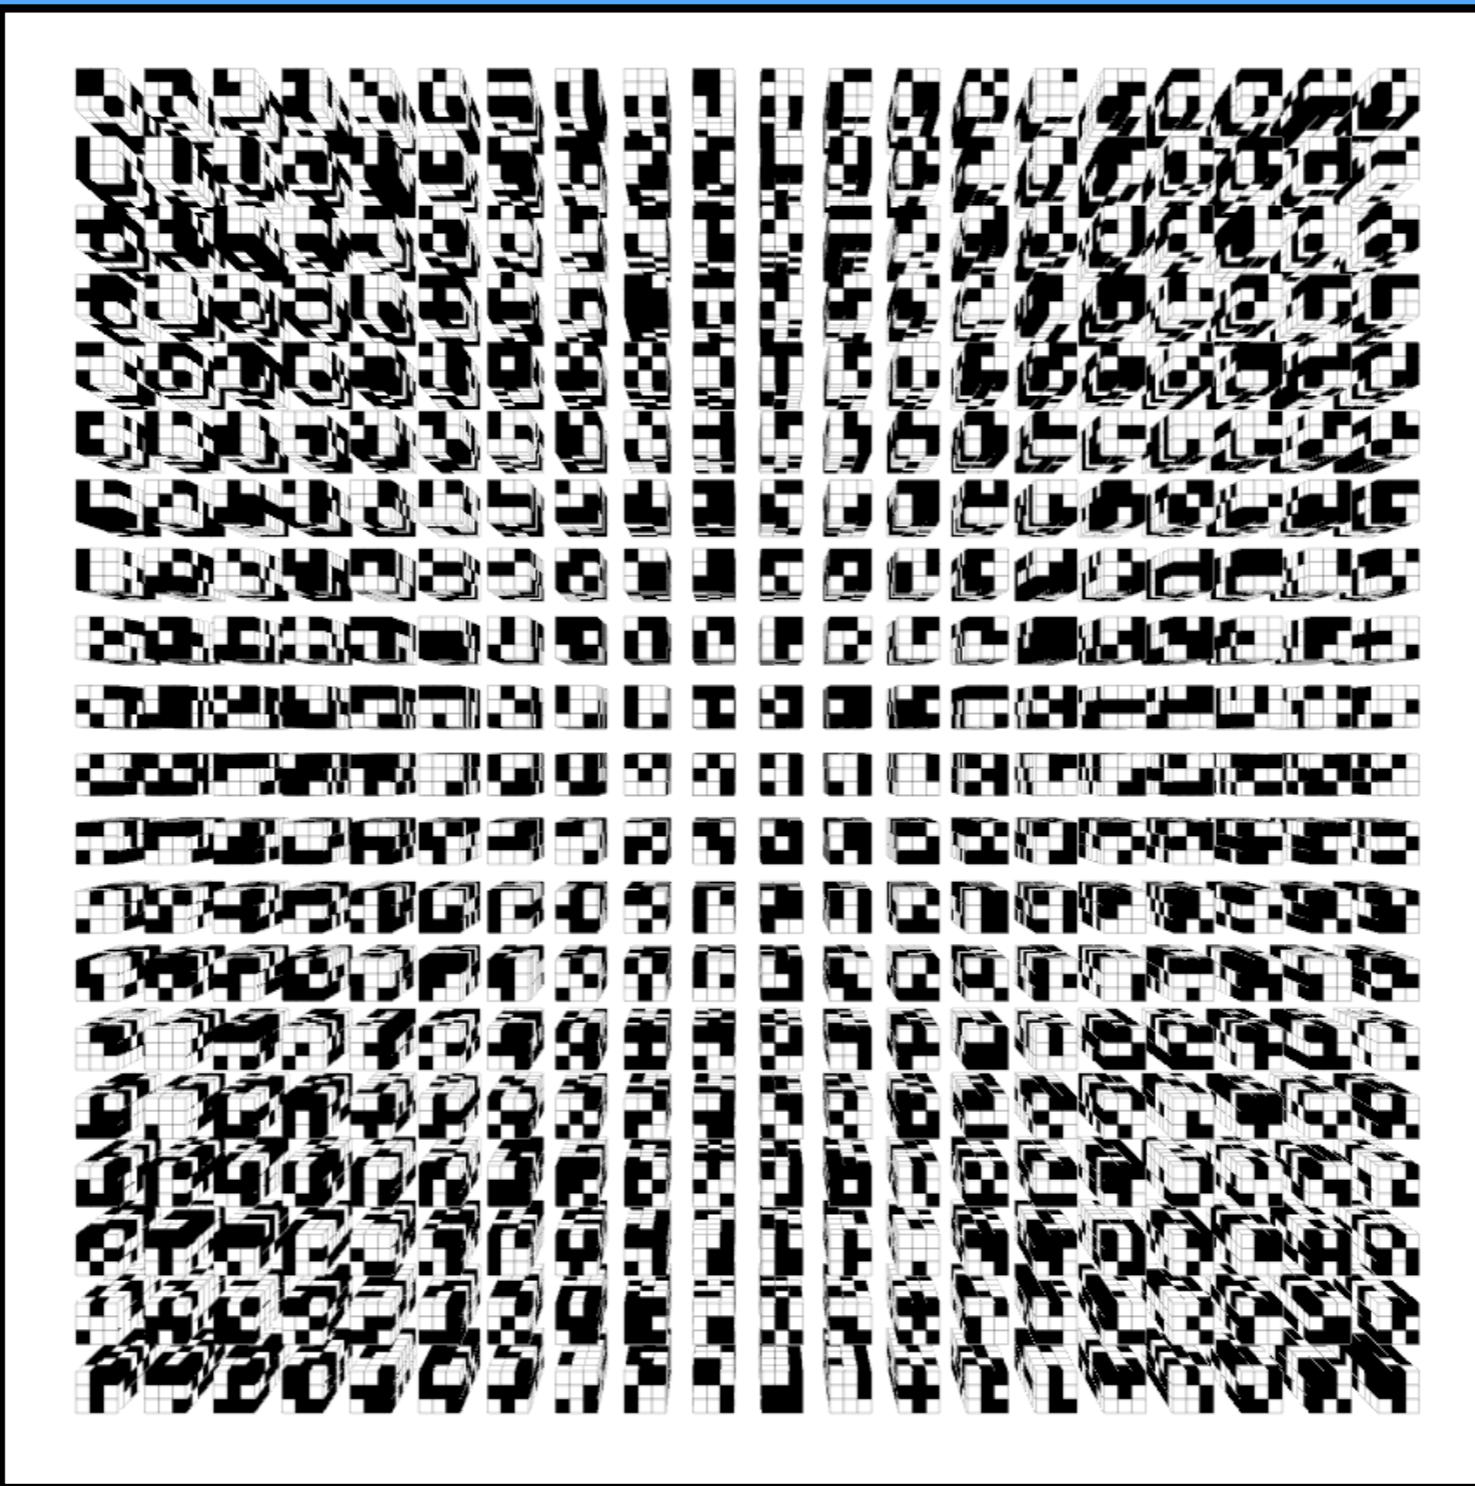
TWITTER

WE NEED TO STUDY
Resilience in
Complex
Systems

Complicated System



Complex System



COMPLICATED ≠ **COMPLEX**

“Complex systems run in degraded mode.”
“Complex systems run as broken systems.”

- RICHARD COOK

“Counterintuitive. That’s [Jay] Forrester’s word to describe complex systems.

Leverage points are not intuitive. Or if they are, we intuitively use them backward, systematically worsening whatever problems we are trying to solve.”

- DONELLA MEADOWS

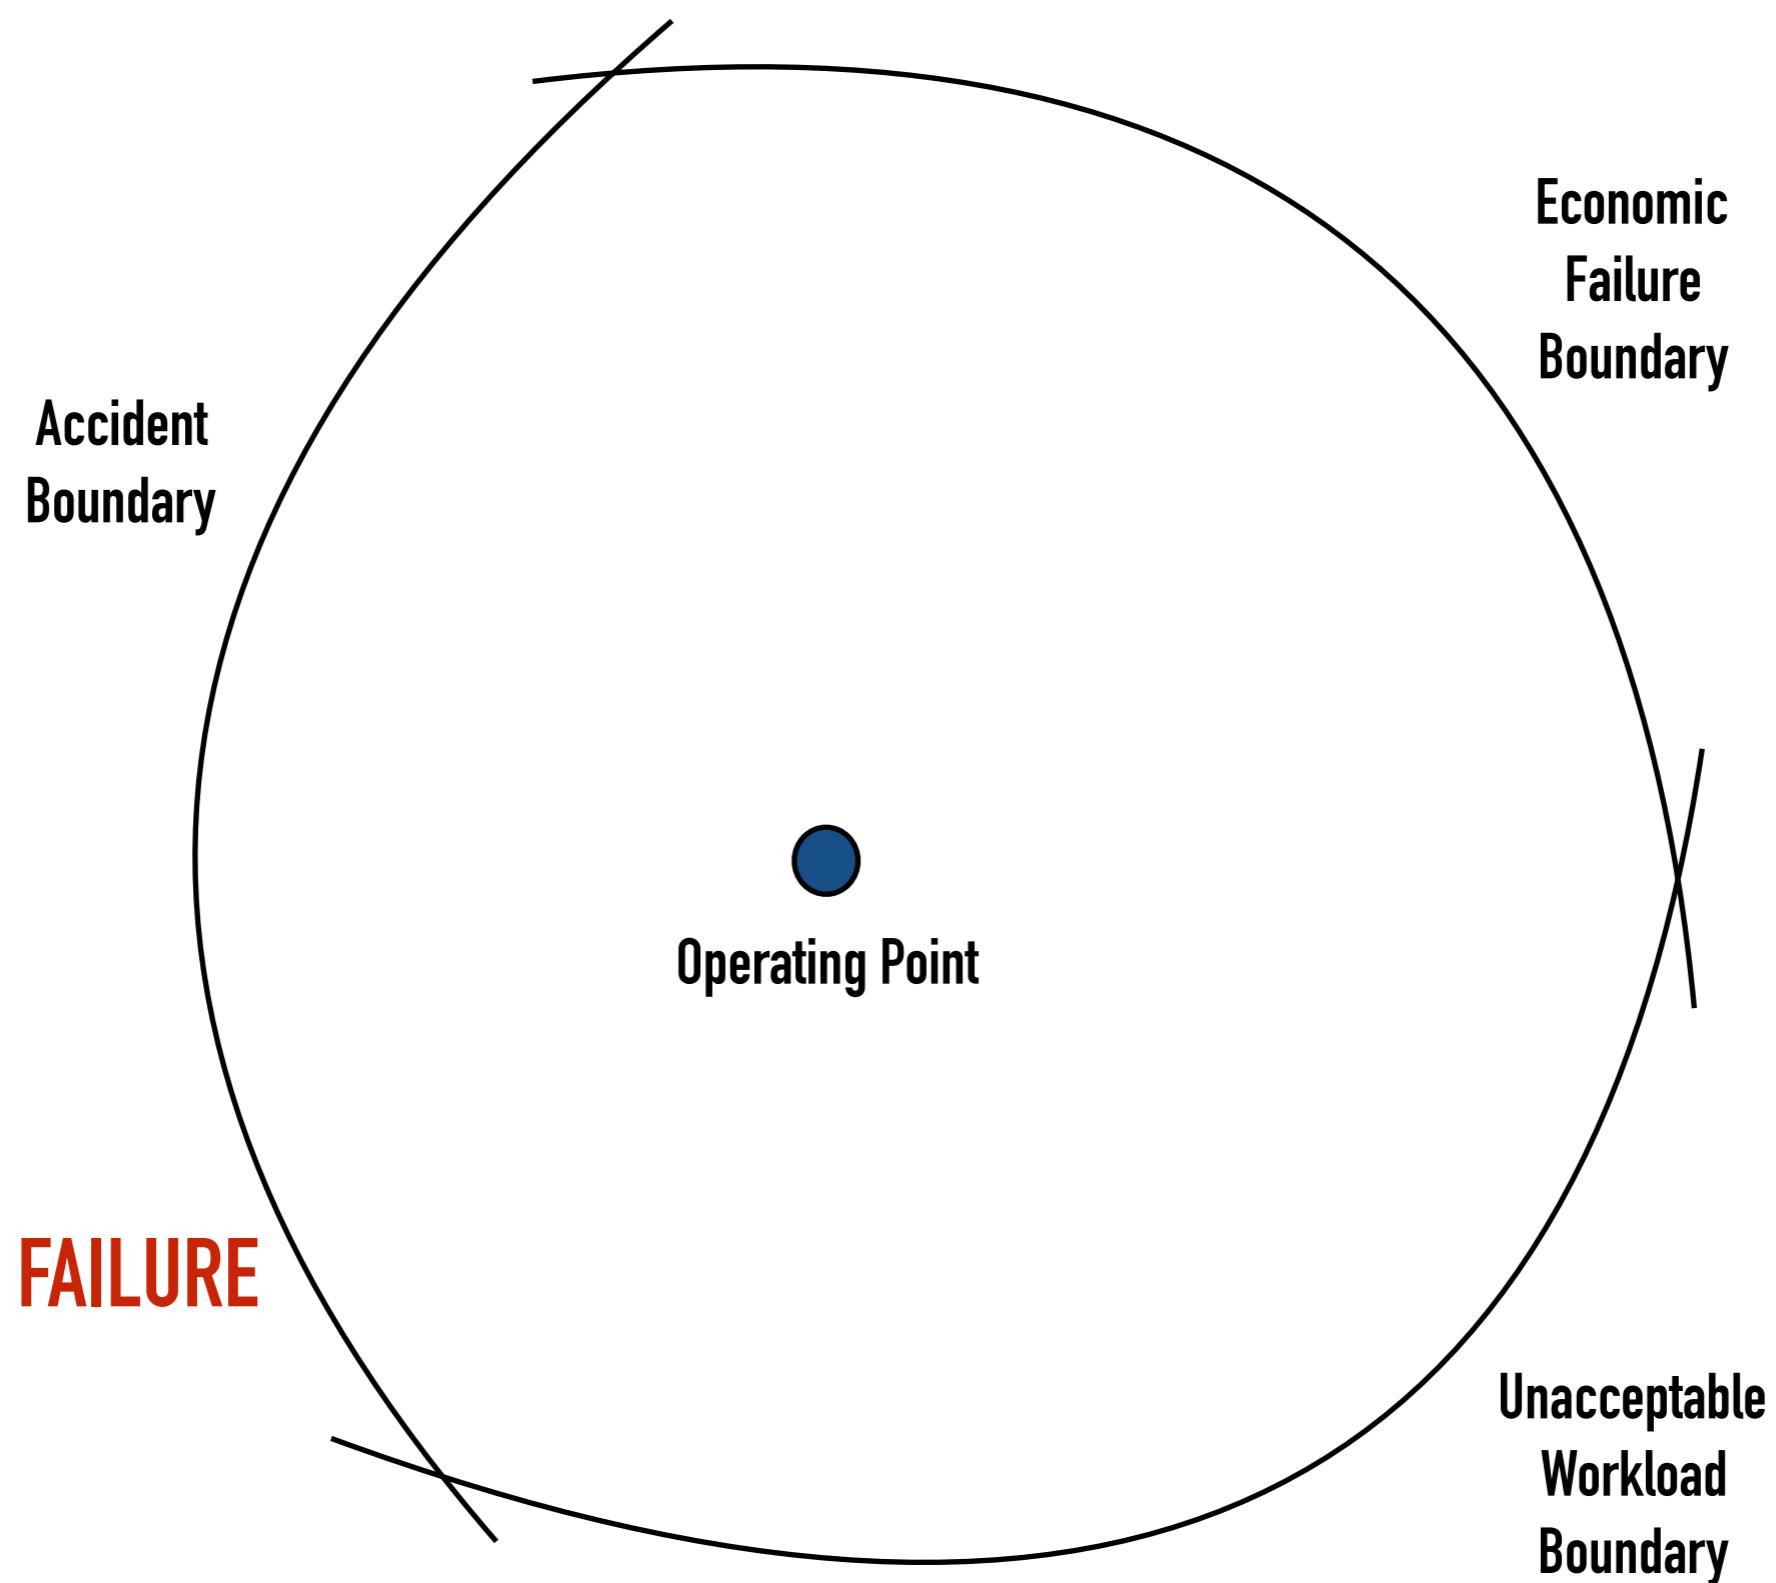
“Humans should not be involved in setting timeouts.”

“Human involvement in complex systems is the biggest source of trouble.”

- BEN CHRISTENSEN, NETFLIX

Humans Generally
Make Things Worse

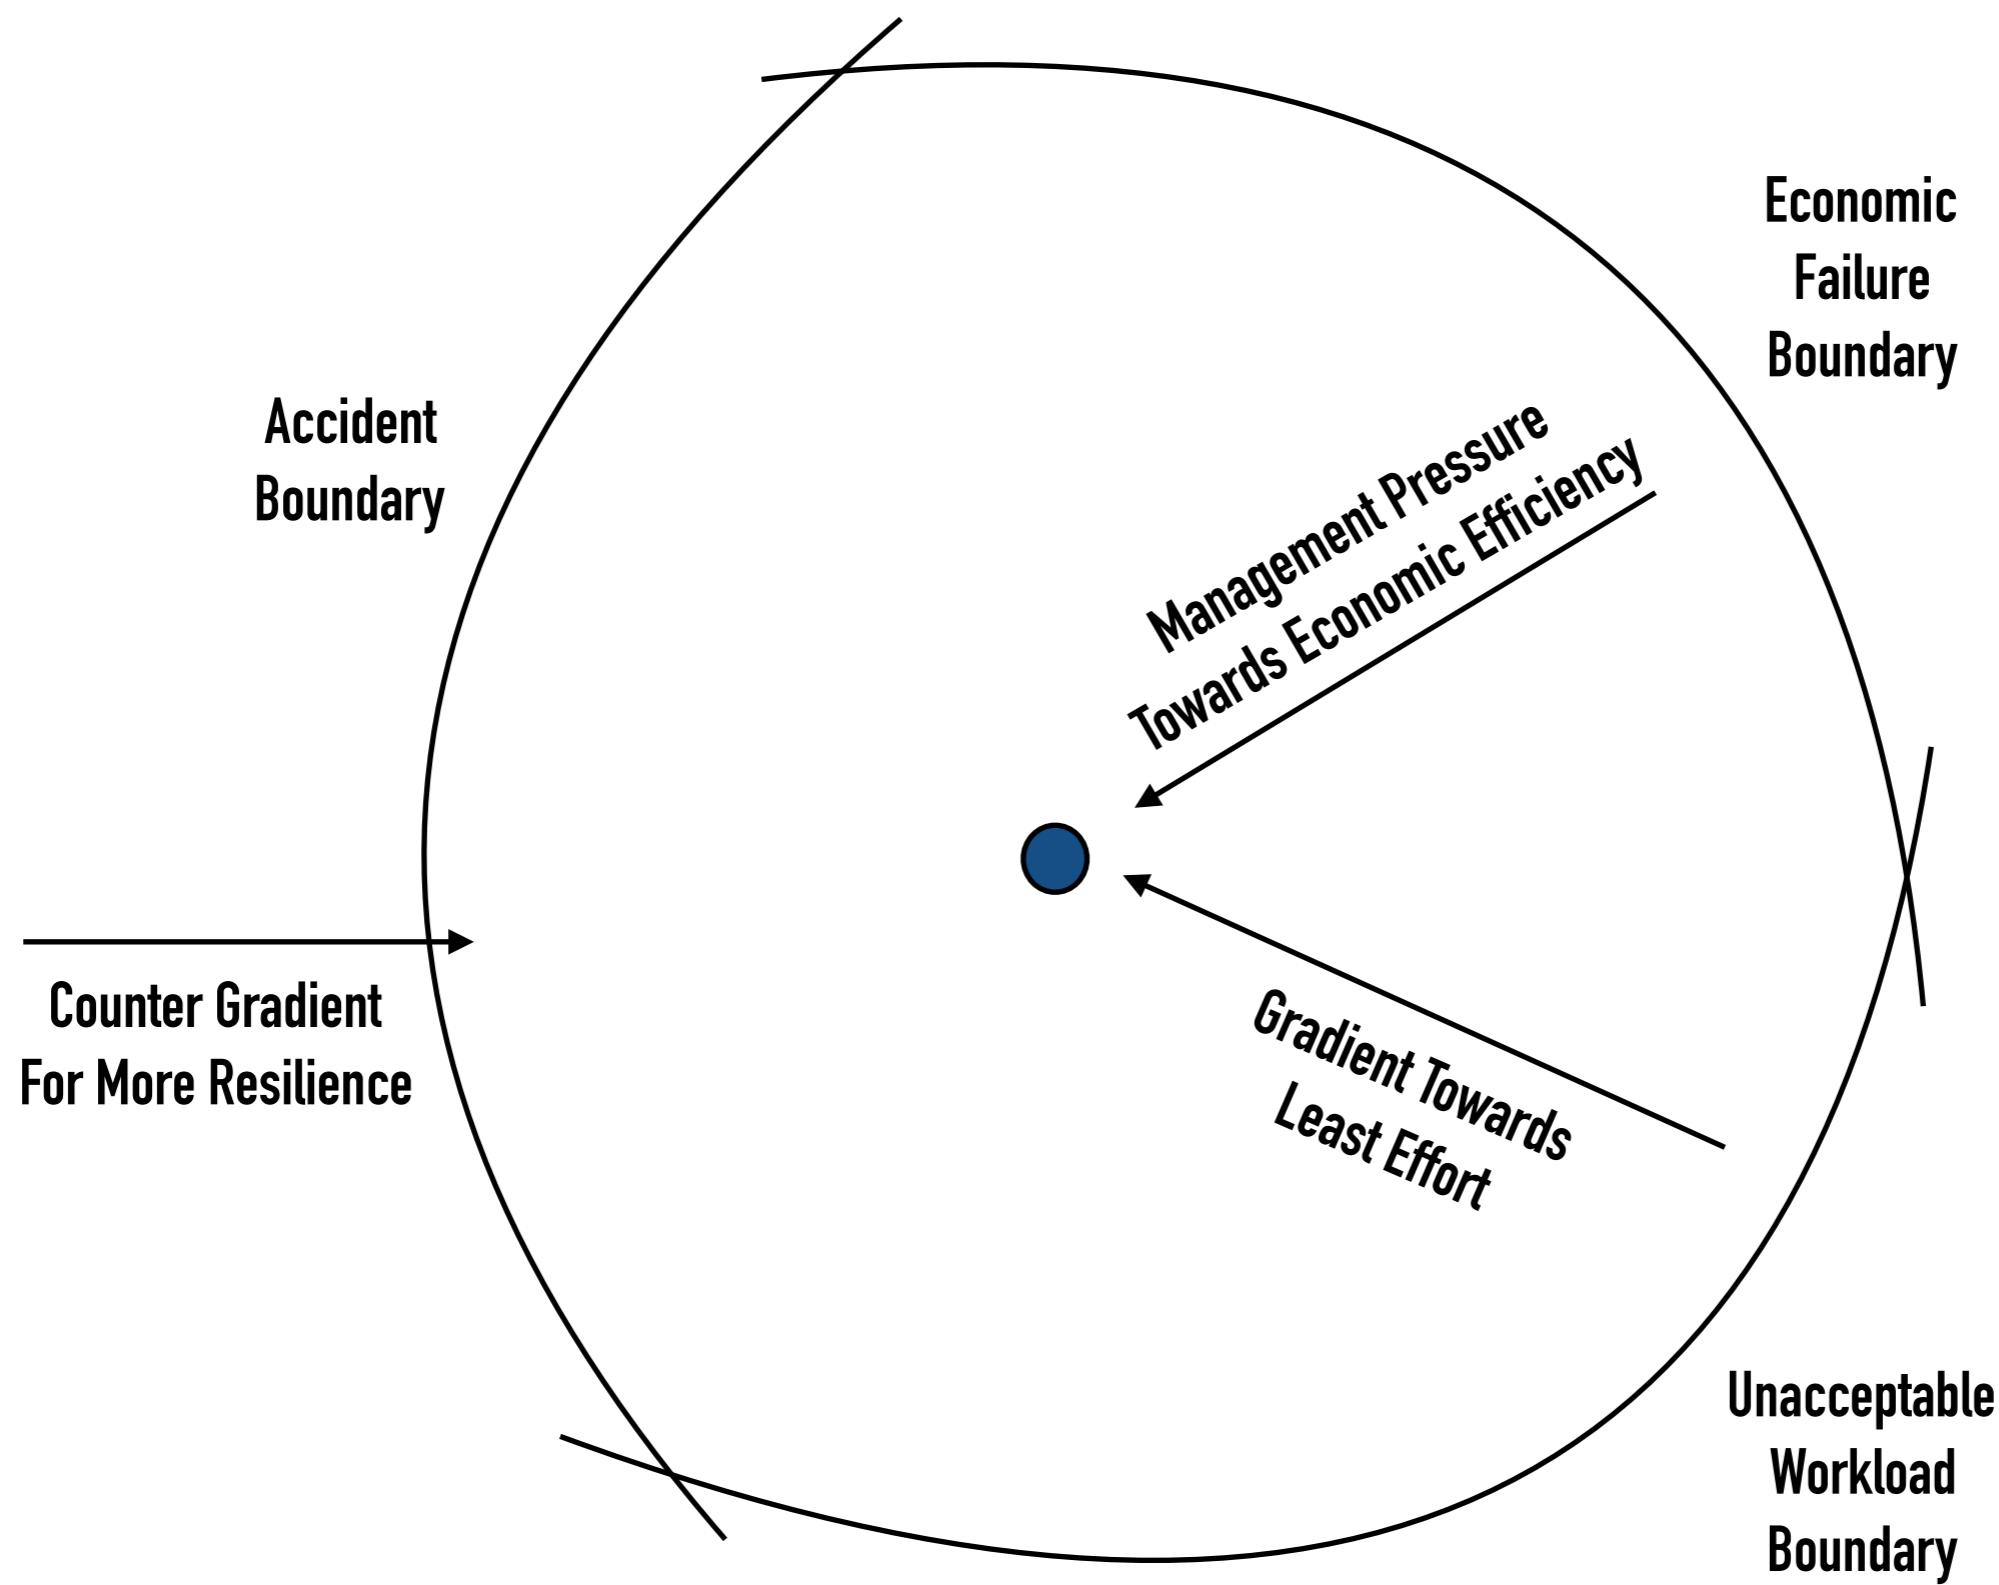
Operating at the Edge of Failure



"Going solid": a model of system dynamics and consequences for patient safety - R Cook, J Rasmussen

Resilience in complex adaptive systems: Operating at the Edge of Failure - Richard Cook - Talk at Velocity NY 2013

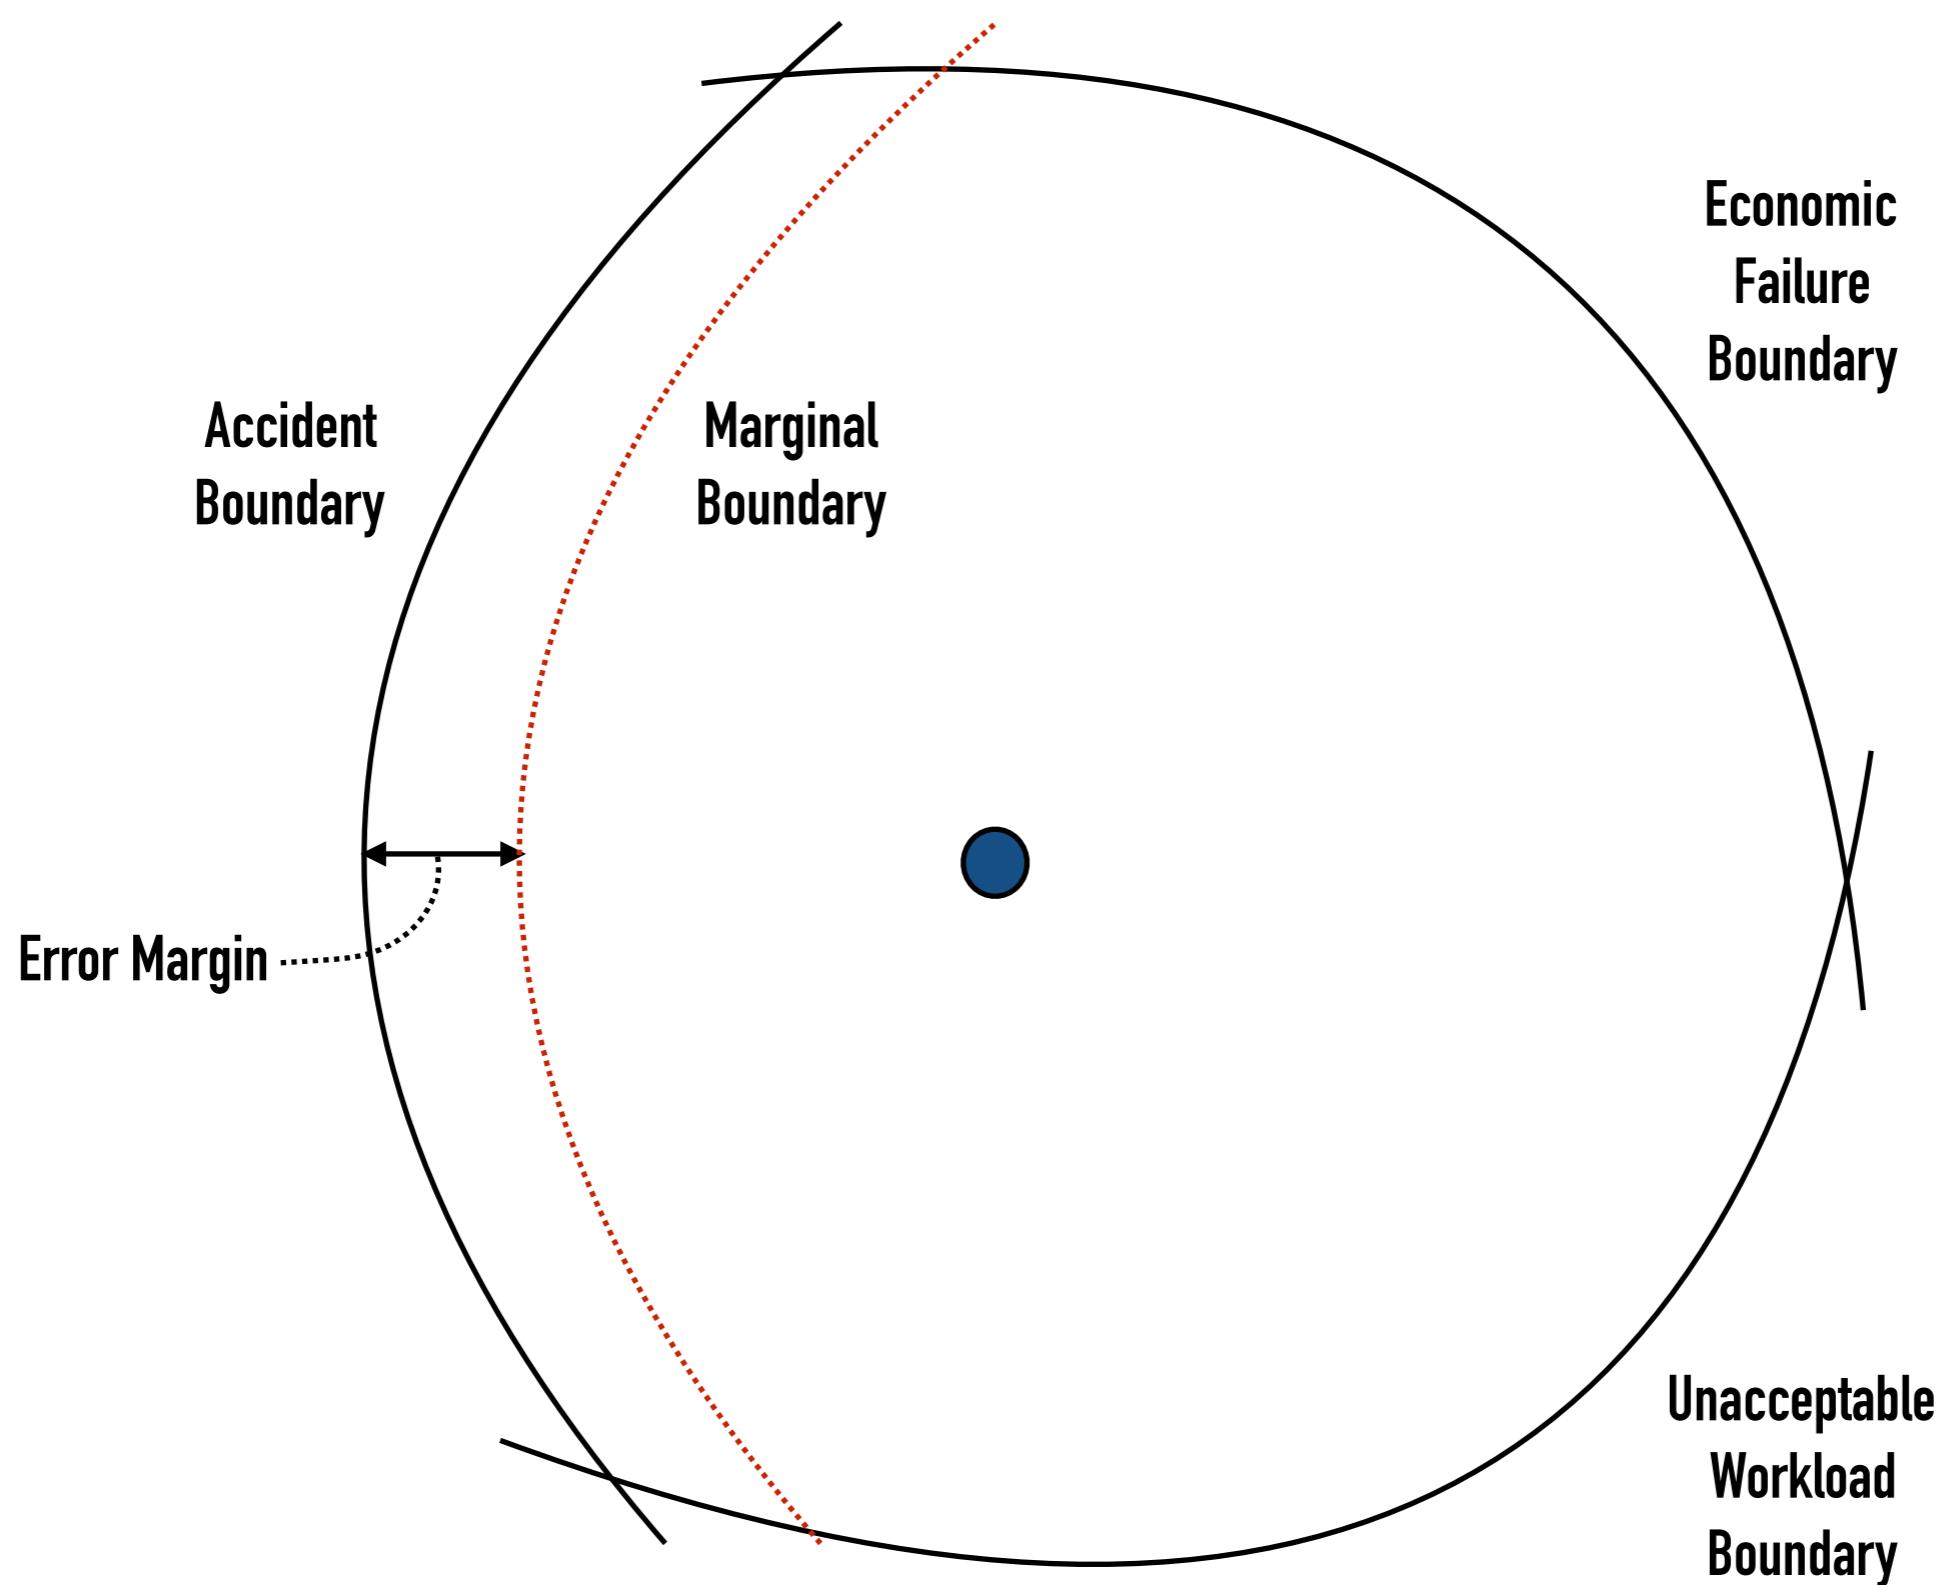
Operating at the Edge of Failure



"Going solid": a model of system dynamics and consequences for patient safety - R Cook, J Rasmussen

Resilience in complex adaptive systems: Operating at the Edge of Failure - Richard Cook - Talk at Velocity NY 2013

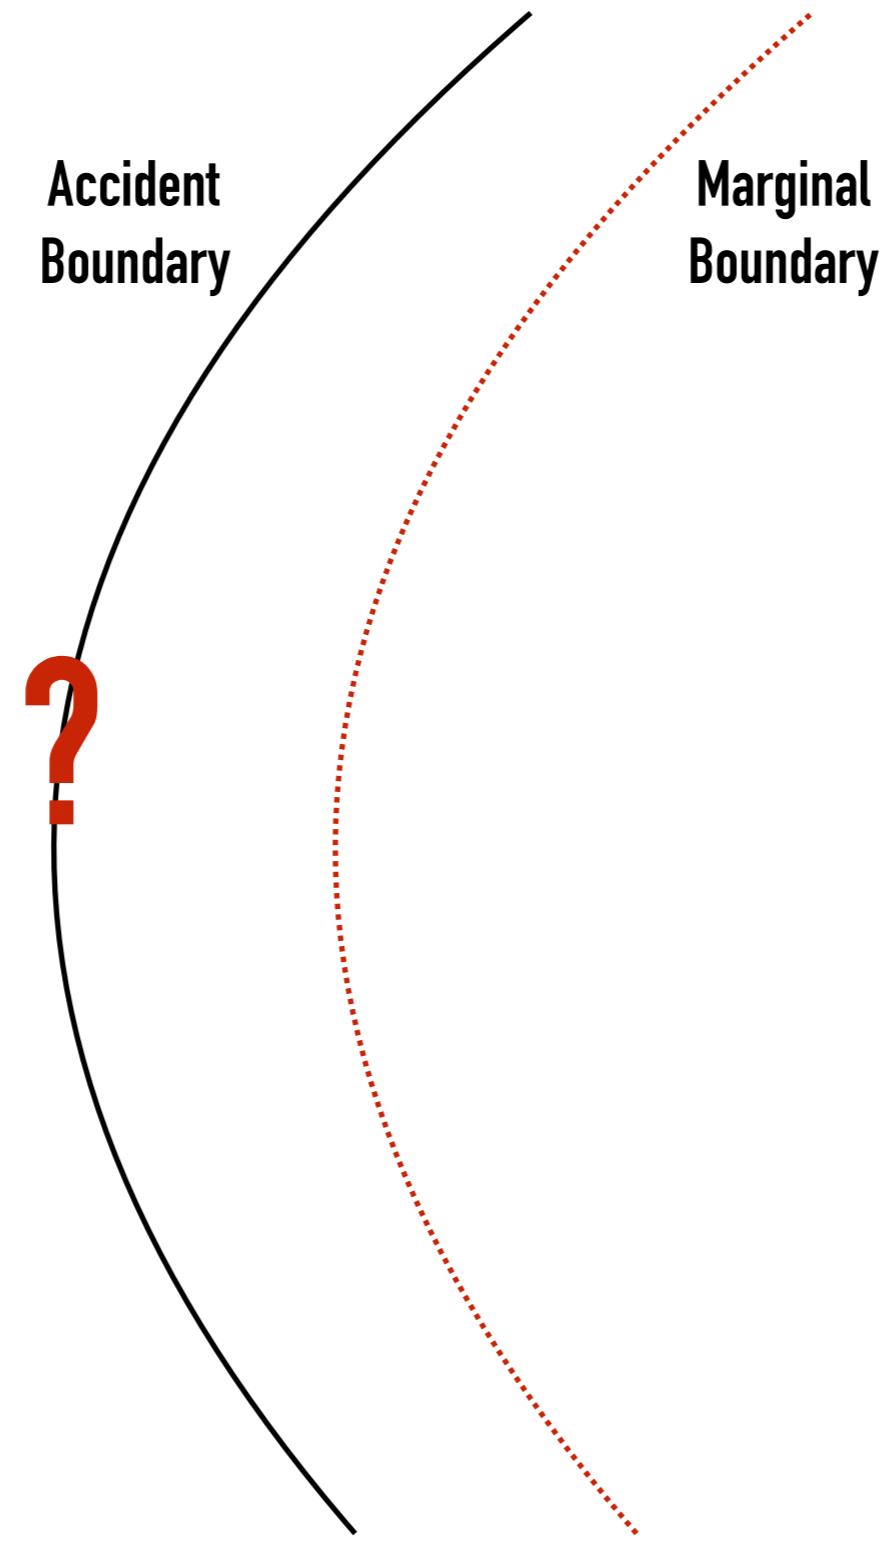
Operating at the Edge of Failure



"Going solid": a model of system dynamics and consequences for patient safety - R Cook, J Rasmussen

Resilience in complex adaptive systems: Operating at the Edge of Failure - Richard Cook - Talk at Velocity NY 2013

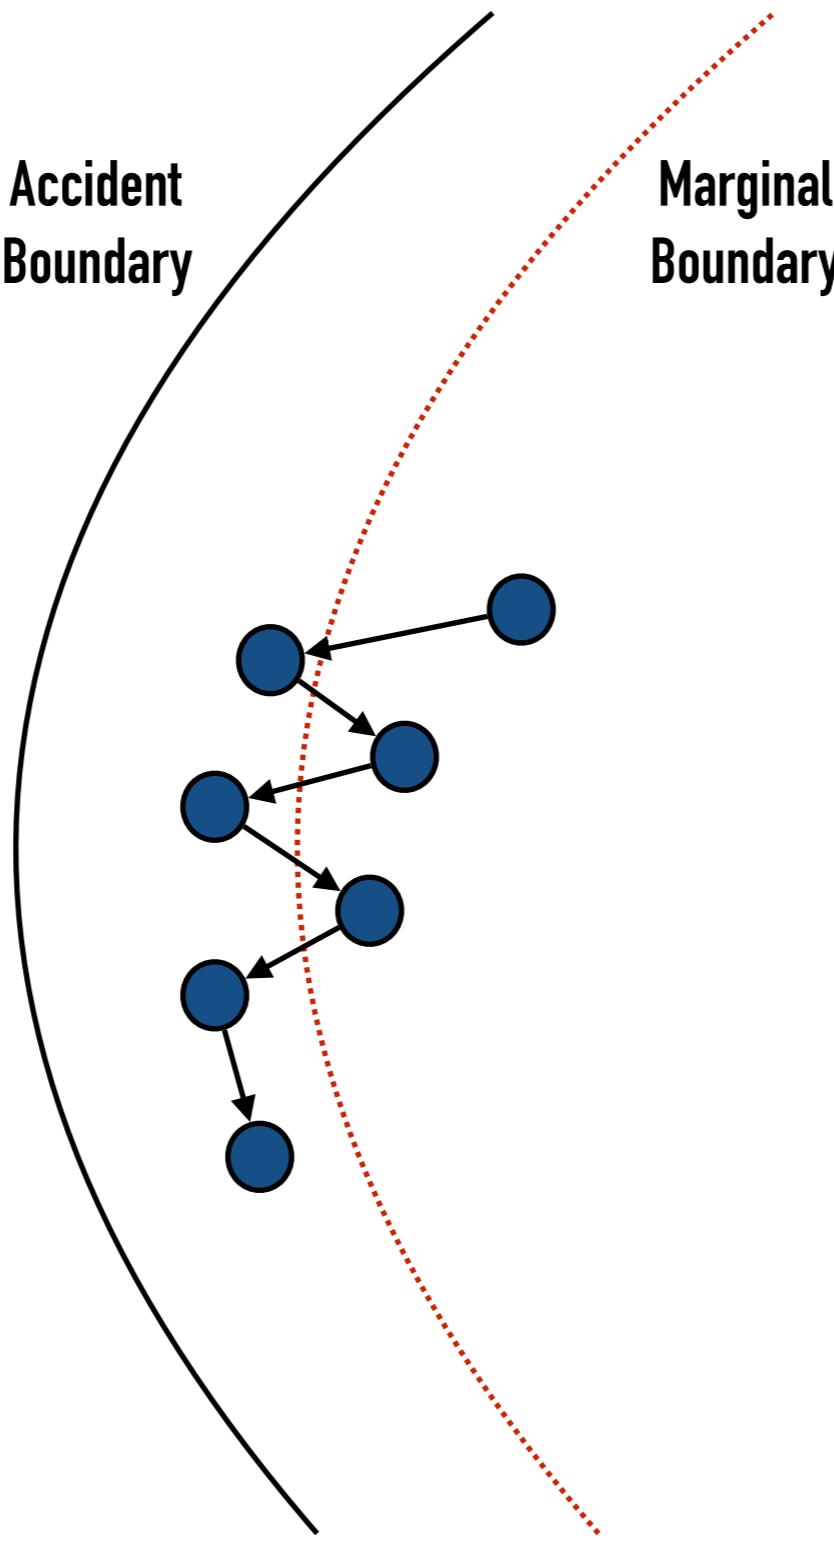
Operating at the Edge of Failure



"Going solid": a model of system dynamics and consequences for patient safety - R Cook, J Rasmussen

Resilience in complex adaptive systems: Operating at the Edge of Failure - Richard Cook - Talk at Velocity NY 2013

Operating at the Edge of Failure



"Going solid": a model of system dynamics and consequences for patient safety - R Cook, J Rasmussen

Resilience in complex adaptive systems: Operating at the Edge of Failure - Richard Cook - Talk at Velocity NY 2013

Embrace
Failure



Resilience
is by
Design

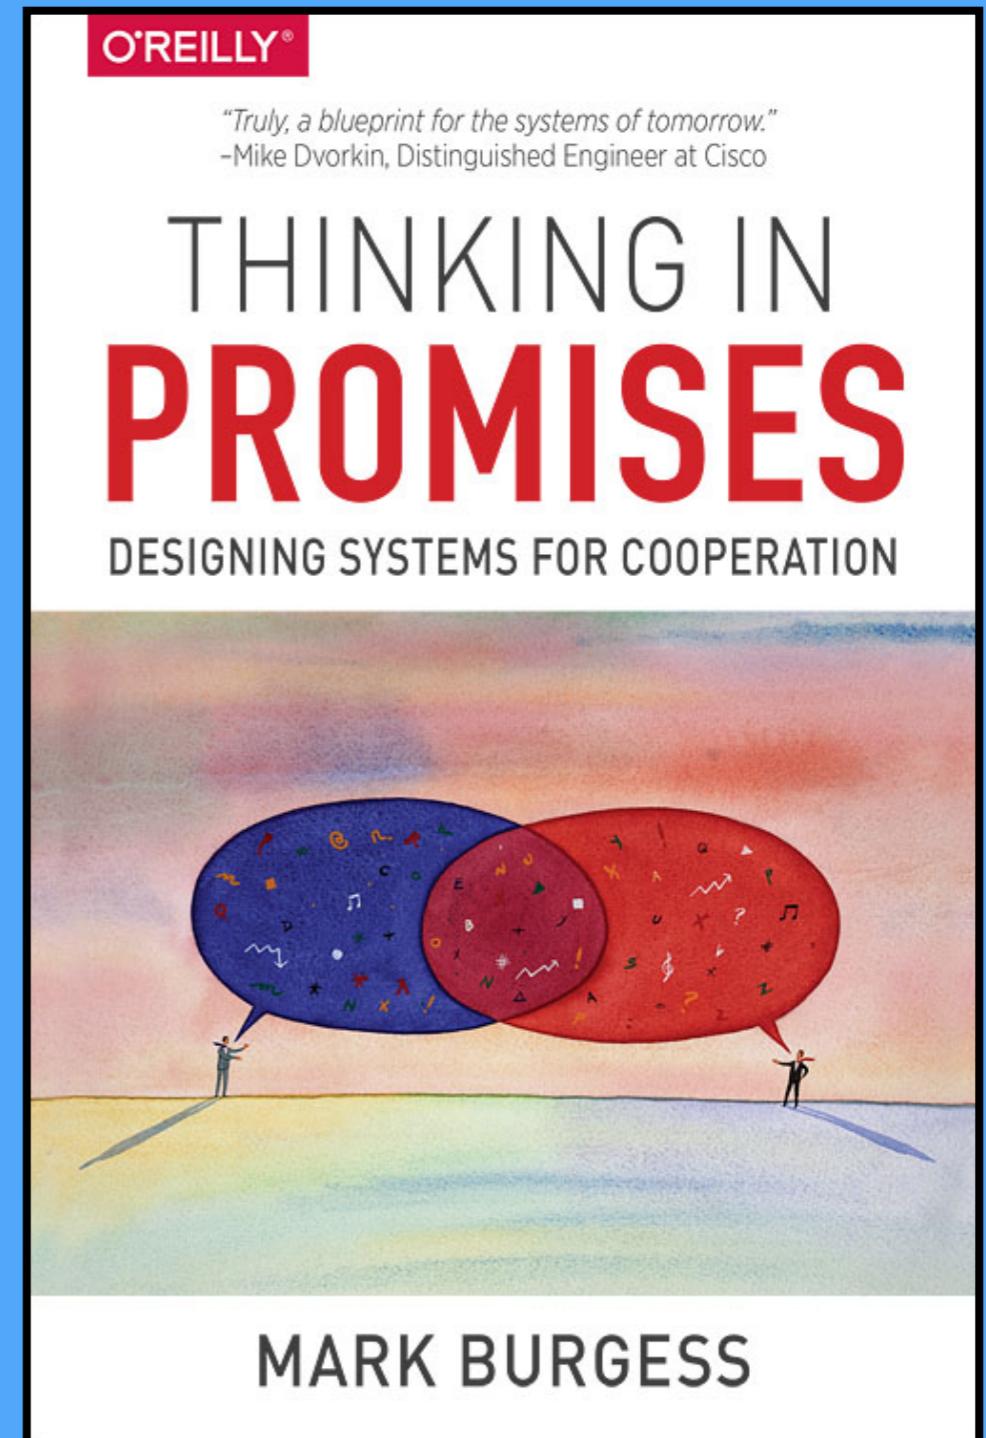
Photo courtesy of FEMA/Joselyne Augustino

**“Autonomy makes information local,
leading to greater certainty and stability.”**

- MARK BURGESS

Promise Theory

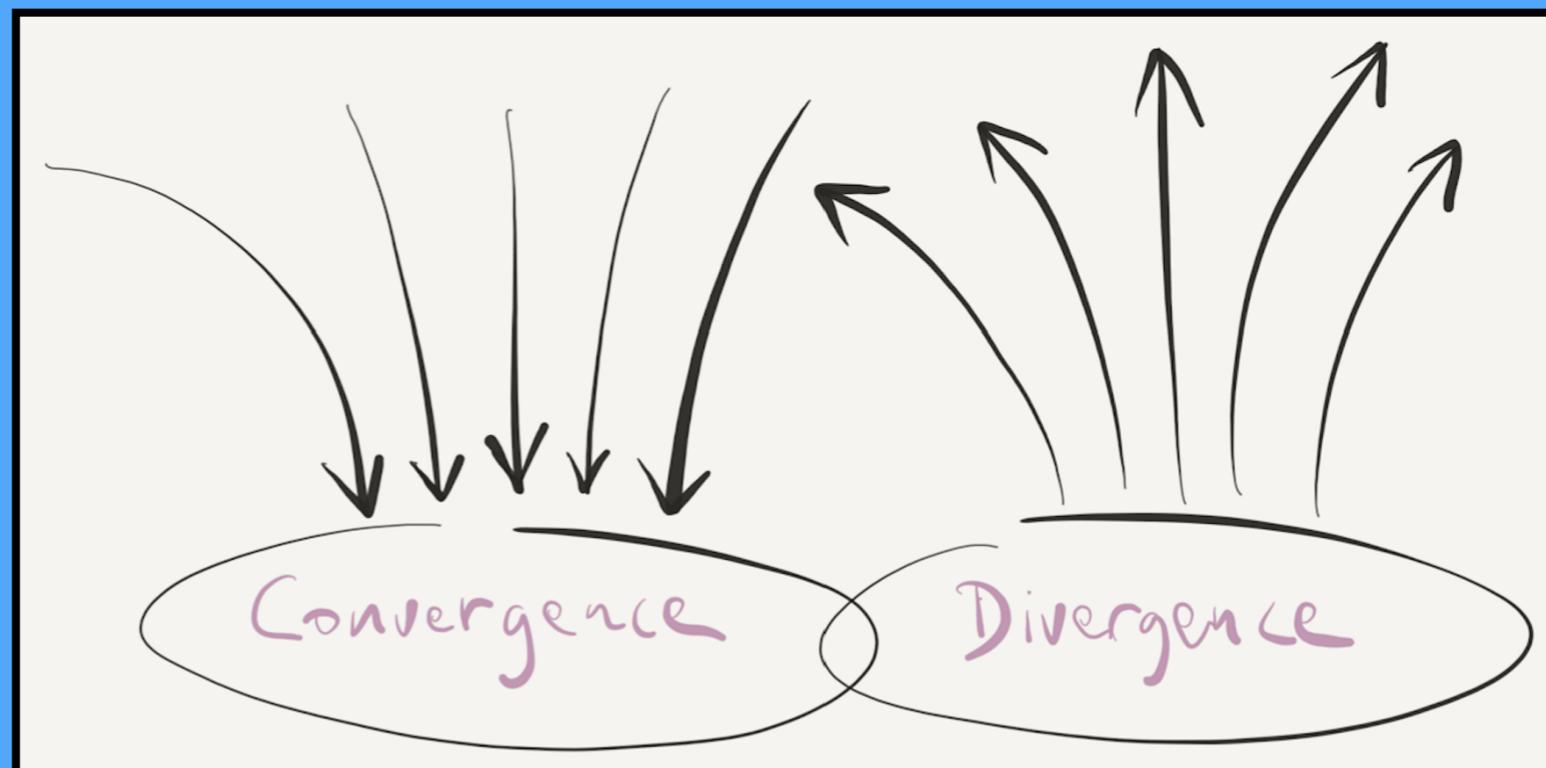
THINK IN
PROMISES
NOT
COMMANDS



Promise Theory

PROMISES CONVERGE TOWARDS
A DEFINITE OUTCOME FROM
UNPREDICTABLE BEGINNINGS
⇒ IMPROVED STABILITY

COMMANDS DIVERGE INTO
UNPREDICTABLE OUTCOMES FROM
DEFINITE BEGINNINGS
⇒ DECREASED STABILITY



Resilience in Biological Systems



Meerkats

Puppies! Now that I've got your attention, complexity theory - Nicolas Perony, TED talk

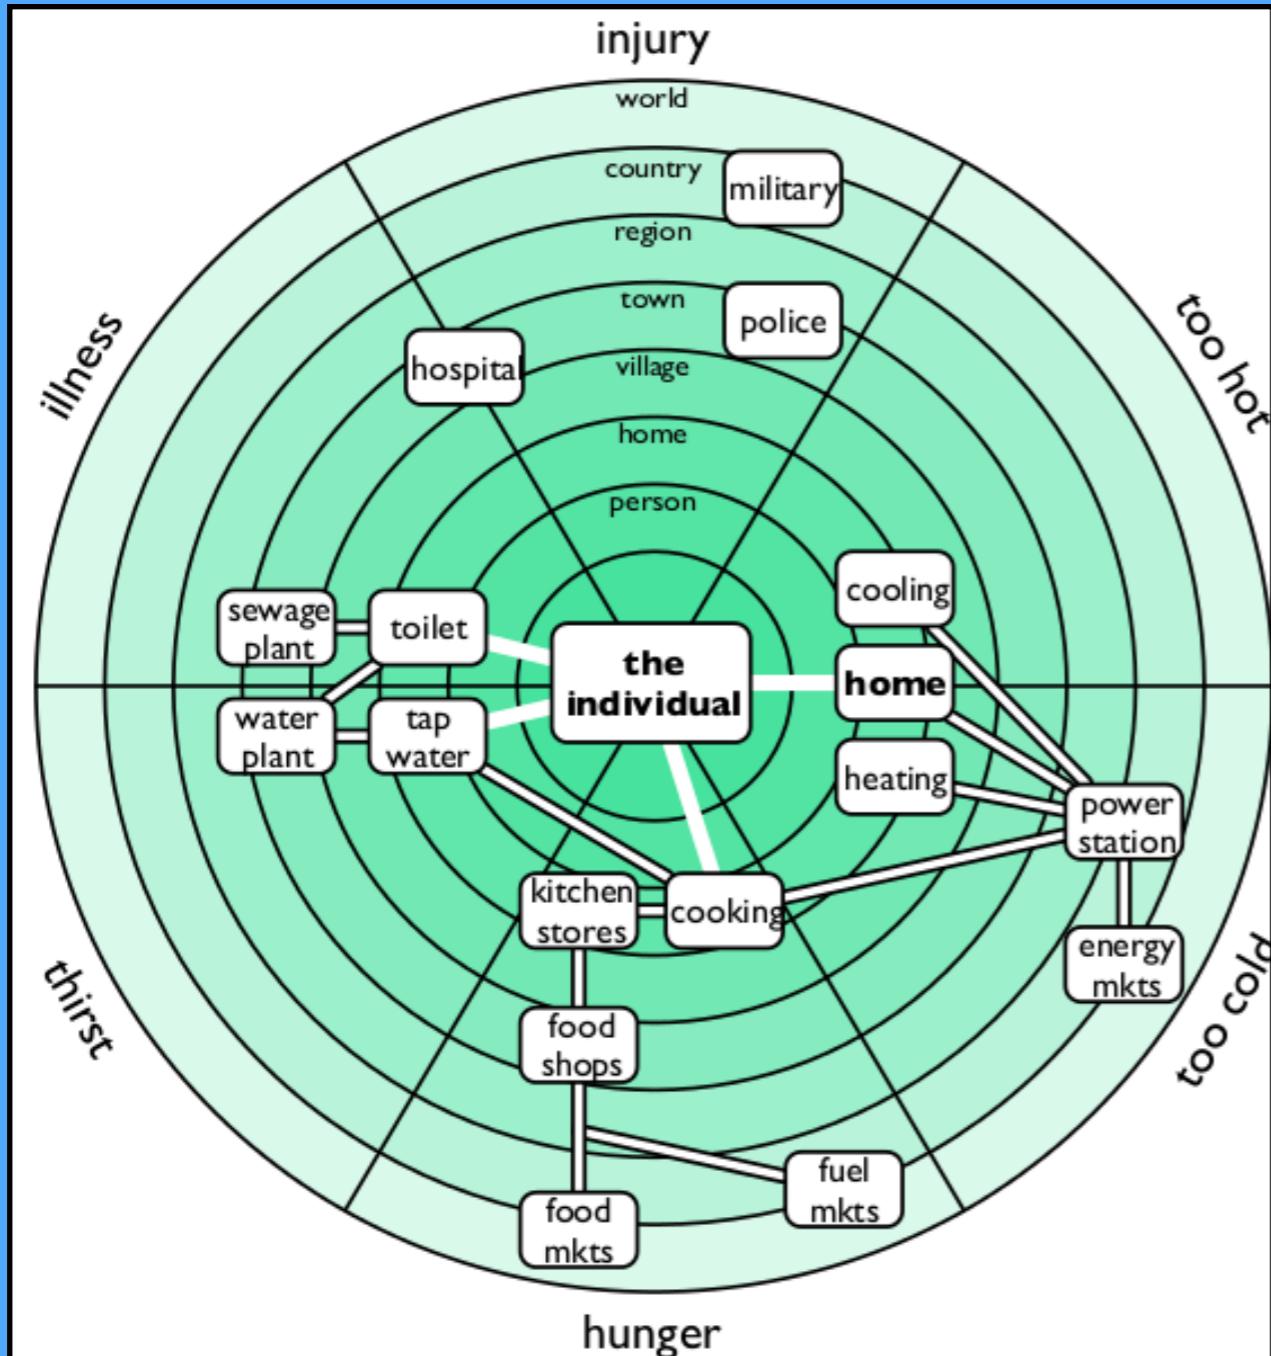
“In three words, in the animal kingdom,
simplicity leads to complexity
which leads to resilience.”

- NICOLAS PERONY

Resilience in Social Systems

Dealing in Security

UNDERSTANDING VITAL SERVICES, AND HOW THEY KEEP YOU SAFE



- 1 INDIVIDUAL (YOU)**
- 6 WAYS TO DIE**
- 3 SETS OF ESSENTIAL SERVICES**
- 7 LAYERS OF PROTECTION**

WHAT WE CAN LEARN FROM

Resilience in Biological and Social Systems

- 1. FEATURE DIVERSITY AND REDUNDANCY**
- 2. INTER-CONNECTED NETWORK STRUCTURE**
- 3. WIDE DISTRIBUTION ACROSS ALL SCALES**
- 4. CAPACITY TO SELF-ADAPT & SELF-ORGANIZE**

Applying resilience thinking: Seven principles for building resilience in social-ecological systems - Reinette Biggs et. al.

Toward Resilient Architectures 1: Biology Lessons - Michael Mehaffy, Nikos A. Salingaros

Resilience in Computer Systems

Our Disaster Recovery Plan Goes Something Like This...



DILBERT
By Scott Adams

We Need To
Manage
Failure

Not Try To Avoid It

Let It
Crash

Crash Only Software

STOP = CRASH SAFELY
START = RECOVER FAST

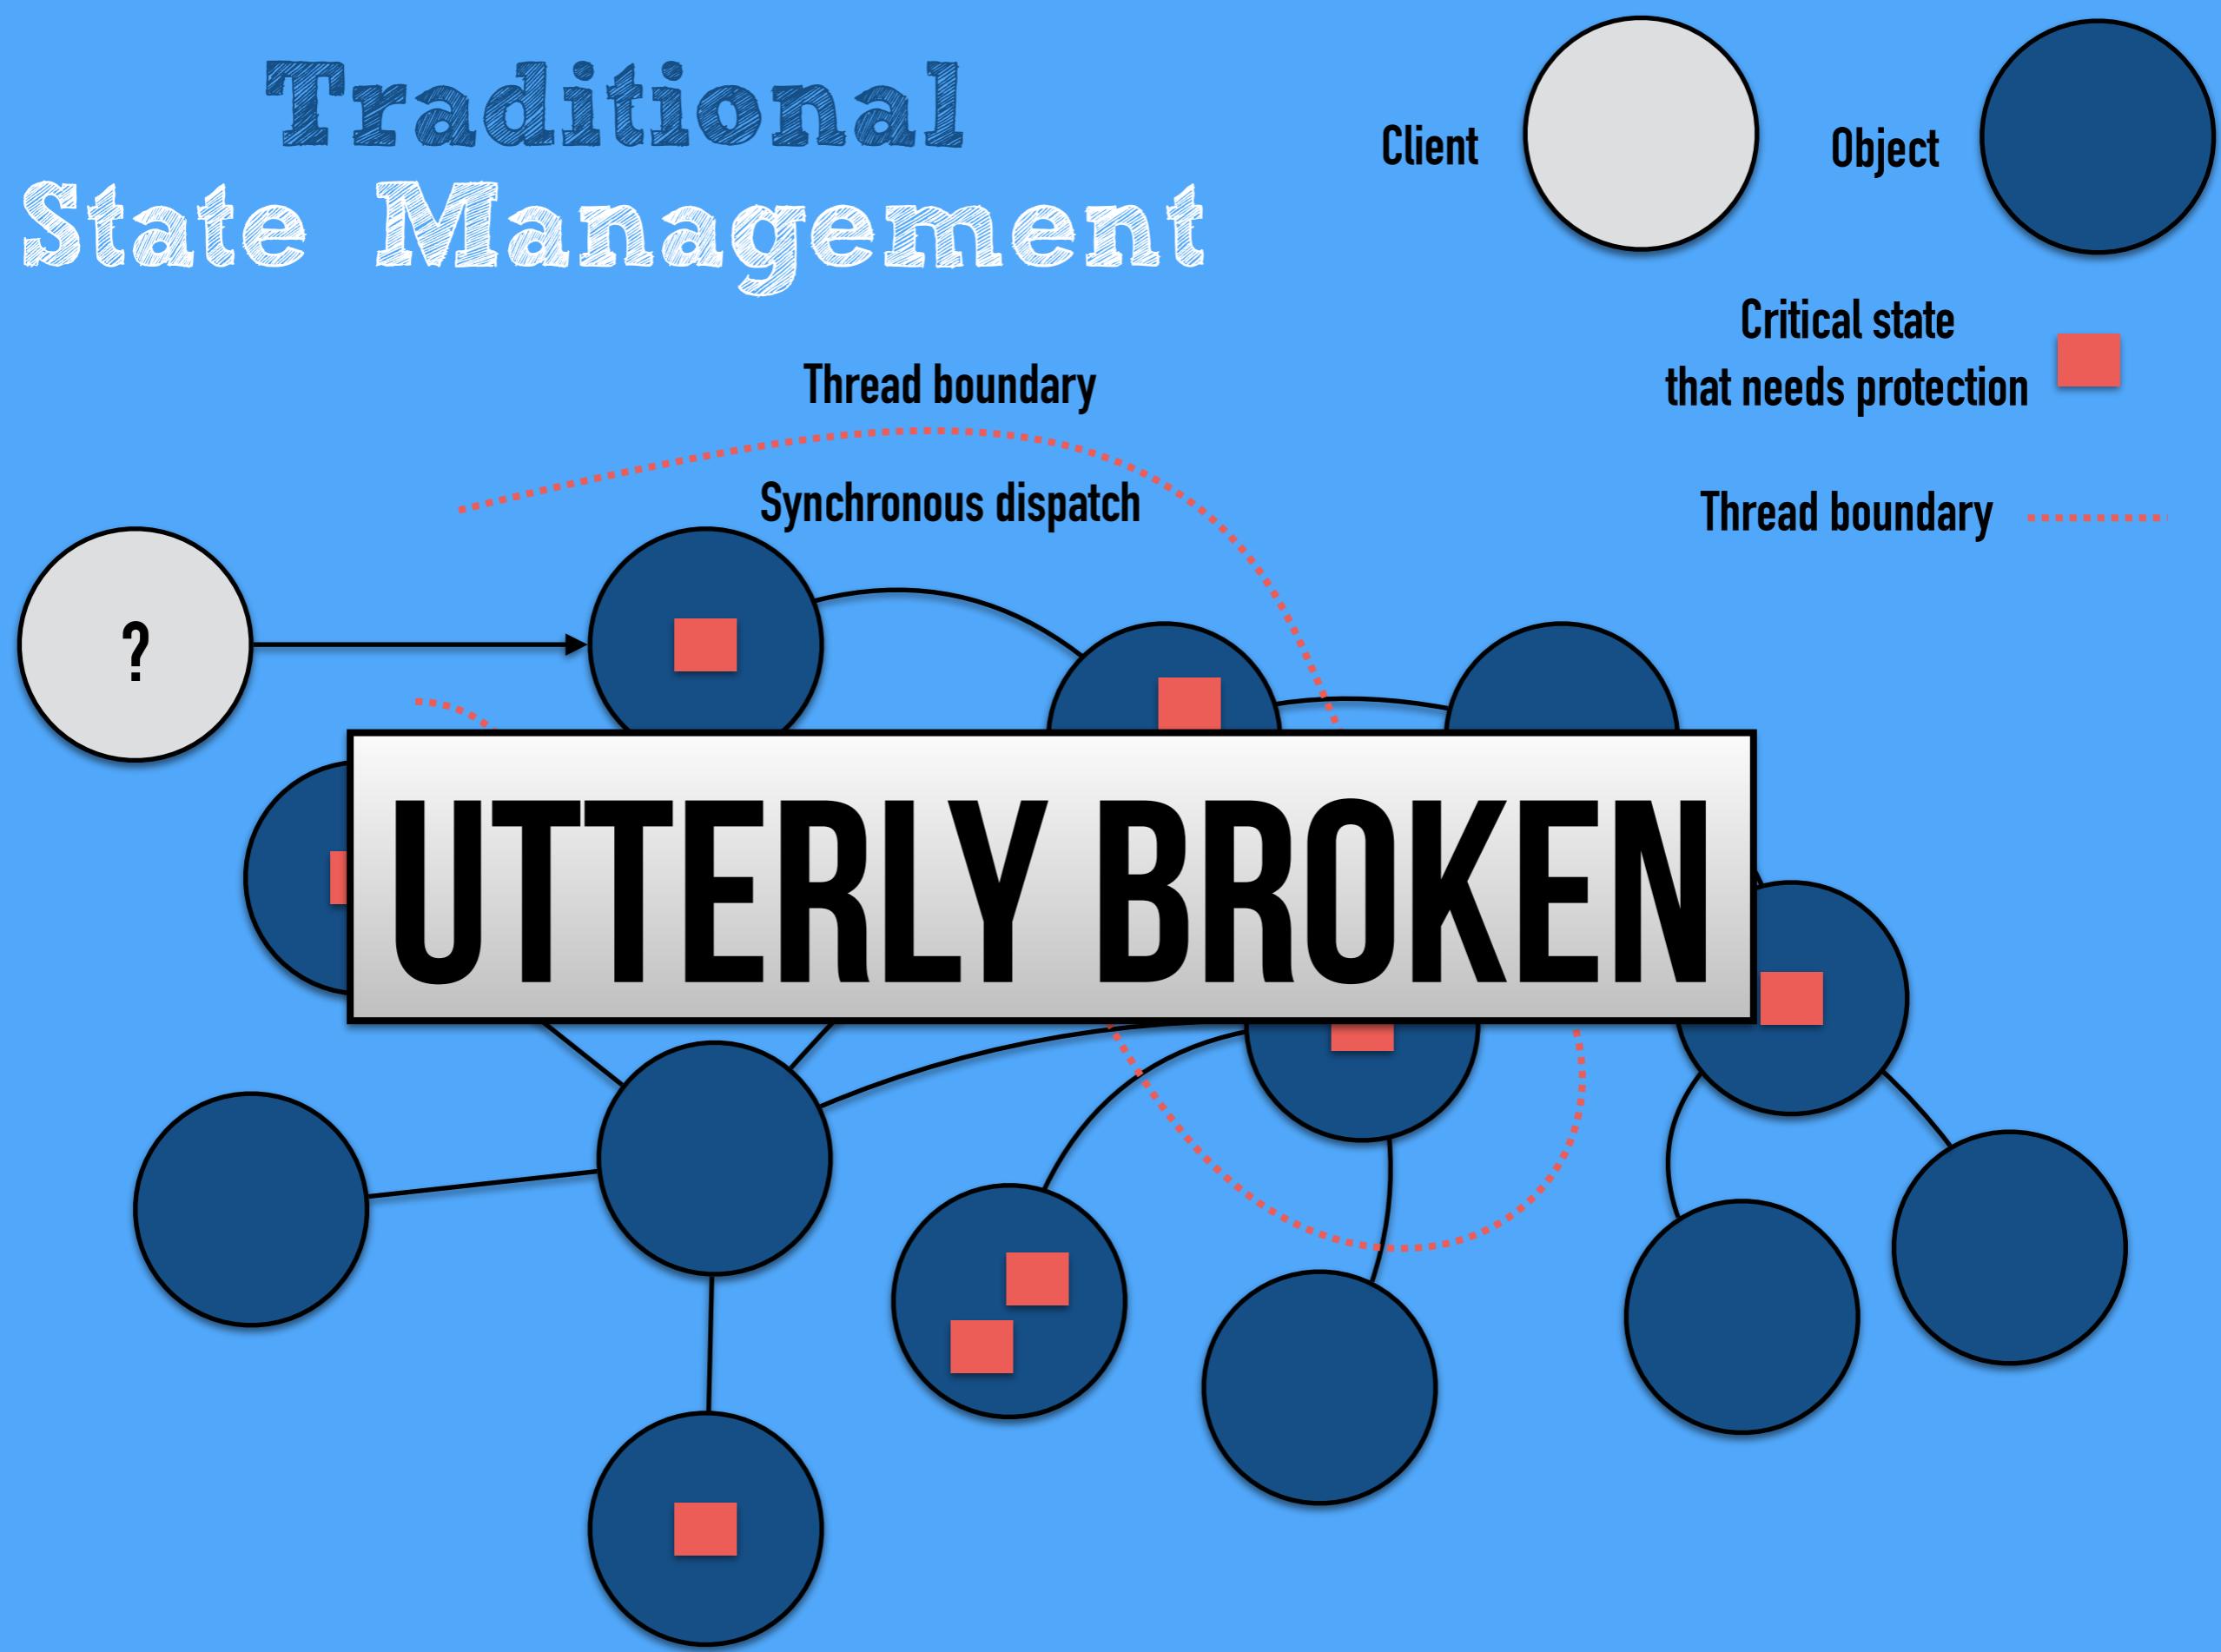
Recursive Restartability

TURNING THE CRASH-ONLY SLEDGEHAMMER INTO A SCALPEL



Recursive Restartability: Turning the Reboot Sledgehammer into a Scalpel - George Candea, Armando Fox

Traditional State Management



**“Accidents come from relationships
not broken parts.”**

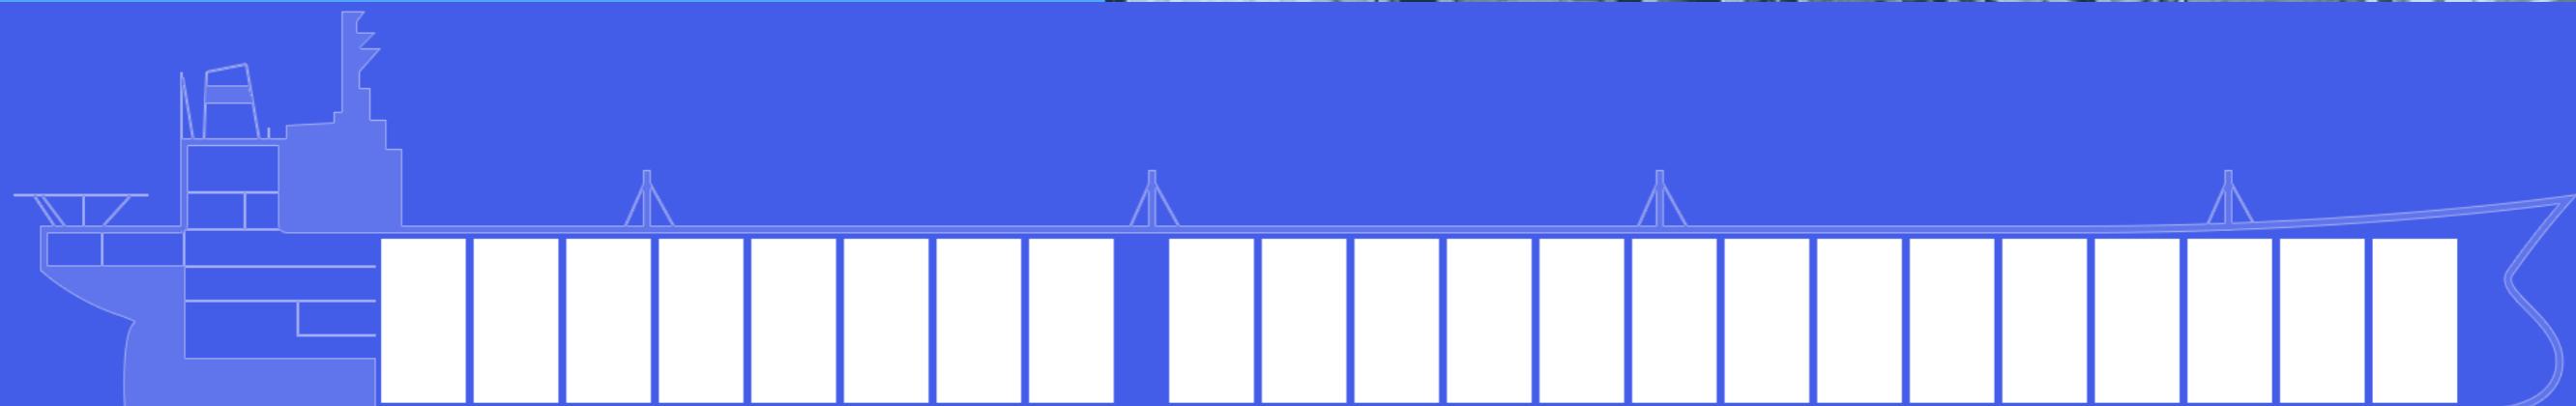
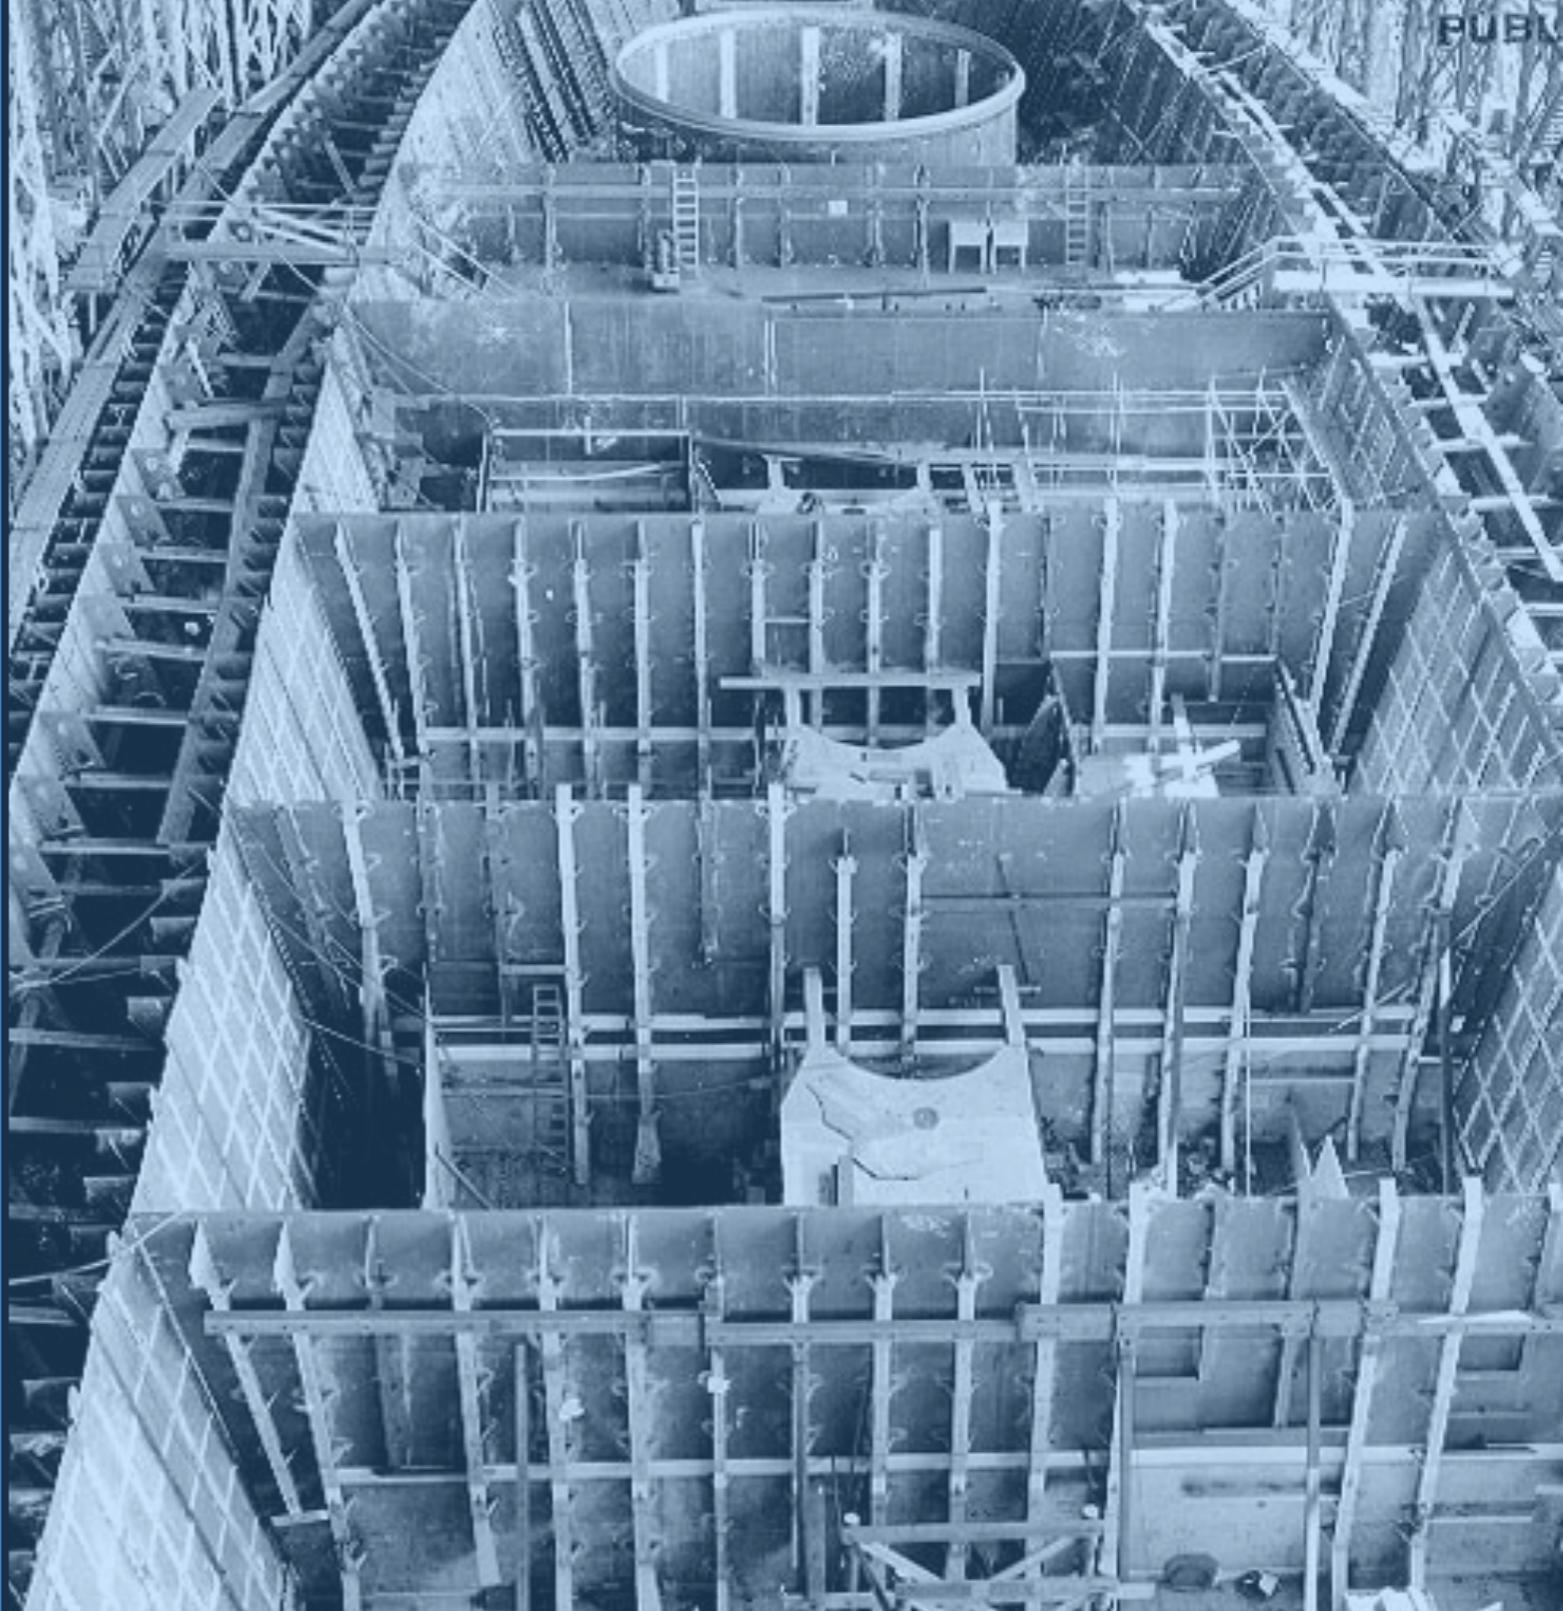
- SIDNEY DEKKER

REQUIREMENTS FOR A Sane Failure Model

FAILURES NEED TO BE

1. CONTAINED—AVOID CASCADING FAILURES
2. REIFIED—AS MESSAGES
3. SIGNALLED—ASYNCHRONOUSLY
4. OBSERVED—BY 1-N
5. MANAGED—OUTSIDE FAILED CONTEXT

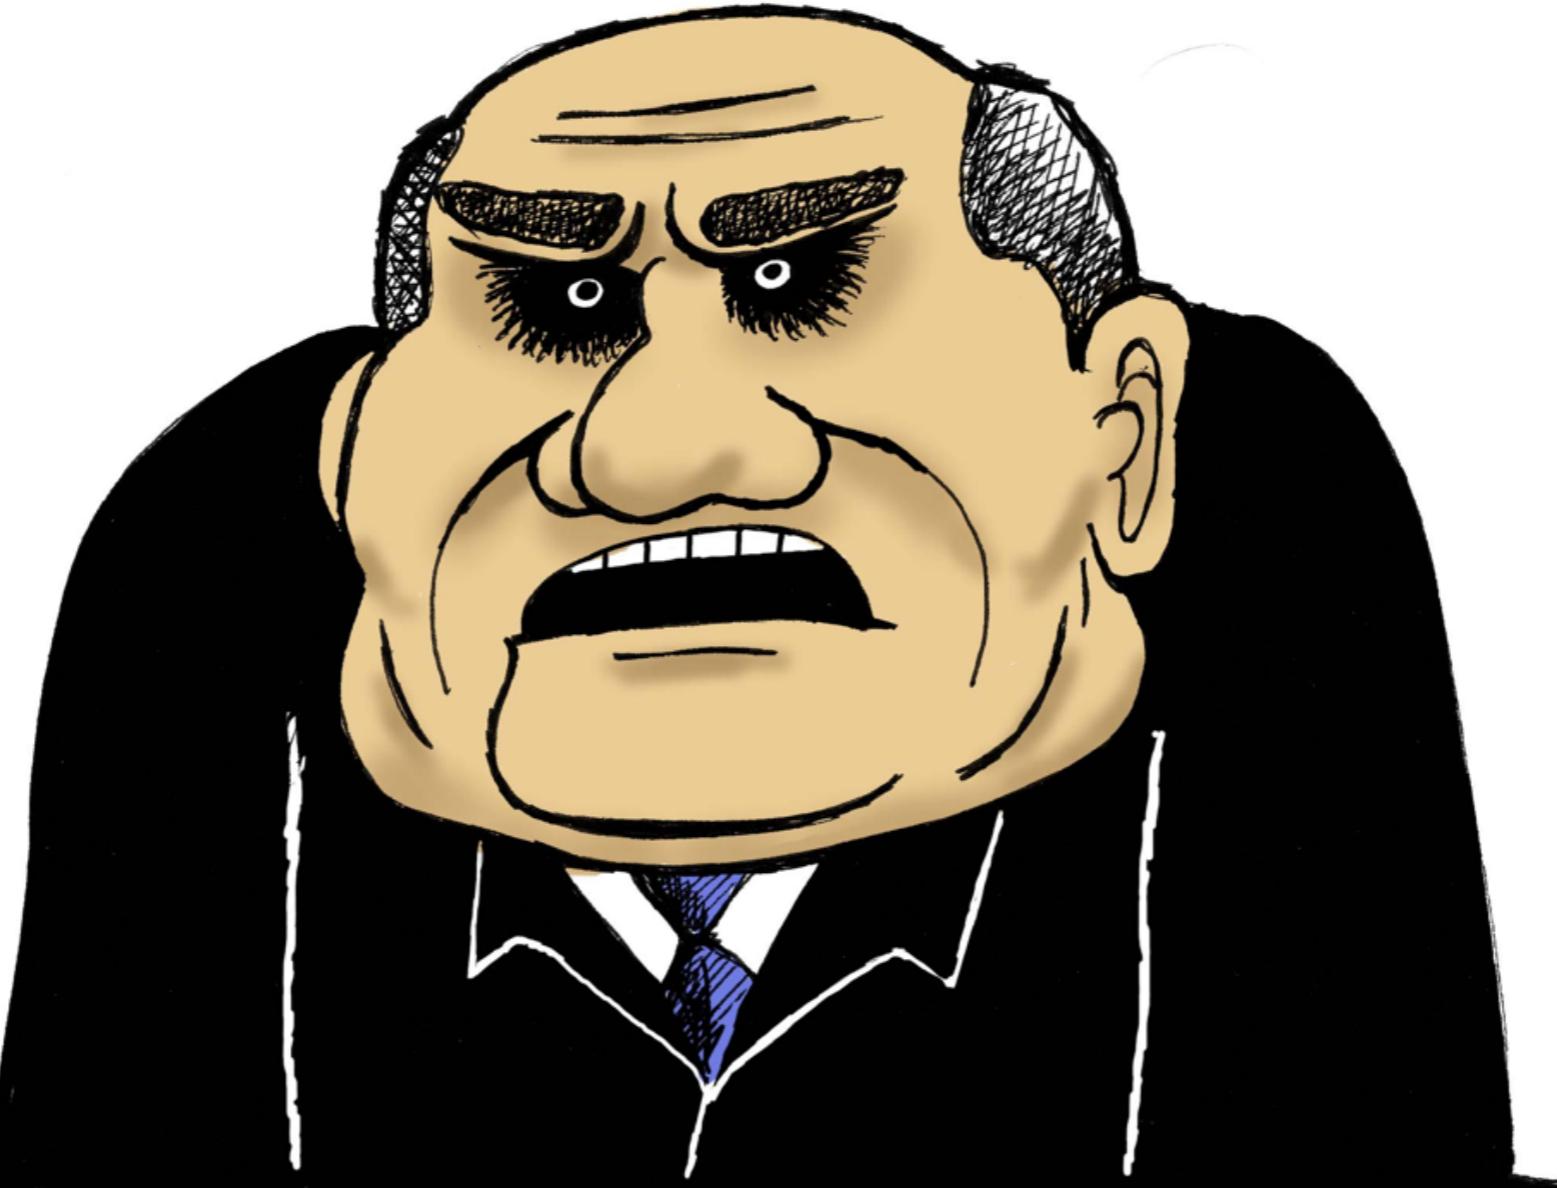
Bulkhead Pattern



Enter Supervision

Adrian S. Bruce ©2004 www.artandtechnology.com.au

BRUCE 12/04



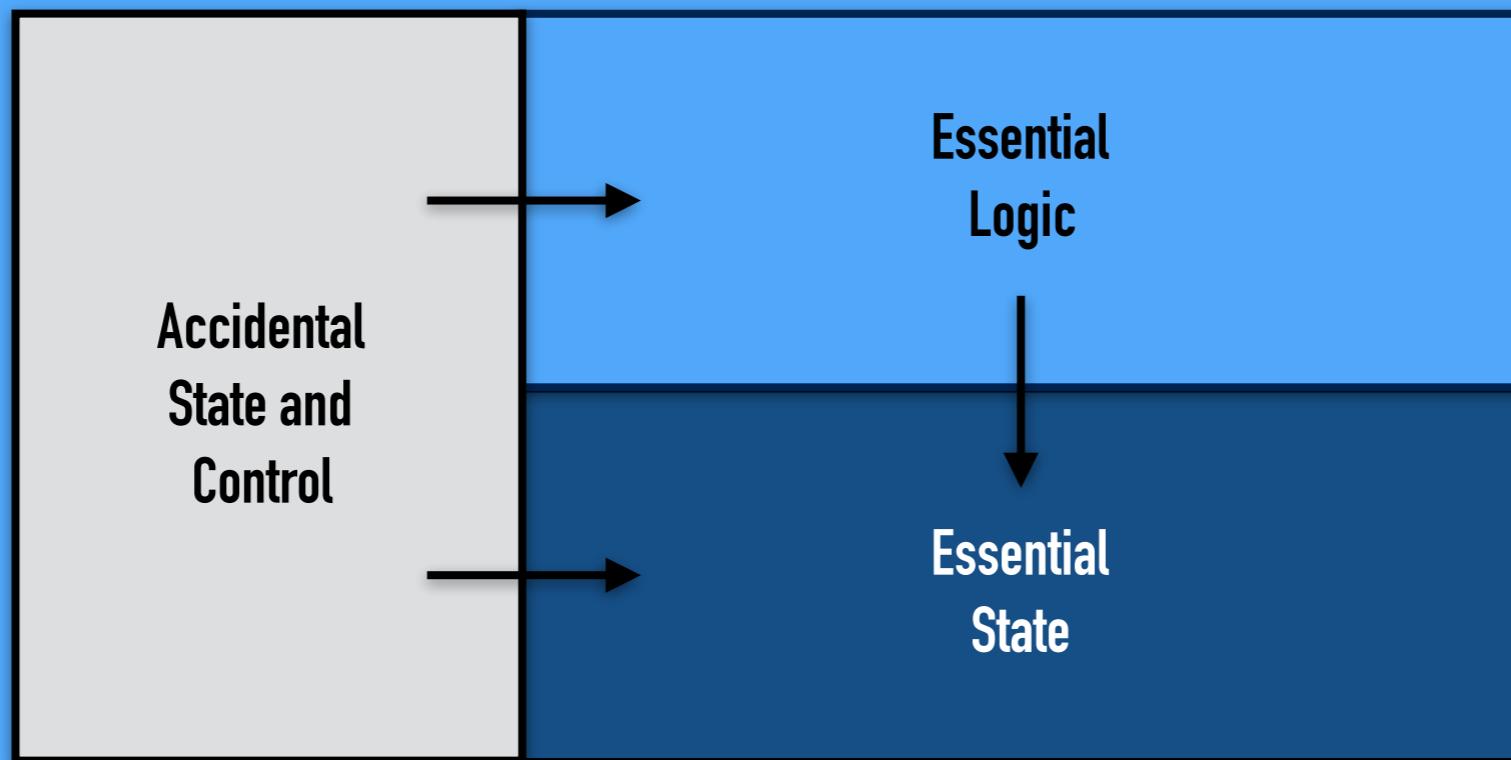
**THE BEATINGS WILL CONTINUE
UNTIL MORALE IMPROVES**

WE NEED A WAY OUT OF THE State Tar Pit

- INPUT DATA
- DERIVED DATA

Critical

WE NEED A WAY OUT OF THE State Tar Pit



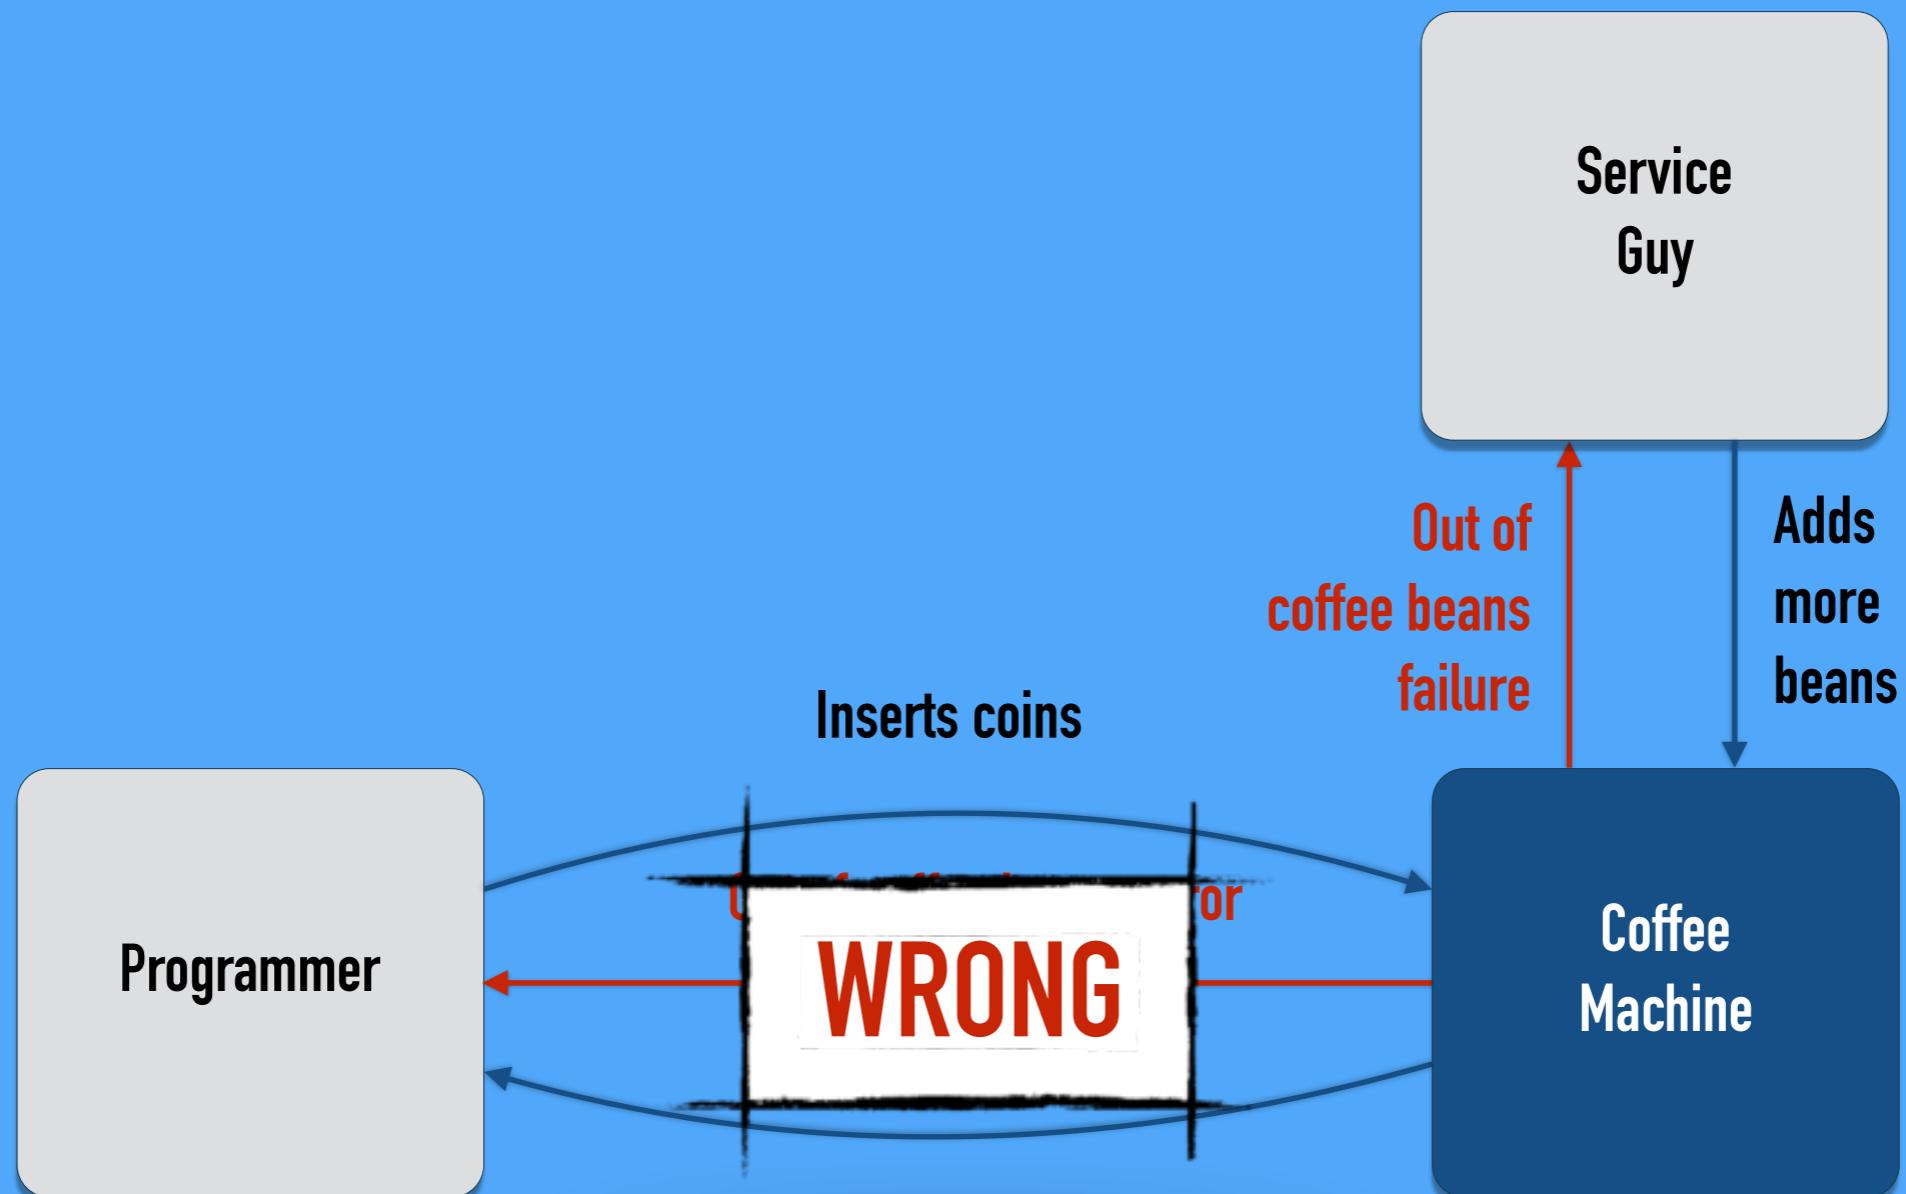


THE Vending Machine Pattern

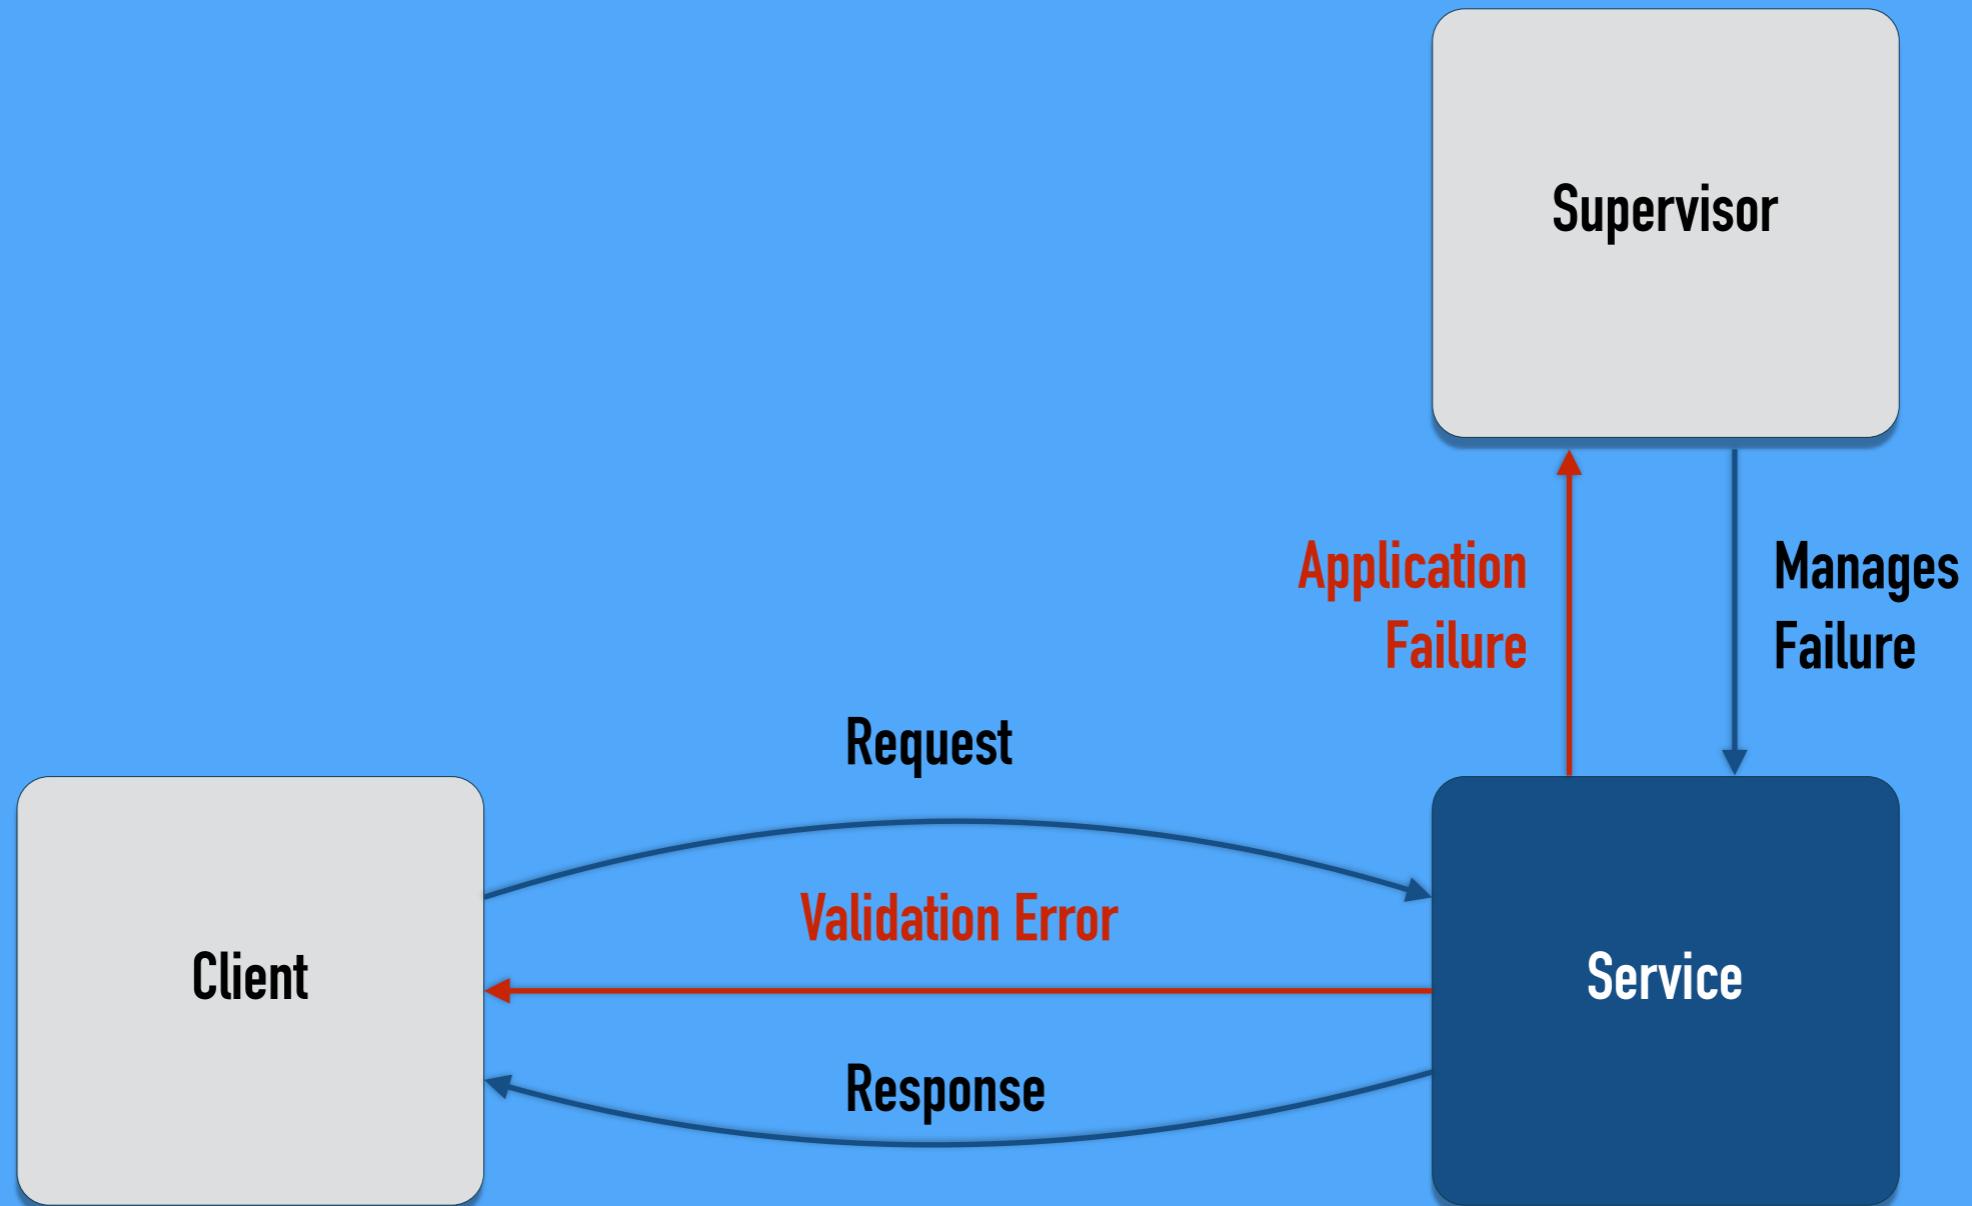
Think Vending Machine



Think Vending Machine



Think Vending Machine





ERROR kernel Pattern

ONION-LAYERED STATE & FAILURE MANAGEMENT

Making reliable distributed systems in the presence of software errors - Joe Armstrong
On Erlang, State and Crashes - Jesper Louis Andersen

Onion Layered State Management

Client

Object

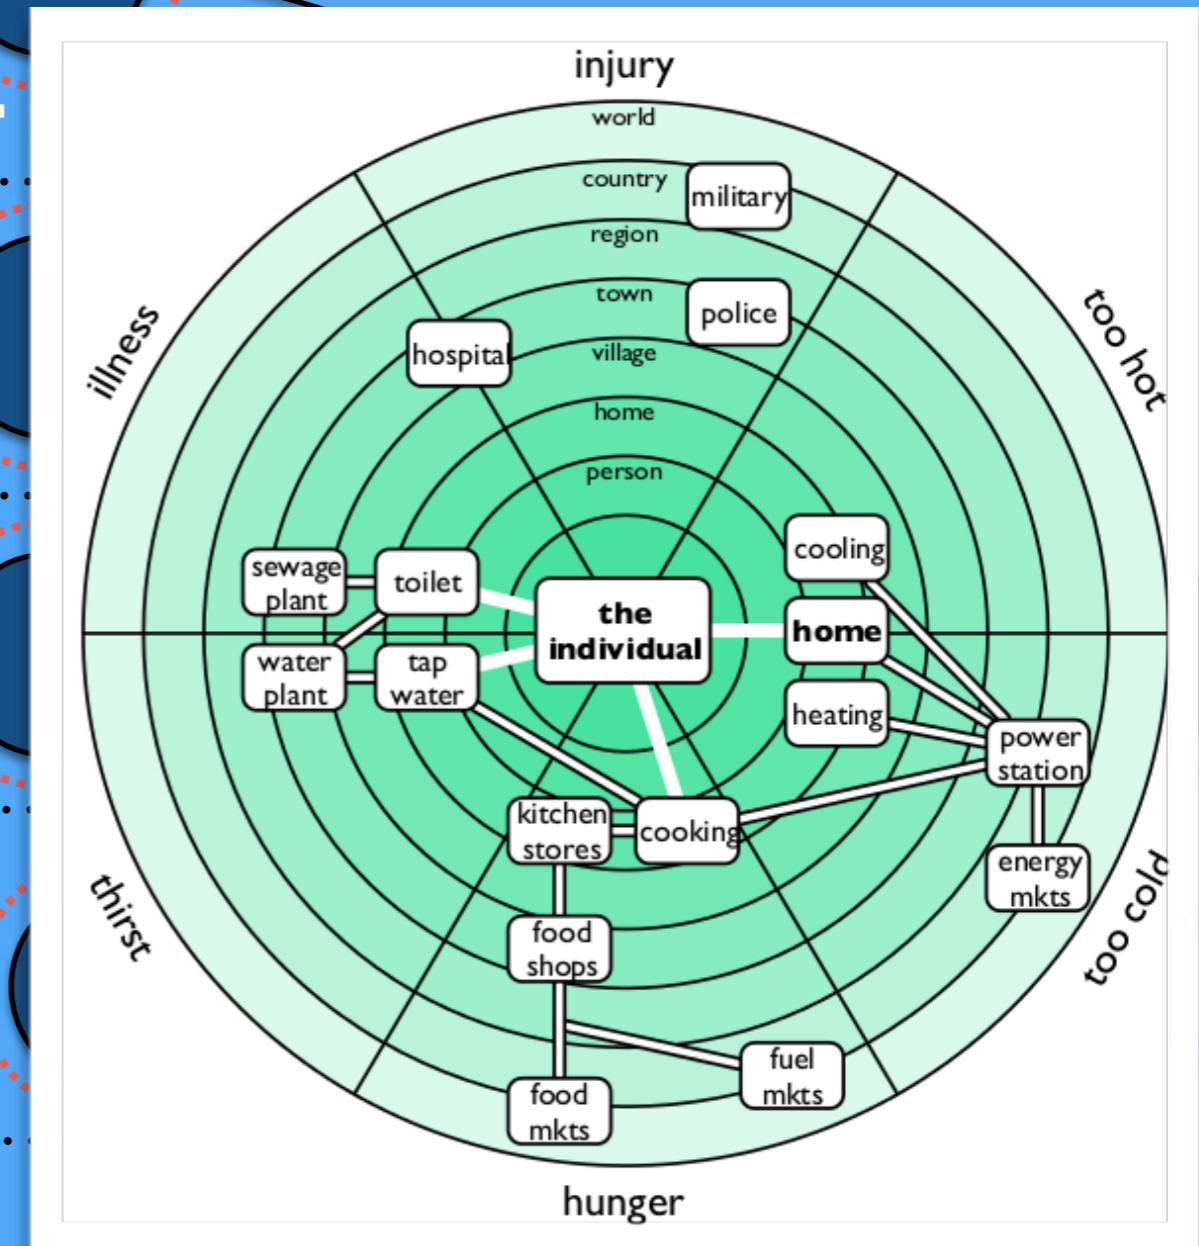
Supervision



Error Kernel

Critical state
that needs protection

Thread boundary



Demo Time

LET'S MODEL A RESILIENT VENDING MACHINE, IN AKKA

Demo Runner

```
object VendingMachineDemo extends App {  
  
    val system = ActorSystem("vendingMachineDemo")  
    val coffeeMachine = system.actorOf(Props[CoffeeMachineManager], "coffeeMachineManager")  
    val customer = Inbox.create(system) // emulates the customer  
  
    ... // test runs  
  
    system.shutdown()  
}
```

Test Happy Path

```
// Insert 2 coins and get an Espresso
customer.send(coffeeMachine, Coins(2))
customer.send(coffeeMachine, Selection(Espresso))
val Beverage(coffee1) = customer.receive(5.seconds)
println(s"Got myself an $coffee1")
assert(coffee1 == Espresso)
```

Test User Error

```
customer.send(coffeeMachine, Coins(1))
customer.send(coffeeMachine, Selection(Latte))
val NotEnoughCoinsError(message) = customer.receive(5.seconds)
println(s"Got myself a validation error: $message")
assert(message == "Please insert [1] coins")
```

Test System Failure

```
// Insert 1 coin (had 1 before) and try to get my Latte
// Machine should:
//   1. Fail
//   2. Restart
//   3. Resubmit my order
//   4. Give me my coffee
customer.send(coffeeMachine, Coins(1))
customer.send(coffeeMachine, TriggerOutOfCoffeeBeansFailure)
customer.send(coffeeMachine, Selection(Latte))
val Beverage(coffee2) = customer.receive(5.seconds)
println(s"Got myself a $coffee2")
assert(coffee2 == Latte)
```

Protocol

```
// Coffee types
trait CoffeeType
case object BlackCoffee extends CoffeeType
case object Latte extends CoffeeType
case object Espresso extends CoffeeType

// Commands
case class Coins(number: Int)
case class Selection(coffee: CoffeeType)
case object TriggerOutOfCoffeeBeansFailure

// Events
case class CoinsReceived(number: Int)

// Replies
case class Beverage(coffee: CoffeeType)

// Errors
case class NotEnoughCoinsError(message: String)

// Failures
case class OutOfCoffeeBeansFailure(customer: ActorRef,
                                     pendingOrder: Selection,
                                     nrOfInsertedCoins: Int) extends Exception
```

CoffeeMachine

```
class CoffeeMachine extends Actor {  
    val price = 2  
    var nrOfInsertedCoins = 0  
    var outOfCoffeeBeans = false  
    var totalNrOfCoins = 0  
  
    def receive = { ... }  
  
    override def postRestart(failure: Throwable): Unit = { ... }  
}
```

CoffeeMachine

```
def receive = {
    case Coins(nr) =>
        nrOfInsertedCoins += nr
        totalNrOfCoins += nr
        println(s"Inserted [$nr] coins")
        println(s"Total number of coins in machine is [$totalNrOfCoins)")

    case selection @ Selection(coffeeType) =>
        if (nrOfInsertedCoins < price)
            sender.tell(NotEnoughCoinsError(
                s"Insert [{price - nrOfInsertedCoins}] coins"), self)
        else {
            if (outOfCoffeeBeans)
                throw new OutOfCoffeeBeansFailure(sender, selection, nrOfInsertedCoins)
            println(s"Brewing your $coffeeType")
            sender.tell(Beverage(coffeeType), self)
            nrOfInsertedCoins = 0
        }

    case TriggerOutOfCoffeeBeansFailure =>
        outOfCoffeeBeans = true
}
```

CoffeeMachine

```
override def postRestart(failure: Throwable): Unit = {
  println(s"Restarting coffee machine...")
  failure match {
    case OutOfCoffeeBeansFailure(customer, pendingOrder, coins) =>
      nrOfInsertedCoins = coins
      outOfCoffeeBeans = false
      println(s"Resubmitting pending order $pendingOrder")
      context.self.tell(pendingOrder, customer)
  }
}
```

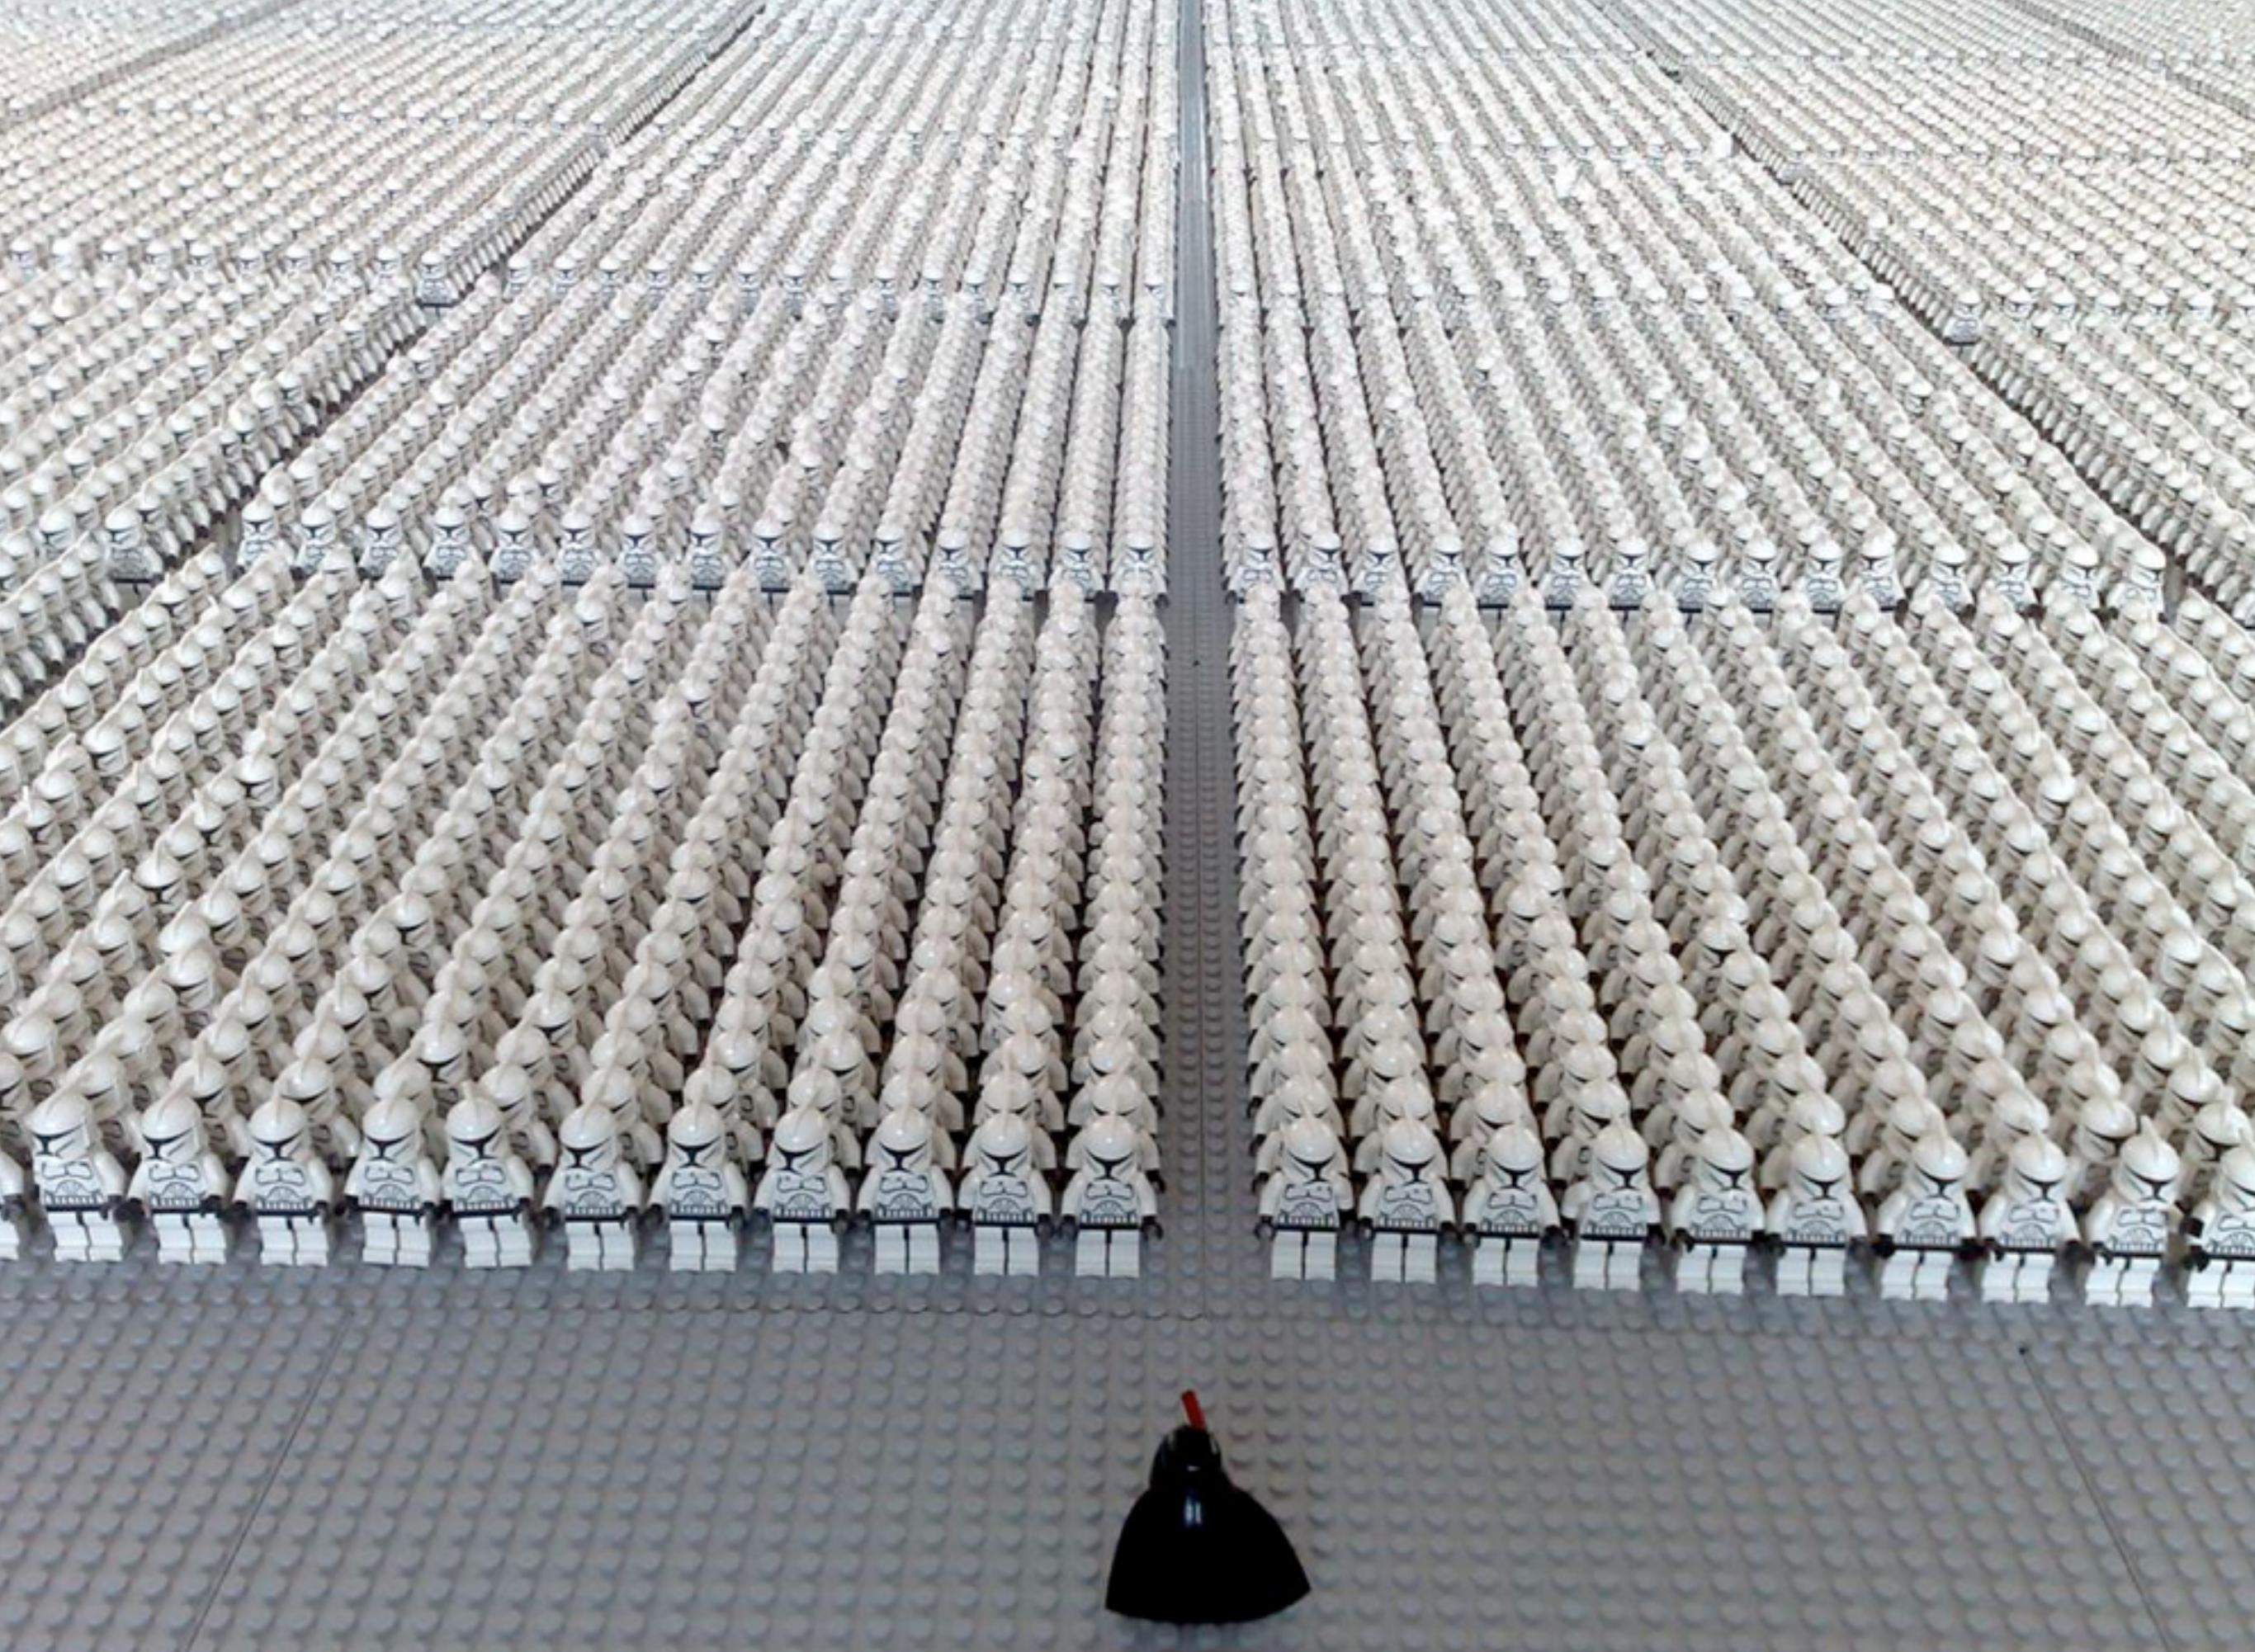
Supervisor

```
class CoffeeMachineManager extends Actor {  
    override val supervisorStrategy =  
        OneForOneStrategy(maxNrOfRetries = 10, withinTimeRange = 1.minute) {  
            case e: OutOfCoffeeBeansFailure =>  
                println(s"ServiceGuy notified: $e")  
                Restart  
            case _: Exception =>  
                Escalate  
        }  
  
    // to simplify things he is only managing 1 single machine  
    val machine = context.actorOf(  
        Props[coffeeMachine], name = "coffeeMachine")  
  
    def receive = {  
        case request => machine.forward(request)  
    }  
}
```

**SO..... ARE WE DONE?
SORRY... BUT NOT REALLY.**

We can not
keep putting
all eggs in the
same basket





A man with dark, curly hair and a mustache, wearing sunglasses and a light-colored suit, is smiling and giving a thumbs-up gesture. He is standing outdoors with a large, spiky cactus visible on the right side of the frame. The background shows a dry, arid landscape under a clear sky.

The Network
is Reliable
NOT
Really

HERE, WE ARE LIVING IN THE
Looming
Shadow of
Impossibility
Theorems

FLP: CONSENSUS IS IMPOSSIBLE
CAP: CONSISTENCY IS IMPOSSIBLE



Towards Resilient Distributed Systems

ISOLATION

- Autonomous Microservices
- Resilient Protocols
- Virtualization

DATA RESILIENCE

- Eventual & Causal Consistency
- Event Logging
- Flow Control / Feedback Control

EMBRACE THE NETWORK

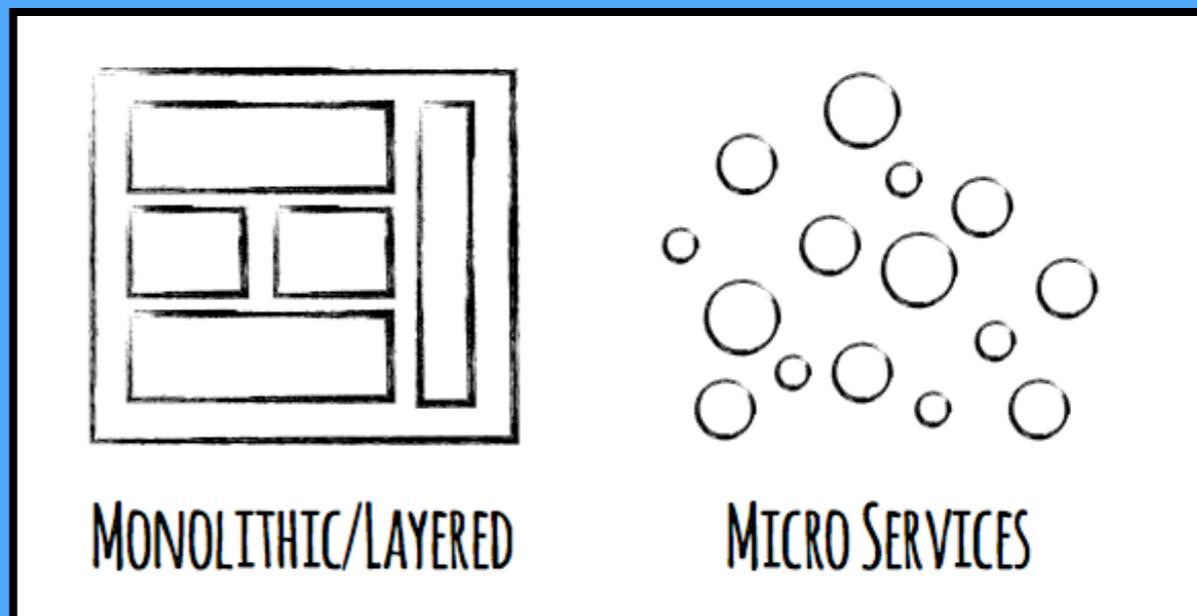
- Asynchronicity
- Location Transparency

SELF-HEALING

- Decentralized Architectures
- Gossip Protocols
- Failure Detection

Microservices

- 1. AUTONOMY**
- 2. ISOLATION**
- 3. MOBILITY**
- 4. SINGLE RESPONSIBILITY**
- 5. EXCLUSIVE STATE**

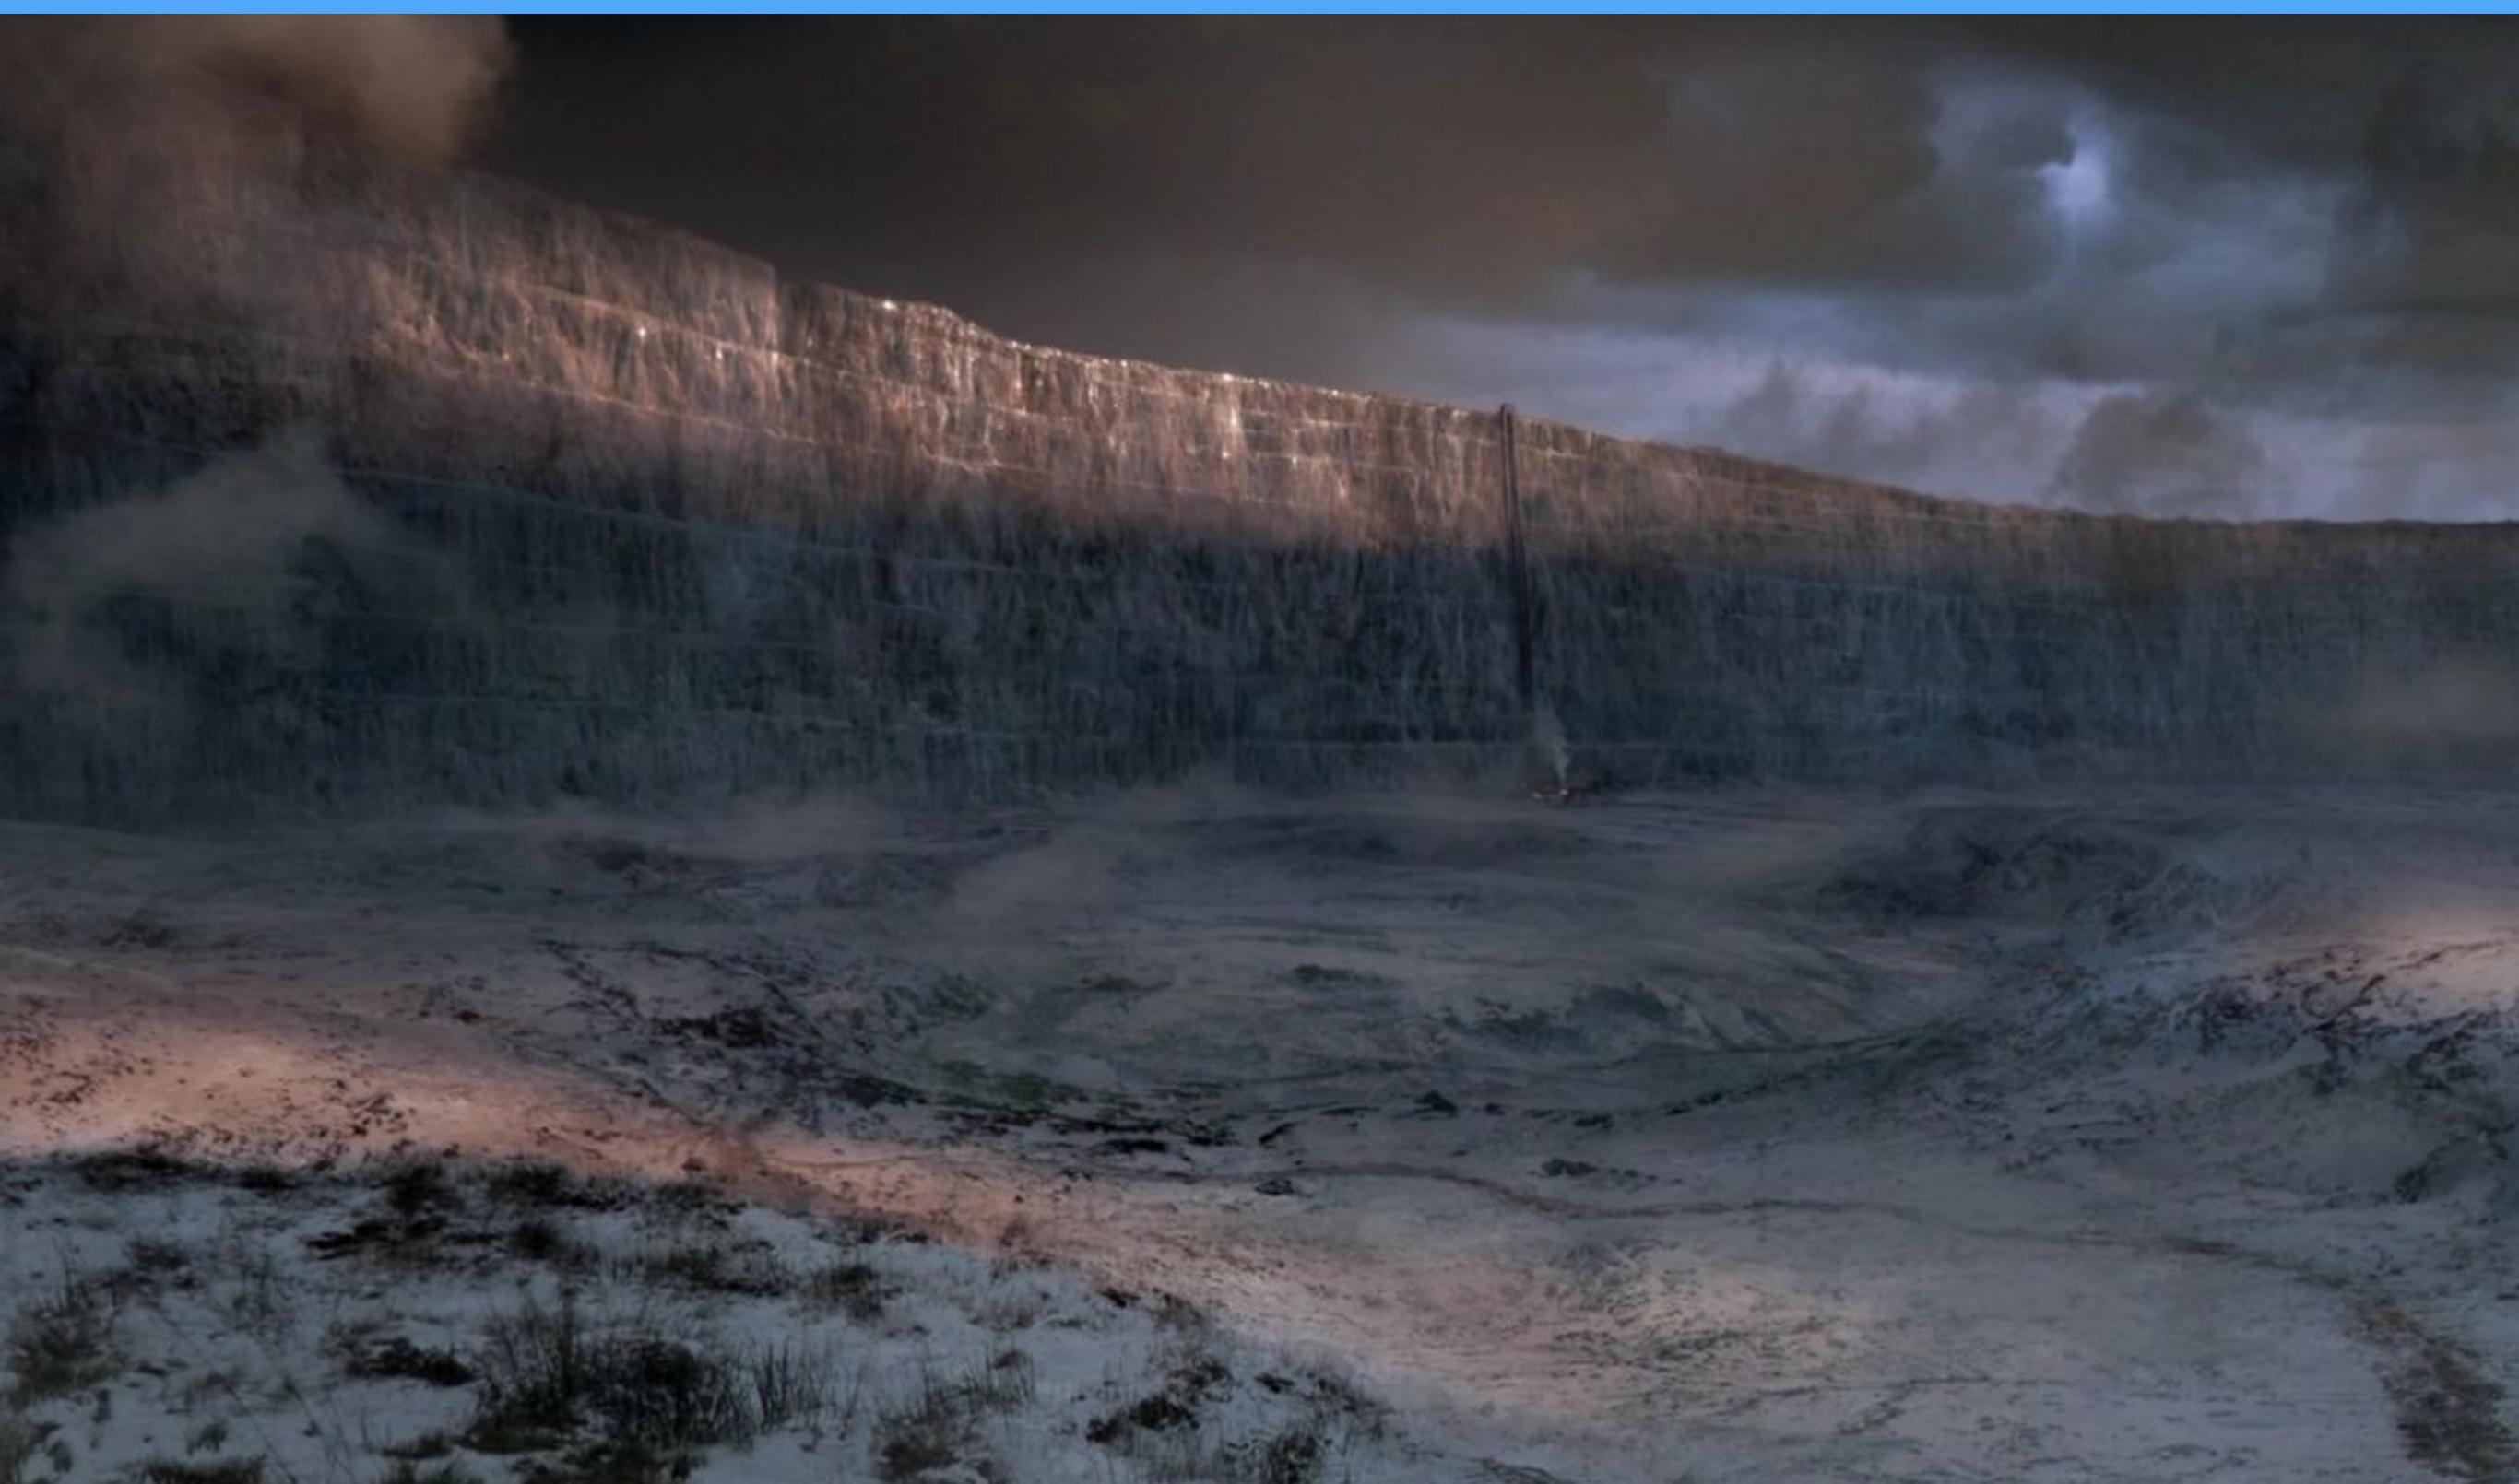


Apply Promise Theory

A portrait of a man with short brown hair and a light beard, wearing a dark blue turtleneck sweater. He is looking directly at the camera with a neutral expression. The background is a soft-focus view of what appears to be a modern interior space with blue and purple lighting.

AN AUTONOMOUS SERVICE
CAN ONLY PROMISE
ITS OWN BEHAVIOR

We need to decompose the system using
Consistency Boundaries



Inside Data

OUR CURRENT PRESENT—STATE

Outside Data

BLAST FROM THE PAST—FACTS

Between Services

HOPE FOR THE FUTURE—COMMANDS

**WITHIN the Consistency Boundary
we can have STRONG CONSISTENCY**



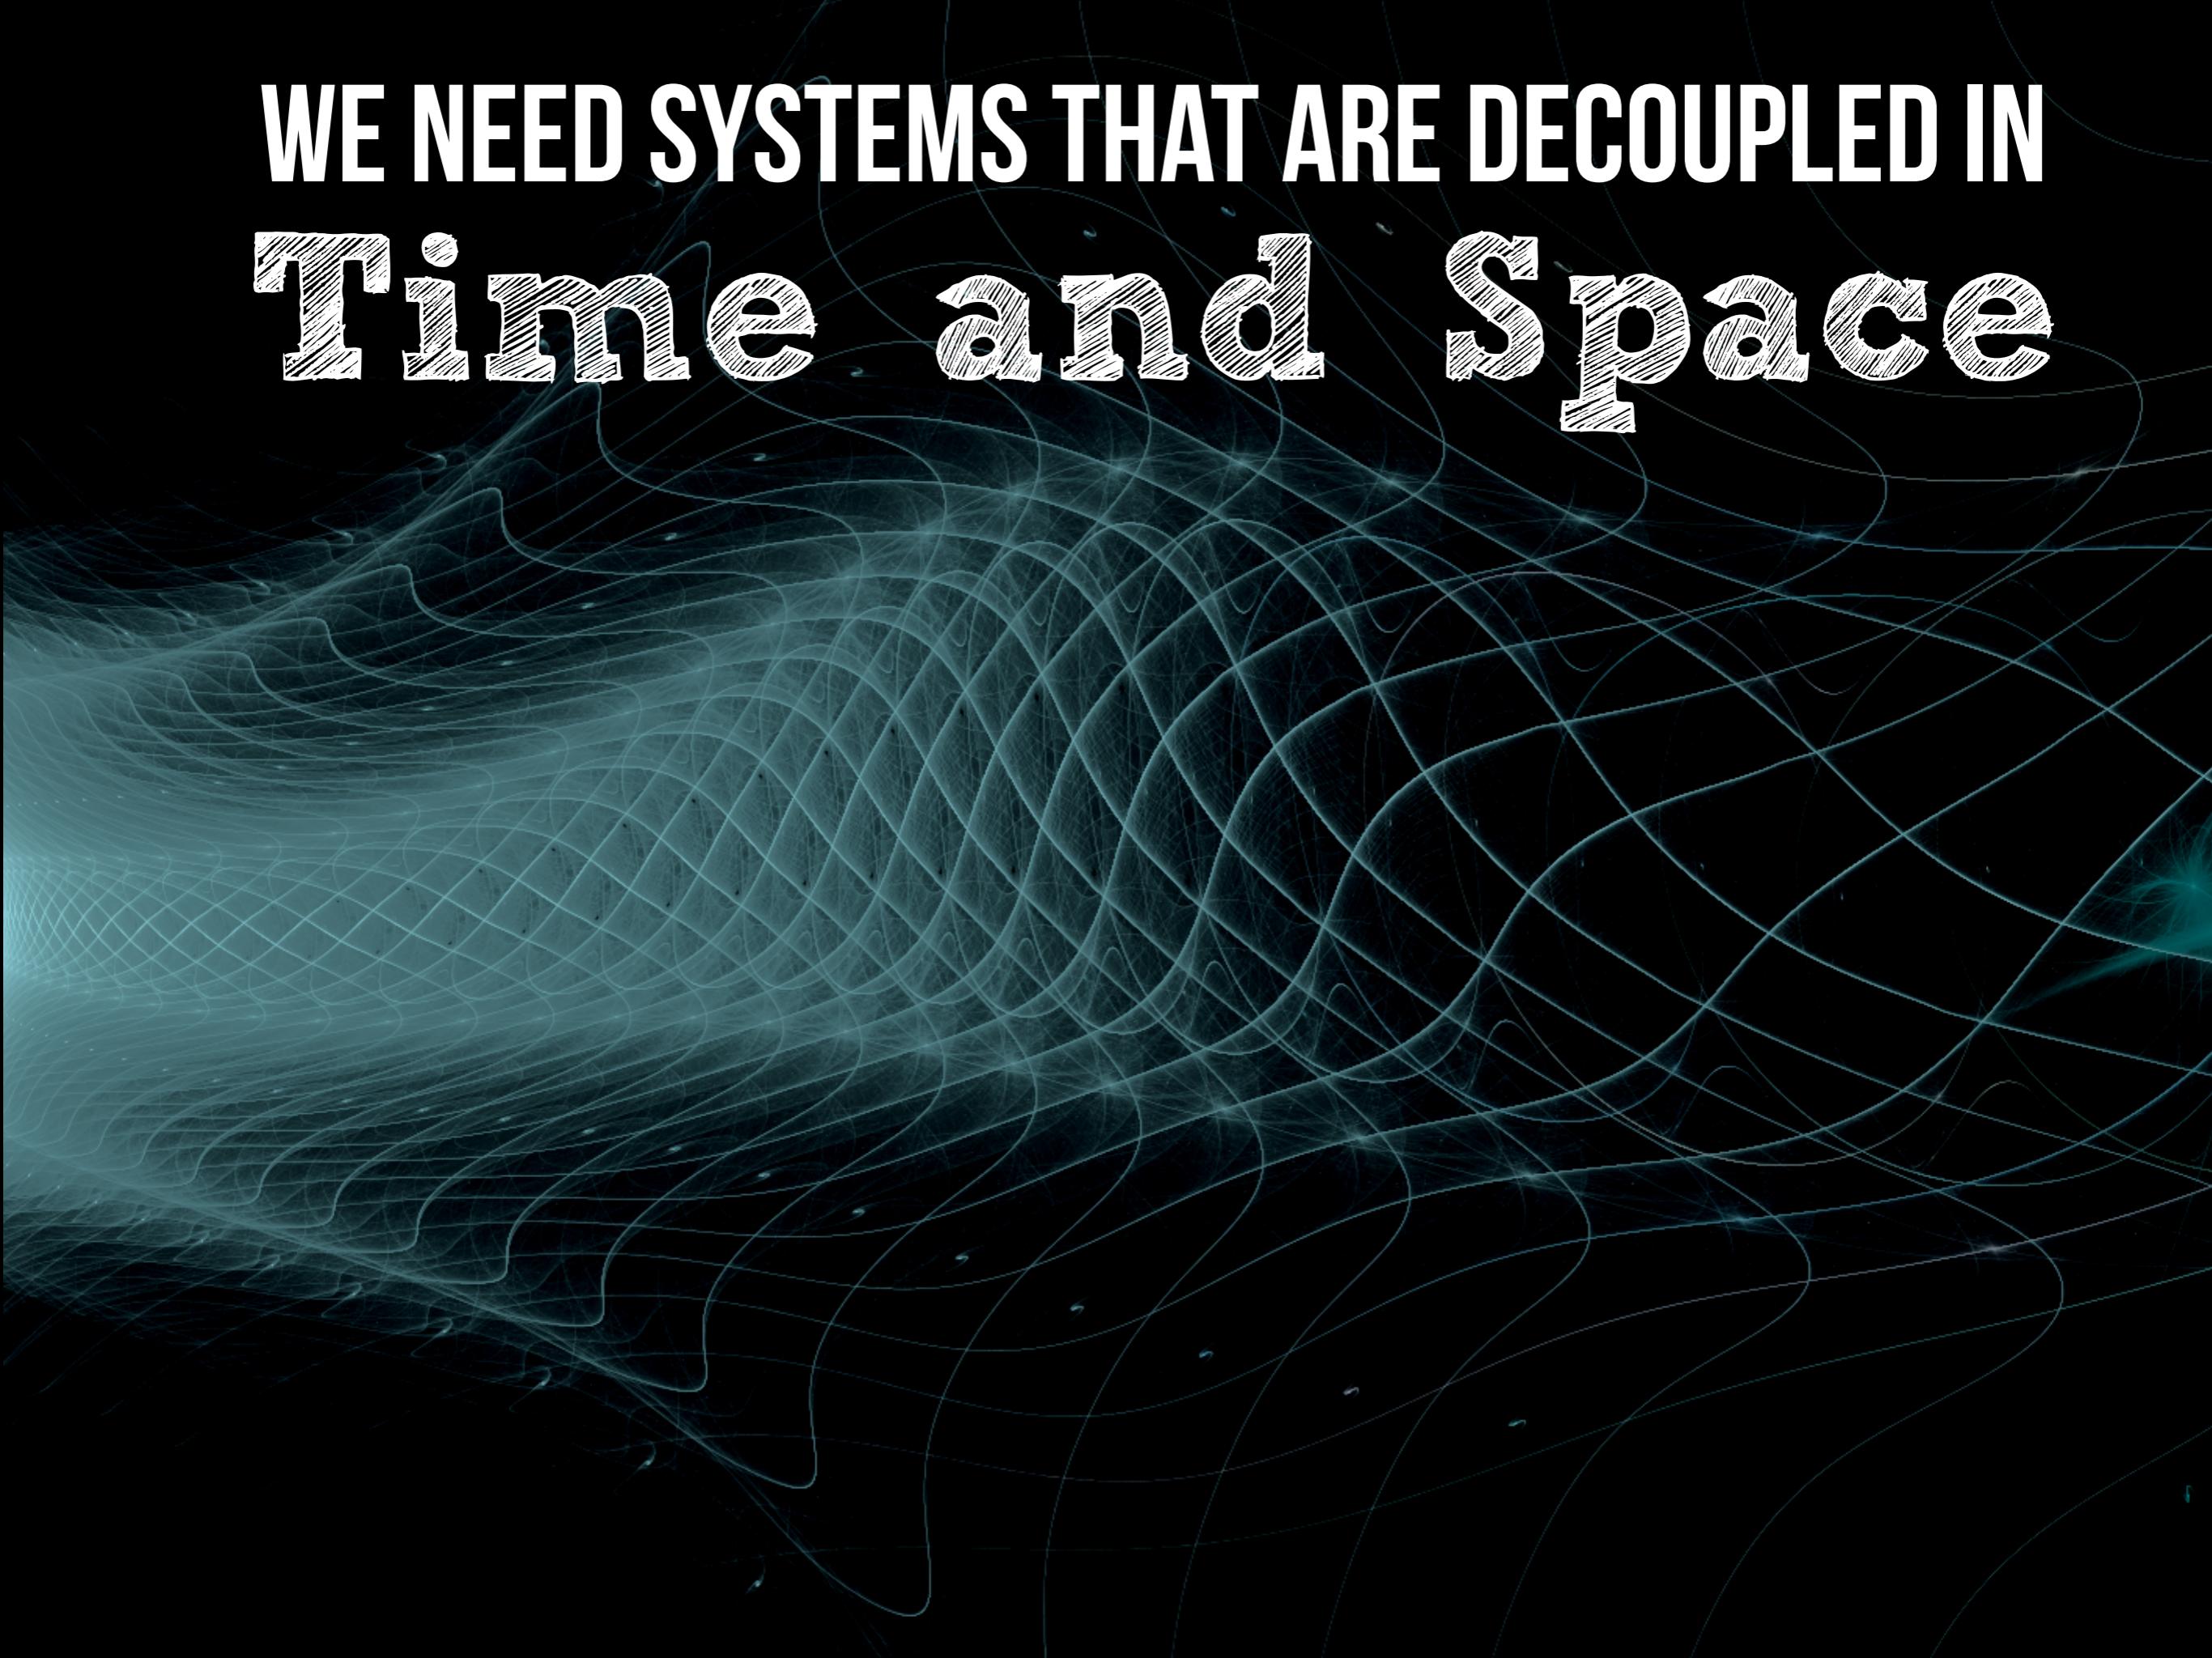
BETWEEN

Consistency
Boundaries
it is a

ZOO



WE NEED SYSTEMS THAT ARE DECOUPLED IN
time and space



Embrace the Network

- GO ASYNCHRONOUS
- MAKE DISTRIBUTION FIRST CLASS
- LEARN FROM THE MISTAKES OF RPC, EJB & CORBA
- LEVERAGE LOCATION TRANSPARENCY
- ACTOR MODEL DOES IT RIGHT

Location Transparency

ONE COMMUNICATION ABSTRACTION
ACROSS ALL DIMENSIONS OF SCALE

CORE \Rightarrow SOCKET \Rightarrow CPU \Rightarrow

CONTAINER \Rightarrow SERVER \Rightarrow RACK \Rightarrow

DATA CENTER \Rightarrow GLOBAL

Resilient Protocols

DEPEND ON

- ASYNCHRONOUS COMMUNICATION
- EVENTUAL CONSISTENCY

ARE TOLERANT TO

- MESSAGE LOSS
- MESSAGE REORDERING
- MESSAGE DUPLICATION

EMBRACE ACID 2.0

- ASSOCIATIVE
- COMMUTATIVE
- IDEMPOTENT
- DISTRIBUTED

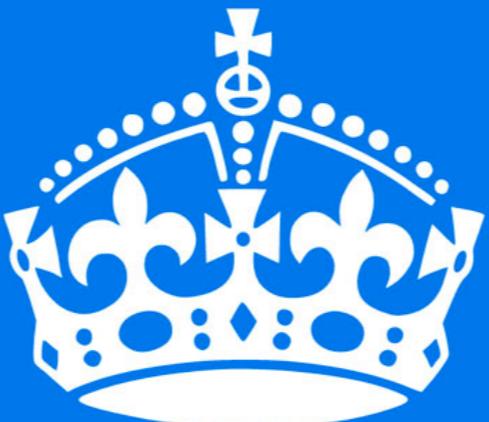
“To make a system of interconnected components crash-only, it must be designed so that components can tolerate the crashes and temporary unavailability of their peers. This means we require: [1] strong modularity with relatively impermeable component boundaries, [2] timeout-based communication and lease-based resource allocation, and [3] self-describing requests that carry a time-to-live and information on whether they are idempotent.”

- GEORGE CANDEA, ARMANDO FOX

"Software components should be designed such that they can deny service for any request or call. Then, if an underlying component can say No, apps must be designed to take No for an answer and decide how to proceed: give up, wait and retry, reduce fidelity, etc."

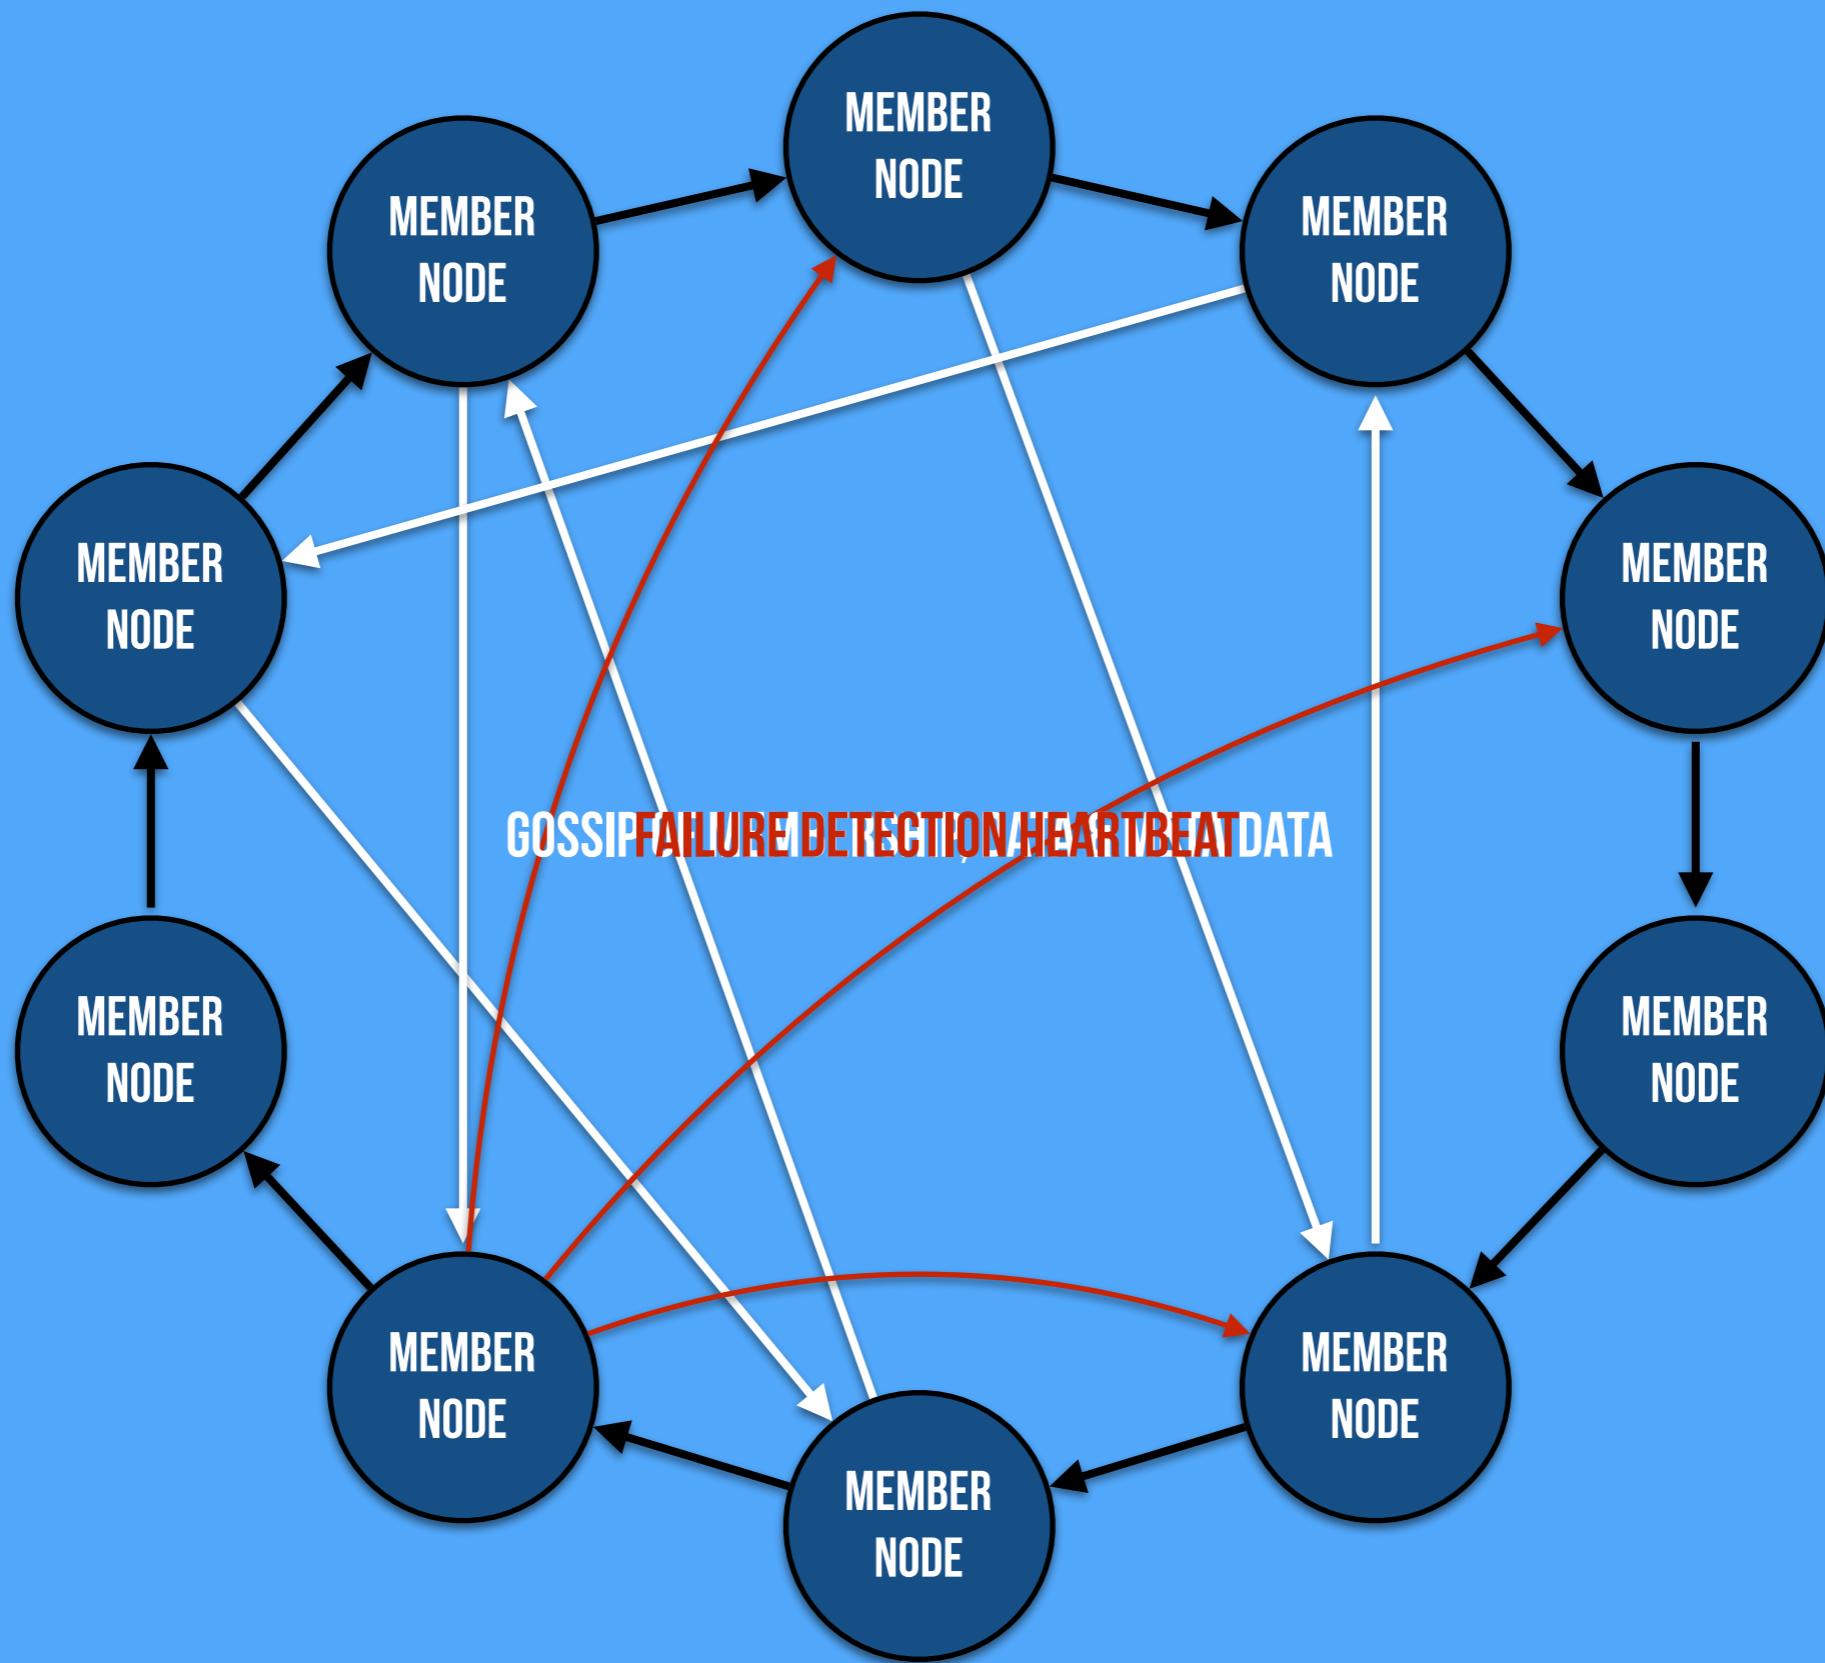
- GEORGE CANDEA, ARMANDO FOX

Services need to learn to accept
NO for an answer



KEEP
CALM
AND
SAY
NO

Decentralized Epidemic Gossip Protocols





**STRONG
Consistency
IS THE WRONG DEFAULT**

**“Two-phase commit is the
anti-availability protocol.”**

- PAT HELLAND

WE HAVE TO RELY ON
**Eventual
Consistency**
BUT RELAX, IT'S HOW THE WORLD WORKS

But I really need
Transactions

**“In general, application developers
simply do not implement large
scalable applications assuming
distributed transactions.”**

- PAT HELLAND

USE A PROTOCOL OF
**Guess.
Apologize.
Compensate.**

**“The truth is the log. The database is a
cache of a subset of the log.”**

- PAT HELLAND

CRUD is DEAD



Event Logging

- WORK WITH FACTS—IMMUTABLE VALUES
- EVENT SOURCING
- DB OF FACTS—KEEP ALL HISTORY
- JUST REPLAY ON FAILURE
- FREE AUDITING, DEBUGGING, REPLICATION
- SINGLE WRITER PRINCIPLE
- AVOIDS OO-RELATIONAL IMPEDIMENT MISMATCH
- CQRS—SEPARATE THE READ & WRITE MODEL

Demo
Time

LET'S MODEL A RESILIENT & EVENT LOGGED VENDING MACHINE, IN AKKA

Event Logged CoffeeMachine

```
// Events
case class CoinsReceived(number: Int)

class CoffeeMachine extends PersistentActor {
    val price = 2
    var nrOfInsertedCoins = 0
    var outOfCoffeeBeans = false
    var totalNrOfCoins = 0

    override def persistenceId = "CoffeeMachine"

    override def receiveCommand: Receive = {
        case Coins(nr) =>
            nrOfInsertedCoins += nr
            println(s"Inserted [$nr] coins")
            persist(CoinsReceived(nr)) { evt =>
                totalNrOfCoins += nr
                println(s"Total number of coins in machine is [$totalNrOfCoins]")
            }
        ...
    }

    override def receiveRecover: Receive = {
        case CoinsReceived(coins) =>
            totalNrOfCoins += coins
            println(s"Total number of coins in machine is [$totalNrOfCoins]")
    }
}
```

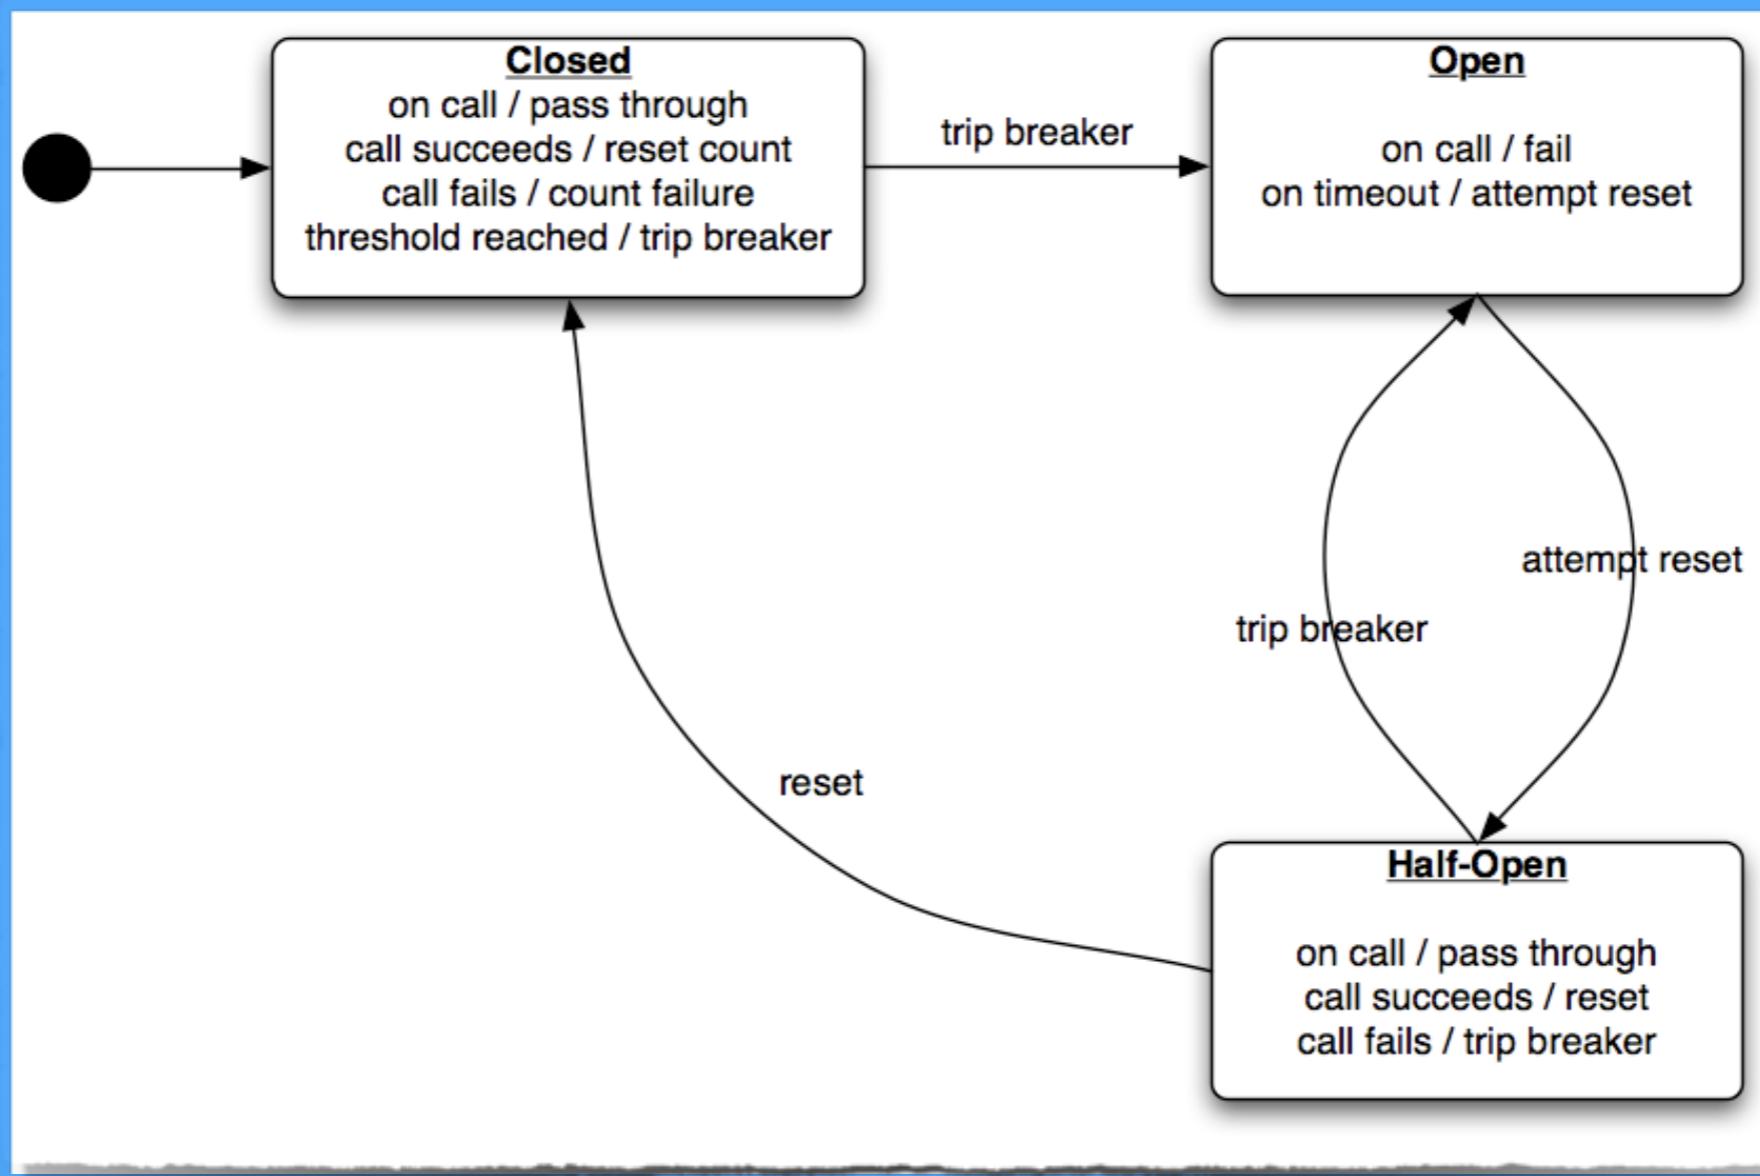
<https://gist.github.com/jboner/1db37eeee3ed3c9422e4>

**“An escalator can never break: it can only
become stairs. You should never see an
Escalator Temporarily Out Of Order sign, just
Escalator Temporarily Stairs. Sorry for the
convenience.”**

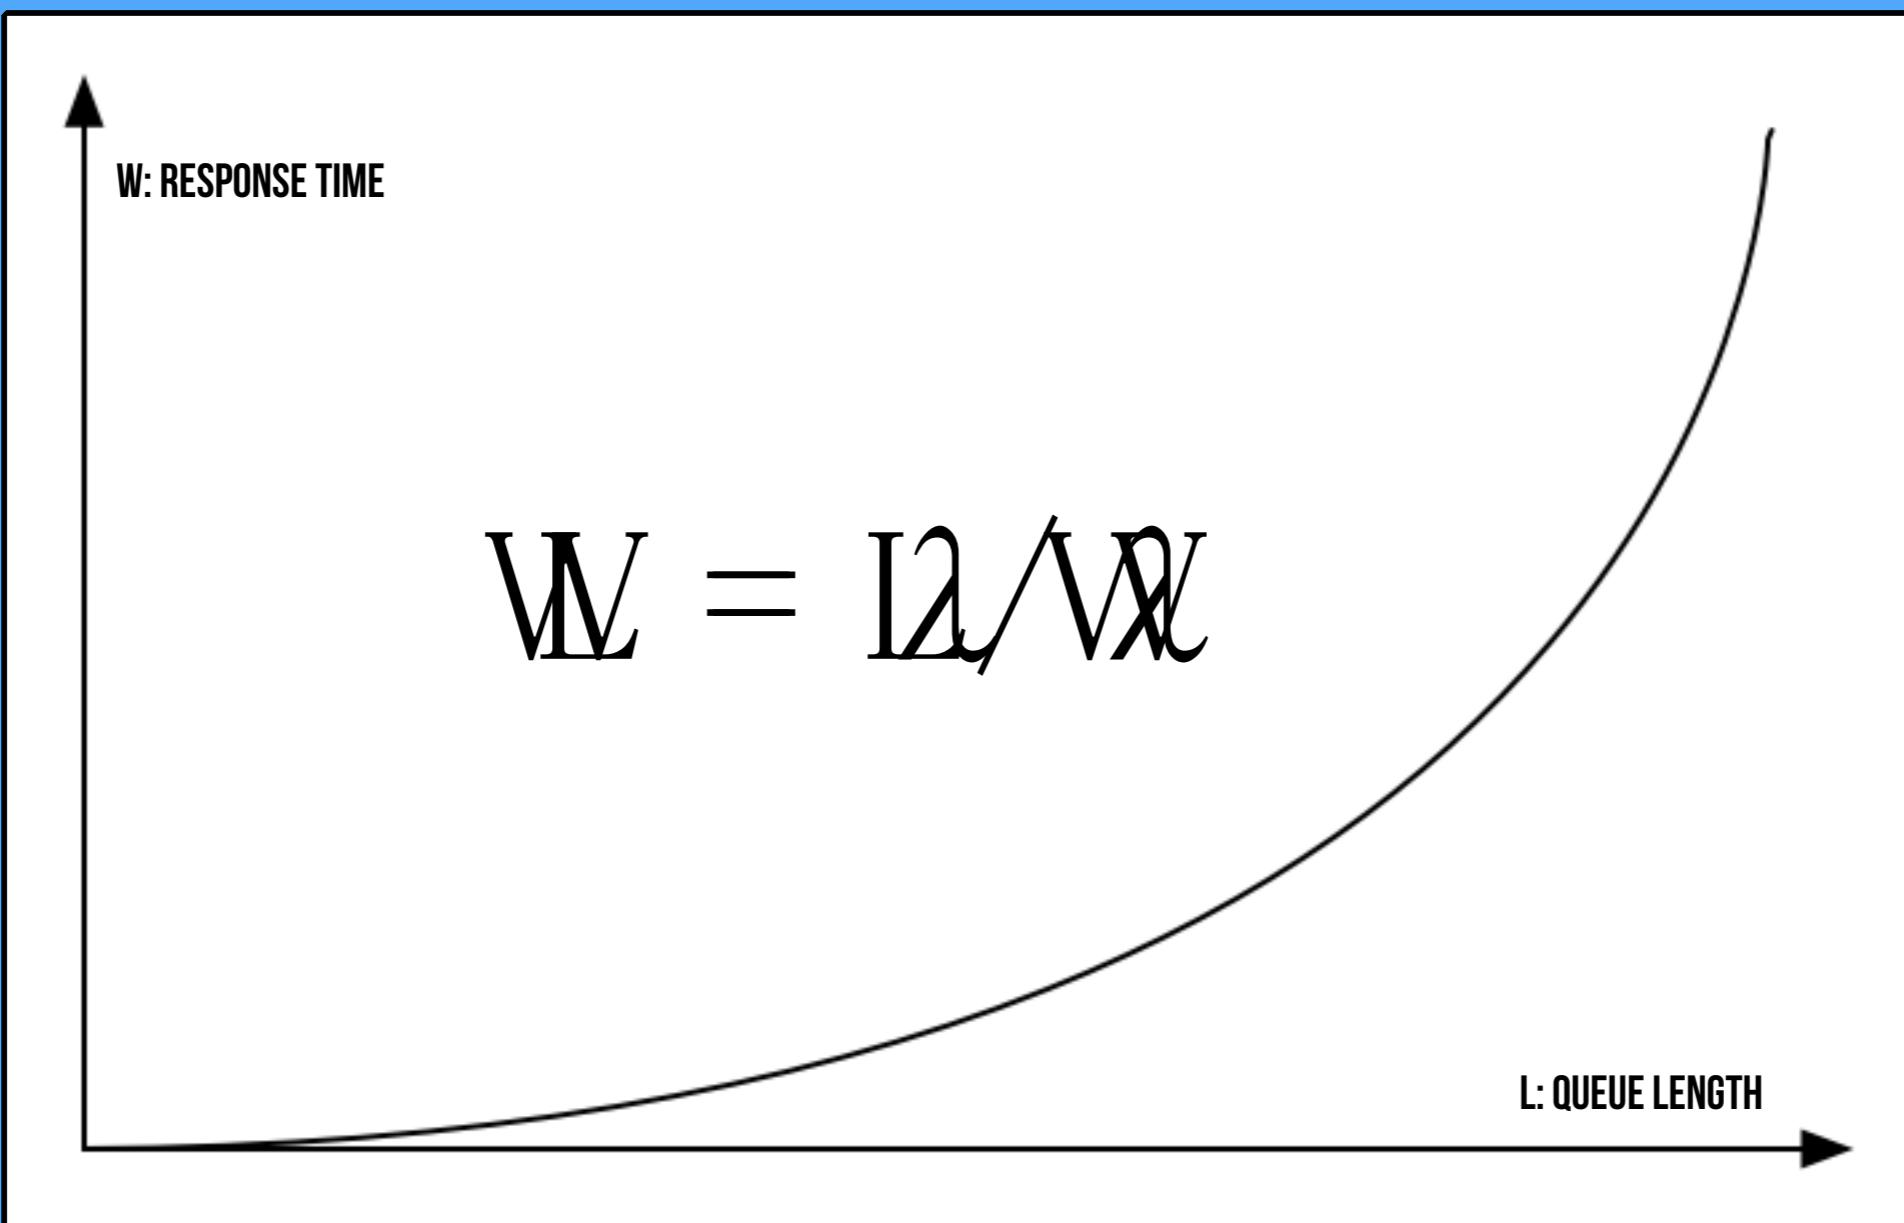
- MITCH HEDBERG

Graceful
Degradation

Circuit Breaker



Little's Law



RESPONSE TIME = QUEUE LENGTH / RESPONSE RATE

Flow Control

ALWAYS APPLY BACKPRESSURE



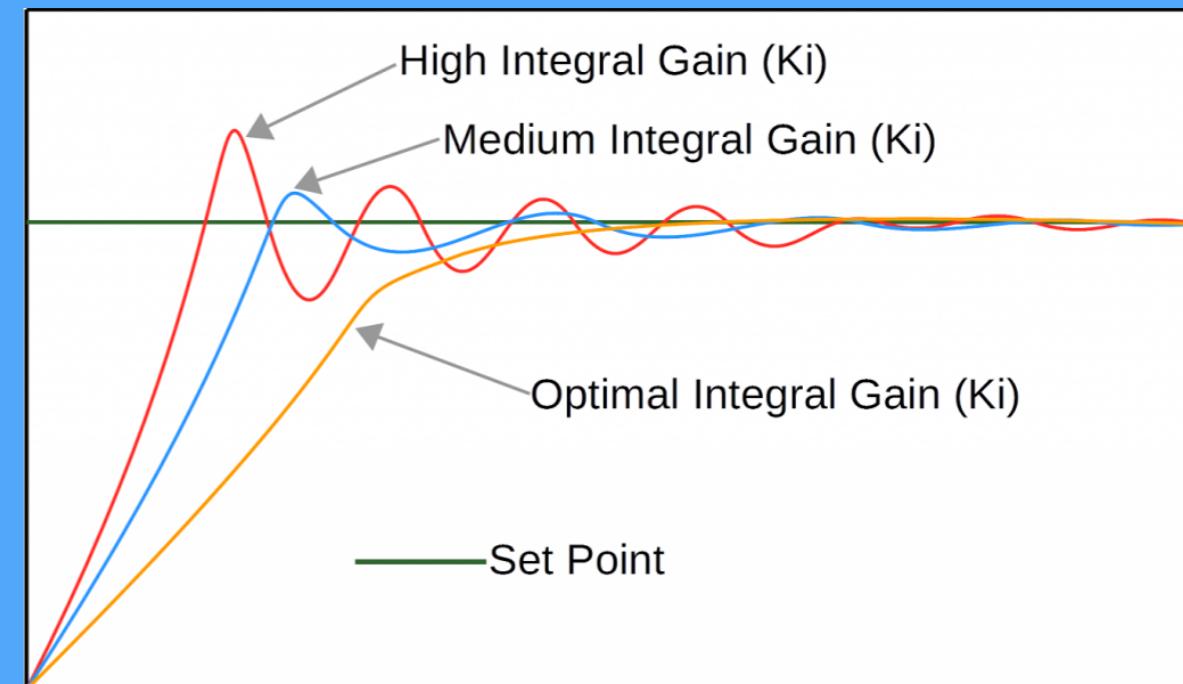
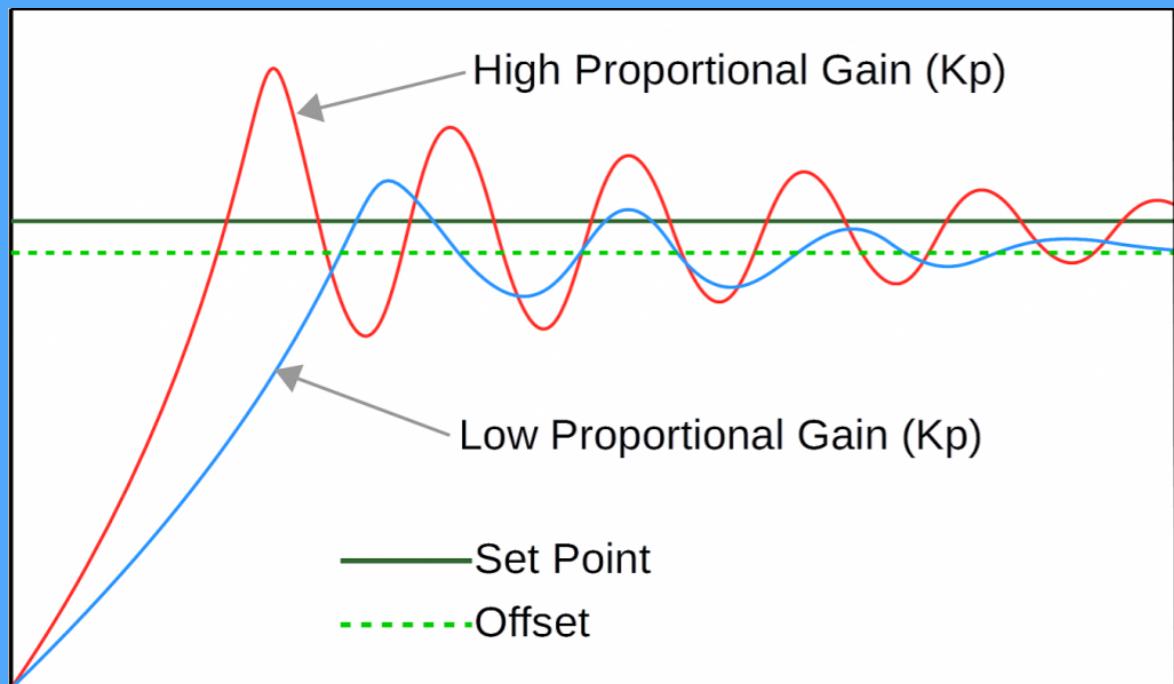
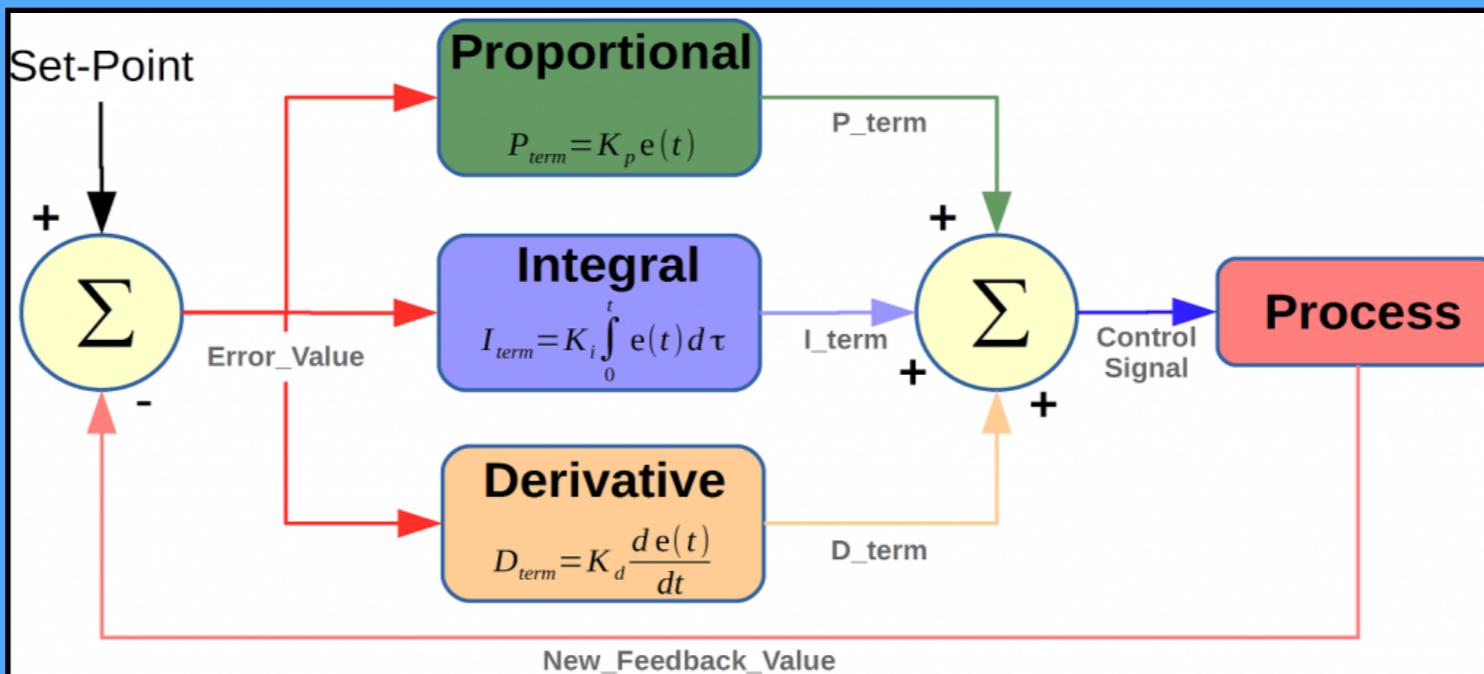
Feedback Control

THE FEEDBACK PRINCIPLE

“Continuously compare the actual output to its desired reference value; then apply a change to the system inputs that counteracts any deviation of the actual output from the reference.”

- PHILIPP K. JANERT

Feedback Control



Influencing a Complex System

Places to Intervene in a Complex System

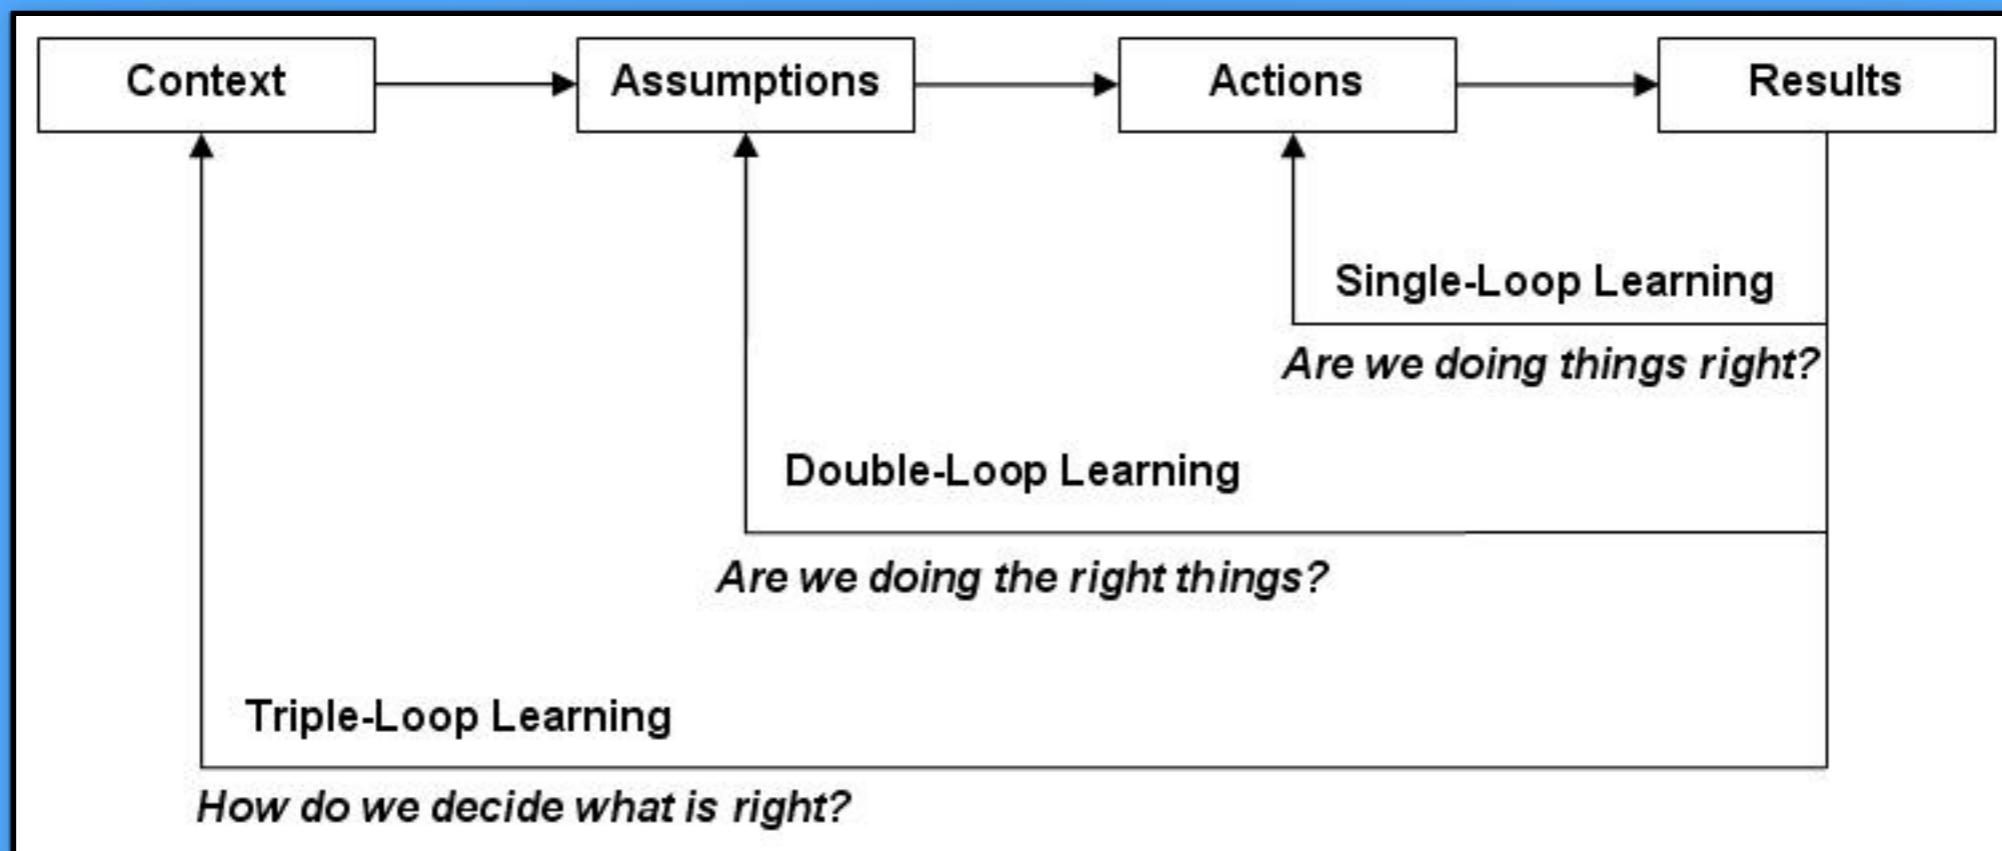
1. The constants, parameters or numbers
2. The sizes of buffers relative to their flows
3. The structure of material stocks and flows
4. The lengths of delays, relative to the rate of system change
5. The strength of negative feedback loops
6. The gain around driving positive feedback loops
7. The structure of information flows
8. The rules of the system
9. The power to add, change, evolve, or self-organize structure
10. The goals of the system
11. The mindset or paradigm out of which the system arises
12. The power to transcend paradigms

Triple Loop Learning

LOOP 1: FOLLOW THE RULES

LOOP 2: CHANGE THE RULES

LOOP 3: LEARN HOW TO LEARN



Testing



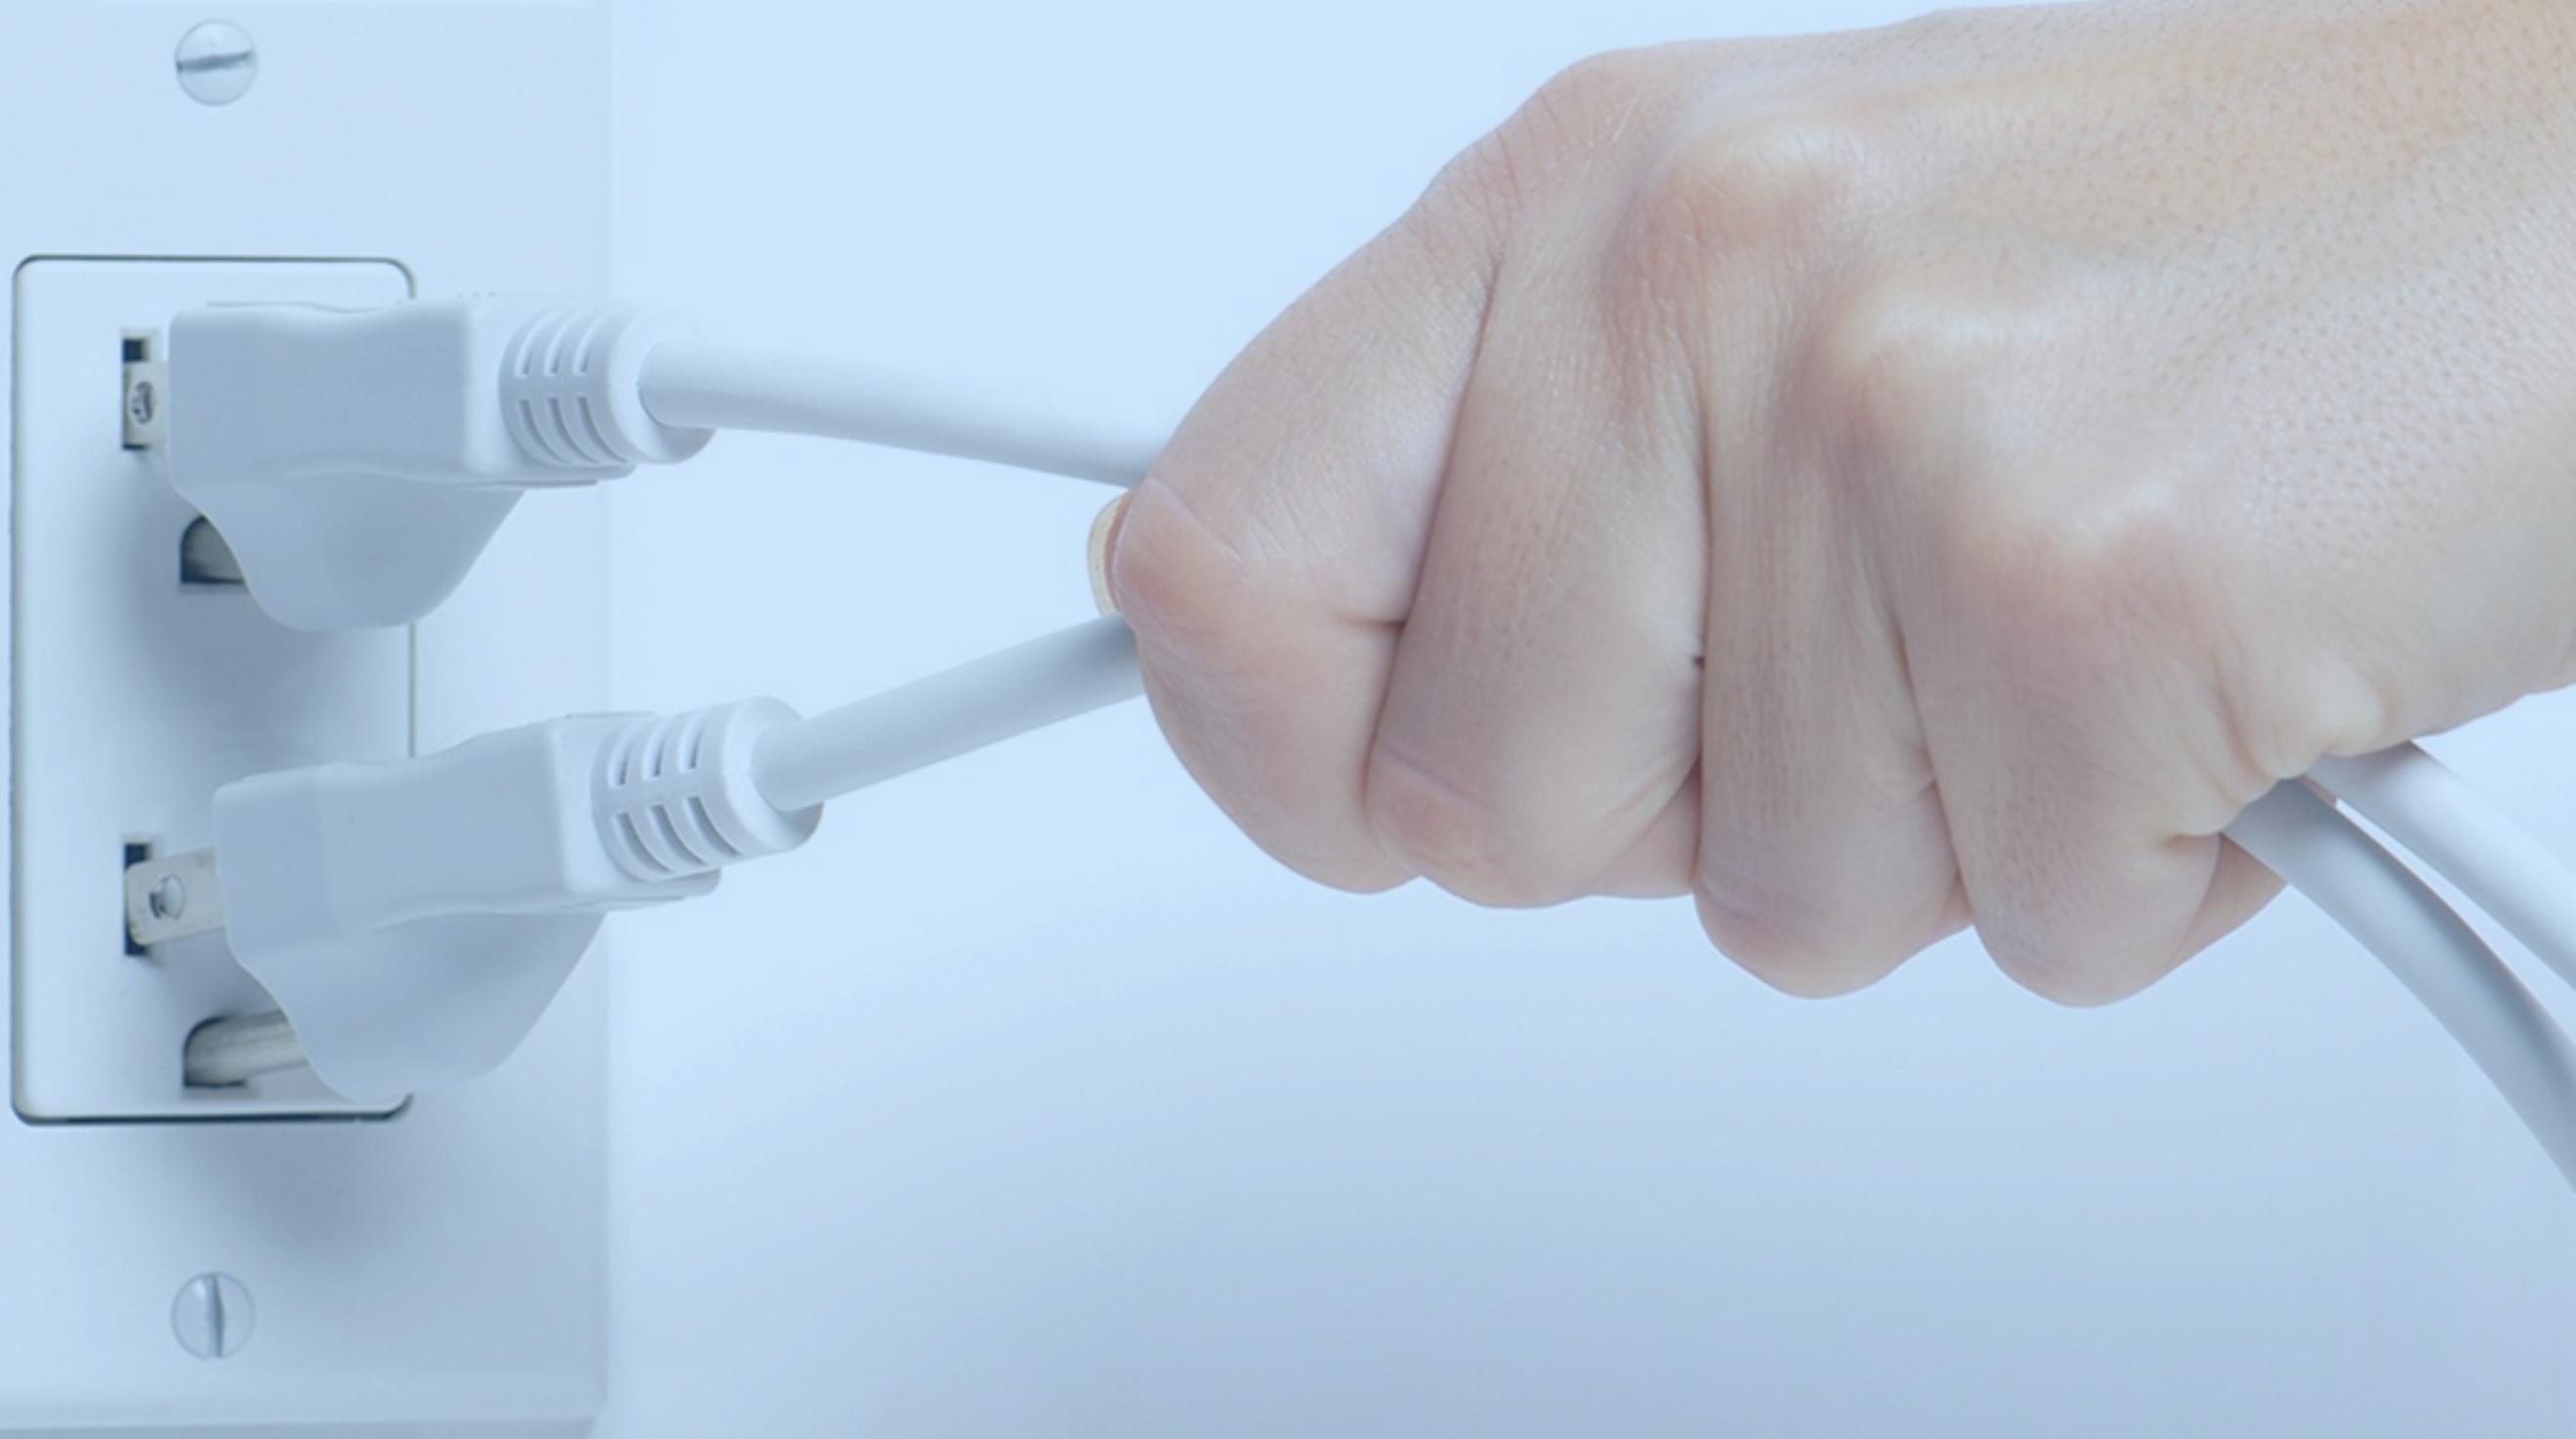
WHAT CAN WE LEARN FROM ARNOLD?



BLOW THINGS UP

**Shoot
Your App
Down**





Pull the Plug
...AND SEE WHAT HAPPENS



Executive Summary

“Complex systems run as broken systems.”

- RICHARD COOK

Resilience
is by
Design



Photo courtesy of FEMA/Joselyne Augustino

**Without Resilience
Nothing Else Matters**

References

- Drift into Failure - <http://www.amazon.com/Drift-into-Failure-Components-Understanding-ebook/dp/B009KOKXKY>
- How Complex Systems Fail - <http://web.mit.edu/2.75/resources/random/How%20Complex%20Systems%20Fail.pdf>
- Leverage Points: Places to Intervene in a System - <http://www.donellameadows.org/archives/leverage-points-places-to-intervene-in-a-system/>
- Going Solid: A Model of System Dynamics and Consequences for Patient Safety - <http://www.ncbi.nlm.nih.gov/pmc/articles/PMC1743994/>
- Resilience in Complex Adaptive Systems: Operating at the Edge of Failure - <https://www.youtube.com/watch?v=PGLYEDpNu60>
- Puppies! Now that I've got your attention, Complexity Theory - https://www.ted.com/talks/nicolas_perony_puppies_now_that_i_ve_got_your_attention_complexity_theory
- How Bacteria Becomes Resistant - <http://www.abc.net.au/science/slab/antibiotics/resistance.htm>
- Towards Resilient Architectures: Biology Lessons - <http://www.metropolismag.com/Point-of-View/March-2013/Toward-Resilient-Architectures-1-Biology-Lessons/>
- Dealing in Security - http://resiliencemaps.org/files/Dealing_in_Security.July2010.en.pdf
- What is resilience? An introduction to social-ecological research - http://www.stockholmresilience.org/download/18.10119fc11455d3c557d6d21/1398172490555/SU_SRC_whatisresilience_sidaApril2014.pdf
- Applying resilience thinking: Seven principles for building resilience in social-ecological systems - <http://www.stockholmresilience.org/download/18.10119fc11455d3c557d6928/1398150799790/SRC+Applying+Resilience+final.pdf>
- Crash-Only Software - https://www.usenix.org/legacy/events/hotos03/tech/full_papers/candea/candea.pdf
- Recursive Restartability: Turning the Reboot Sledgehammer into a Scalpel - http://roc.cs.berkeley.edu/papers/recursive_restartability.pdf
- Out of the Tar Pit - <http://citeseerx.ist.psu.edu/viewdoc/summary?doi=10.1.1.93.8928>
- Bulkhead Pattern - <http://skife.org/architecture/fault-tolerance/2009/12/31/bulkheads.html>
- Making Reliable Distributed Systems in the Presence of Software Errors - http://www.erlang.org/download/armstrong_thesis_2003.pdf
- On Erlang, State and Crashes - <http://jlouisramblings.blogspot.be/2010/11/on-erlang-state-and-crashes.html>
- Akka Supervision - <http://doc.akka.io/docs/akka/snapshot/general/supervision.html>
- Release It!: Design and Deploy Production-Ready Software - <https://pragprog.com/book/mnnee/release-it>
- Feedback Control for Computer Systems - <http://www.amazon.com/Feedback-Control-Computer-Systems-Philipp/dp/1449361692>
- The Network in Reliable - <http://queue.acm.org/detail.cfm?id=2655736>
- Data on the Outside vs Data on the Inside - <https://msdn.microsoft.com/en-us/library/ms954587.aspx>
- Life Beyond Distributed Transactions - <http://adrianmarriott.net/logosroot/papers/LifeBeyondTxns.pdf>
- Immutability Changes Everything - http://cidrdb.org/cidr2015/Papers/CIDR15_Paper16.pdf
- Standing on Distributed Shoulders of Giants - <https://queue.acm.org/detail.cfm?id=2953944>
- Thinking in Promises - <http://shop.oreilly.com/product/0636920036289.do>
- In Search Of Certainty - <http://shop.oreilly.com/product/0636920038542.do>
- Reactive Microservices Architecture - <http://www.oreilly.com/programming/free/reactive-microservices-architecture-orm.csp>
- Reactive Streams - <http://reactive-streams.org>
- Vending Machine Akka Supervision Demo - <https://gist.github.com/jboner/d24c0eb91417a5ec10a6>
- Persistent Vending Machine Akka Supervision Demo - <https://gist.github.com/jboner/1db37eeee3ed3c9422e4>

thank
you

Without Resilience Nothing Else Matters

JONAS BONÉR
CTO Lightbend
@jboner